

A Year Fuzzing XNU Mach IPC



Nguyen Vu Hoang (Peter)

About me

Nguyen Vu Hoang (Peter), Security Researcher
at STAR Labs.

Focusing on macOS/iOS bug hunting and
exploitation.

 Awarded bounties by Apple Security
Platform.

 Found some vulnerabilities in Apple from
userland to kernel-level vulnerabilities

Outline

1. Previous XNU Mach IPC Vulnerabilities
2. Template-based XNU IPC Fuzzer
3. CVE-2022-32894
4. Q&A
5. References

1. Previous XNU Mach IPC Vulnerabilities



imgflip.com

CVE-2020-27932 (1/2)

Issue 2107: XNU kernel type confusion in turnstiles

Reported by ianbeer@google.com on Fri, Oct 30, 2020, 2:56 AM GMT+8

Project Member

```
union {
    ipc_kobject_t kobject;           <-- a kobject (like iokit userclient, host_notify port)
    ipc_importance_task_t imp_task;
    ipc_port_t sync_inheritor_port;   <-- the destination port to which a special_reply_port was sent
    struct knote *sync_inheritor_knote;
    struct turnstile *sync_inheritor_ts;
} kdata;
-----,
} kdata;
```

There are a couple of weird edges cases which the turnstile code overlooked; specifically it's possible to turn a regular, user-owned port into a port with a kobject! Specifically, by calling `host_request_notification` on a `send_once` right we can change the type of a port to `IKOT_HOST_NOTIFY` and get the `kobject` field set to a `struct host_notify_entry` pointer.

CVE-2020-27932 (2/2)

```
101         entry = (host_notify_t)zalloc(host_notify_zone);
102         if (entry == NULL) {
103             return KERN_RESOURCE_SHORTAGE;
161    95             ip_lock(port);
161    96             if (!ip_active(port) || port->ip_tempowner || port->ip_specialreply || Added the check
161    97                 ip_is_kolabeled(port) || ip_kotype(port) != IKOT_NONE) {
161    98                 ip_unlock(port);
161    99
161   100                 lck_mtx_unlock(&host_notify_lock);
161   101                 zfree(host_notify_zone, entry);
161   102
161   103             return KERN_FAILURE;
162   104         }
162 117         entry->port = port; assign port->kobject with entry pointer
118         ipc_kobject_set_atomically(port, (ipc_kobject_t)entry, IKOT_HOST_NOTIFY);
119         ip_unlock(port);
120     }
```

CVE-2021-1782

```
...  
switch (command) {  
case MACH_VOUCHER_ATTR_REDEEM:  
  
    /* redeem of previous values is the value */  
    if (0 < prev_value_count) {  
        elem = (user_data_element_t)prev_values[0];  
  
        user_data_lock();  
        assert(0 < elem->e_made);  
        elem->e_made++;  
        user_data_unlock(); added the lock to prevent race condition  
  
        *out_value = (mach_voucher_attr_value_handle_t)elem;  
        return KERN_SUCCESS;  
    }  
}
```

Before the patch

patch in xnu 7195.81.3

Key takeaways

1. Mach IPC subsystems are easy.
 2. *The union keyword* can store multiple types.
- These could lead



features which is not
in mach port structure to
IU kernel.

2. Template-based XNU IPC Fuzzer



What are the three challenges ?

1. Generating test-case using **MACH IPC syscalls** and their arguments in C language is very hard.
 2. The **code-base is huge**. In order to **reach these code paths**, you have to **satisfy many conditions**.
 3. Mutating **Mach IPC syscalls** and their arguments are not enough to expand the code **coverage** and hard to develop.
- **We need more advanced fuzzer which has ability to generate complex test-case and reach in deeper in Mach IPC.**



What are template-based fuzzer?

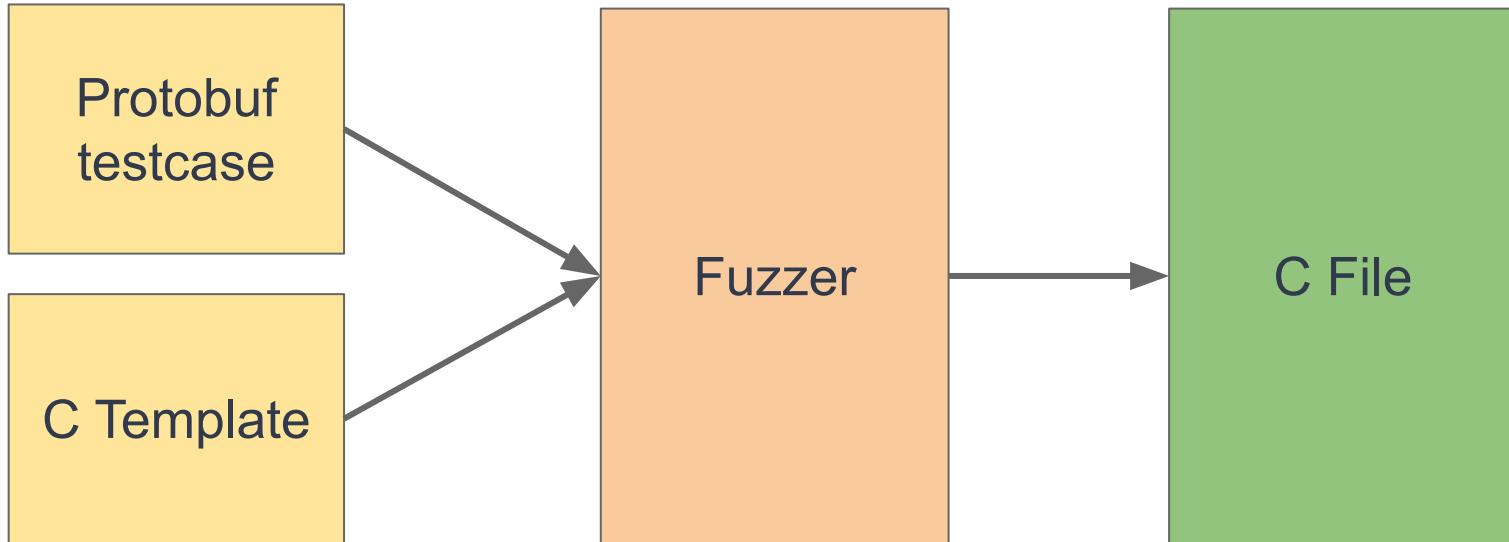


My expectations

- The template based fuzzer have to **easily to modify and edit template in C language**.
- The template design must be **easy to parse and process** (less coding 😊).
- The template able to add some **custom functions** beside mach ipc syscall functions.



Fuzzer Architecture



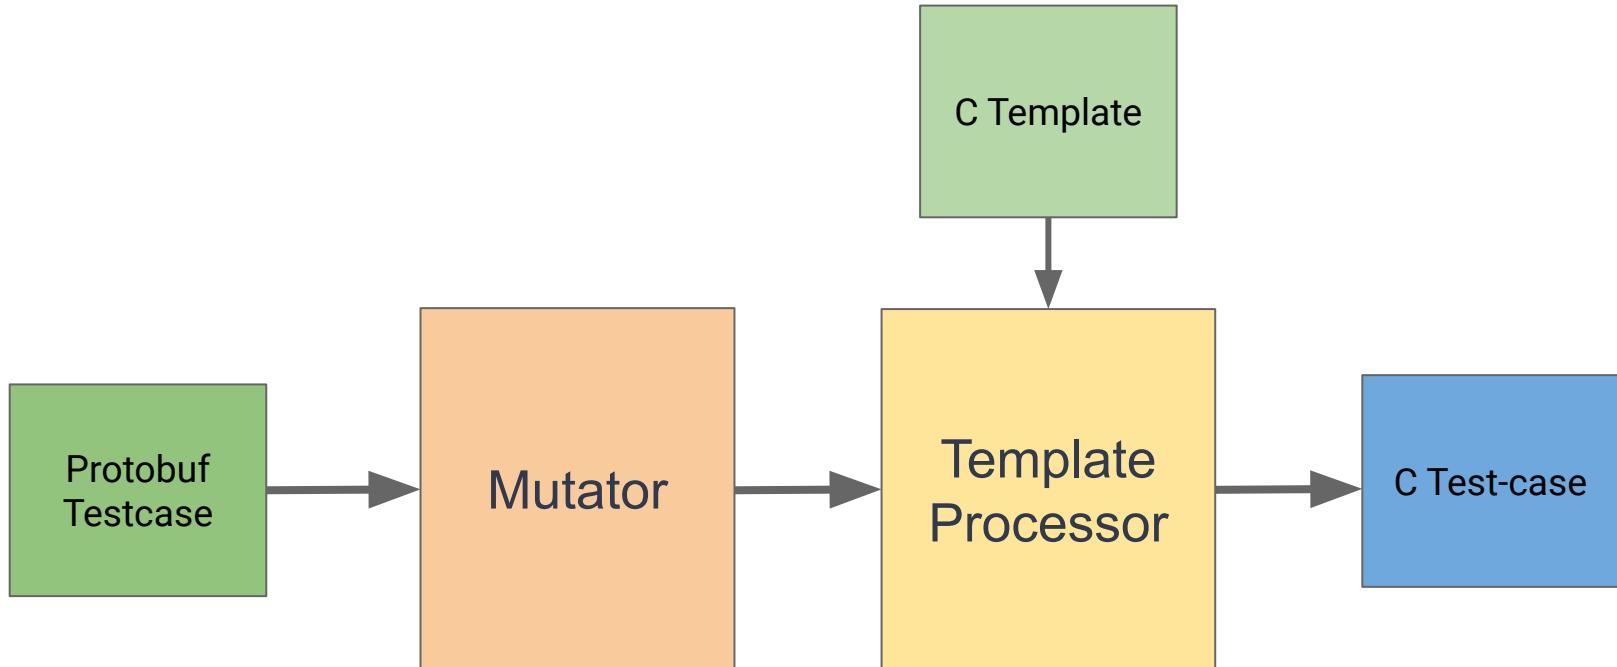
Fuzzer Architecture: Protobuf test-case

```
102     / 49 ✓ message mach_msg {  
103       m 50     string msgh_bits_remote = 1;  
104       m 51     string msgh_bits_local = 2;  
105       m 52     string msgh_bits_voucher = 3;  
106       m 53     string msgh_bits_other = 4;  
107       m 54     uint64 msgh_size = 5;  
108       m 55     uint32 msgh_remote_port = 6;  
109       m 56     uint32 msgh_local_port = 7;  
110       m 57     uint64 msgh_voucher_port = 8;  
111       m 58     int32  msgh_id = 9;                                value = 7;  
112       m 59     uint64 msgh_descriptor_count = 10;  
113       m 60     repeated mach_msg_descriptor descriptors = 11;  
114   } 61     bytes inline_data = 12;  
115 } 62 }
```

Fuzzer Architecture: C Template format

```
86     kern_return_t create_user_data_voucher(mach_port_t *port)
● 196     /*port_declare*/
● 197
● 198     void execute()
● 199     {
● 200         /*GENERATE_CODE*/
● 201     }
● 202
● 203     int main()
● 204     {
● 205         mach_msg_type_number_t my_size;
● 206         /*port_init*/
● 207
● 208         execute();
● 209
● 210         return 0;
● 211     }
```

Fuzzer Architecture: The Fuzzer



Template Processor (1/2)

```
1 /*TEMPLATE_INFO
"defined_functions" : {
  "create_importance_voucher" : {
    "args" : ["port_out:voucher_imp"], → type of this output mach port
    "ret" : "kern_return_t"
  },
  "create_user" : {
    "args" : [],
    "ret" : "kern_return_t"
  },
  "set_thread_"
  "args" : [],
  "ret" : "kern_return_t"
},
"create_bank_voucher" : {
  "args" : ["port_out:voucher_bank"],
  "ret" : "kern_return_t"
},
"custom_send" : {
  "args" : ["port_in:!mk_timer", "port_in:special_reply", "port_in:!mk_timer+!kobject+!pset", "mach_msg_option_one_send", "send_disp"], → types of mach ports are not allow in this argument
  "ret" : "kern_return_t"
},
```

Distinguish mach port by type

- ↳ correctly mutate arguments to reduce the error cases
 - ↳ increase more chances to reach deeper in kernel ⚡

types of mach ports are allow in this argument

types of mach ports are not allow in this argument

input must be mach port

Template Processor (2/2)

```
164 if template_info.aet('defined vars'):  
273 port_25 = thread_get_special_reply_port();  
274 port_26 = mach_thread_self();  
275 mach_port_destruct(task_port_1, port_26, -1, 4660318827490212632);  
276 kr = custom_send(port_3, port_25, port_2, port_3, MACH_SEND_SYNC_BOOTSTRAP_CHECKIN, MACH_MSG_TYPE_MAKE_SEND);  
277 printf("[%s:%d] custom_send: %d (%s)\n", __FUNCTION__, __LINE__, kr, mach_error_string(kr));  
278 kr = custom_send(port_2, port_25, port_2, port_16, MACH_SEND_PROPAGATE_QOS, MACH_MSG_TYPE_MOVE_SEND_ONCE);  
279 printf("[%s:%d] custom_send: %d (%s)\n", __FUNCTION__, __LINE__, kr, mach_error_string(kr));  
280 kr = mach_port_deallocate(task_port_1, port_9);  
281 printf("[%s:%d] mach_port_deallocate: %d (%s)\n", __FUNCTION__, __LINE__, kr, mach_error_string(kr));  
282 kr = thread_set_mach_voucher(mach_thread_self(), port_2);  
283 printf("[%s:%d] thread_set_mach_voucher: %d (%s)\n", __FUNCTION__, __LINE__, kr, mach_error_string(kr));  
284 basic_msg.header.msgh_remote_port = port_7;  
285 basic_msg.header.msgh_local_port = port_13;  
286 basic_msg.header.msgh_voucher_port = MACH_PORT_NULL;  
287 basic_msg.header.msgh_bits = MACH_MSGH_BITS_SET(MACH_MSG_TYPE_MOVE_SEND_ONCE, MACH_MSG_TYPE_MOVE_SEND_ONCE, 0, 0);  
288 basic_msg.header.msgh_id = 24801;  
289 basic_msg.header.msgh_size = sizeof(basic_msg) - sizeof(mach_msg_trailer_t);  
290 kr = mach_msg((mach_msg_header_t *)&basic_msg, MACH_SEND_MSG|MACH_SEND_TIMEOUT|MACH_RCV_MSG|MACH_RCV_TIMEOUT|MACH_RCV_LARGE_IDENTITY|  
    IFY|MACH_RCV_SYNC_WAIT|MACH_RCV_VOUCHER, basic_msg.header.msgh_size, sizeof(basic_msg), port_10, 4000, port_26);  
291 printf("[%s:%d] mach_msg: %d (%s)\n", __FUNCTION__, __LINE__, kr, mach_error_string(kr));  
 000
```

Mutator

Mutator strategies:

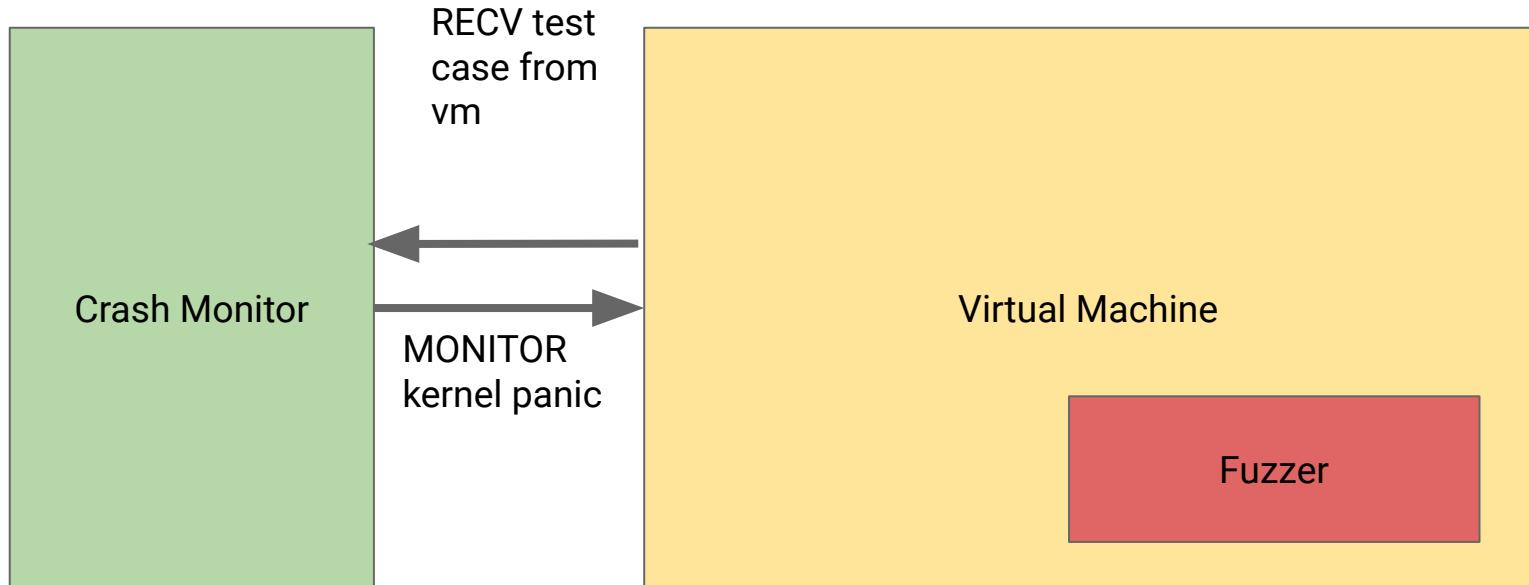
1. Mutating **function arguments** base on **their types** (defined in C Template).
2. Mutating the **order of calling these functions**.
3. Mutating function **bases on their probabilities** defined in C Template.
4. Mutating the **order of each threads** to focus on **race condition cases**.



Mutator

```
for idx, arg_type_name in enumerate(defined_funcs[func_name]['args']):
    func_arg = ExtProtoArgument()
    func_arg.arg_type_name = arg_type_name
1078  v            if template.is_threads_binding:
99 1079              # this template has fixed number of thread.
99 1080              n_thread = len(template.threads_binding)
99
99 1081  v            else:
99 1082              n_thread = random.randint(*DEFAULT_NUM_THREAD_RANGE)
99 1083
99
1084  v            for _ in range(n_thread):
1085              # perform mutation for each thread
1086              self.mutator_thread(template, testcase)
1087
1088      args_init_values = template.args_init_values + recipe_args +
1089      recipe = None
```

How do I handle kernel panic? (1/2)



How do I handle kernel panic? (2/2)

```
**** [IOBluetoothHCIUserClient] [SetBTLPOffWL] -- mBluetoothFamily->mACPIMethods->SetBTLP (OFF) returned 0xE0000  
2C7 error  
☞ Motiva **** [IOBluetoothHCIUserClient] [SetBTRSWL] -- mBluetoothFamily->mACPIMethods->SetBTRS() returned 0xE00002C7 err  
or  
Panic(CPU 0, time 1686178019578498): NMIPi for spinlock acquisition timeout, spinlock: 0xffffffff853190e9b8, spin  
lock owner: 0xffffffff85311bc2c0, current_thread: 0xffffffff85311bc2c0, spinlock_owner_cpu: 0x0  
RAX: 0x00000000000000de, RBX: 0x0000000000000000, RCX: 0x0000000000000830, RDX: 0x0000000000000001  
RSP: 0xfffffff019d69d90, RBP: 0xfffffff019d69d90, RSI: 0x00000000000000de, RDI: 0x0000000000000001  
R8: 0xffffffff85311bc2c0, R9: 0xffffffff8018337a80, R10: 0x0000000000000051, R11: 0xfffffff0153b6088  
R12: 0x00000000000000d0, R13: 0xfffffff015558818, R14: 0x0000000000000001, R15: 0x00000000000000de  
RFL: 0x0000000000000006, RIP: 0xffffffff80174e12af, CS: 0x0000000000000008, SS: 0x0000000000000010  
Panicked task 0xffffffff85311a16c8: 147 threads: pid 0: kernel_task  
S Backtrace (CPU 0), panicked thread: 0xffffffff85311bc2c0, Frame : Return Address  
# 0xfffffff019d69be0 : 0xffffffff80174d735a mach_kernel : _interrupt + 0x18a  
# 0xfffffff019d69c90 : 0xffffffff80172c7d69 mach_kernel : int_from_intstack + 0x38  
- 0xfffffff019d69d90 : 0xffffffff80174e2f5a mach_kernel : _lapic_send_ipi + 0x6a  
# 0xfffffff019d69dc0 : 0xffffffff80174e5a5a mach_kernel : _cpu_interrupt + 0x5a  
k 0xfffffff019d69de0 : 0xffffffff801735f454 mach_kernel : _thread_go + 0xe4  
# 0xfffffff019d69e10 : 0xffffffff80173a95b5 mach_kernel : _waitq_wakeup64_thread + 0xf5  
# 0xfffffff019d69e60 : 0xffffffff8017be01fa mach_kernel : __ZN10IOWorkLoop19signalWorkAvailableEv + 0x3a  
# 0xfffffff019d69e80 : 0xffffffff8017be232b mach_kernel : __ZN22IOInterruptEventSource23normalInterruptOccurredEPvP  
9IOServicei + 0x7b  
# 0xfffffff019d69ea0 : 0xffffffff801a070b2e com.apple.iokit.IOPCIFamily : __ZN32IOPCIMessagedInterruptController15h  
andleInterruptEPvP9IOServicei + 0x158  
# 0xfffffff019d69ef0 : 0xffffffff80185a0cc2 com.apple.driver.AppleACPIPlatform : __ZN23AppleACPIPlatformExpert23dis  
patchGlobalInterruptEi + 0x2e  
# 0xfffffff019d69f00 : 0xffffffff80185a9e59 com.apple.driver.AppleACPIPlatform : __ZN31AppleACPICPUInterruptControl
```



Fuzzing Results

- My fuzzer is written in Python around 2500 lines of code .
- Auditing XNU kernel to create template and prepare for fuzzing and I have found a lot of crash after few days.
- Found a lot of crashes. But there are some valuable crashes: CVE-2022-32894, CVE-2023-32441 (which was patched in iOS 16.6).



Key takeaways

1. **The idea** of fuzzing is the most important factor to discover security vulnerabilities.
2. Stick with **code auditing** beside with fuzzing to improve the mutator.
3. Have a **deep knowledge about the target** before fuzzing and you will find your own bugs.



3. CVE-2022-32894

iOS 15.6.1 and iPadOS 15.6.1

Released August 17, 2022

Kernel

Available for: iPhone 6s and later, iPad Pro (all models), iPad Air 2 and later, iPad 5th generation and later, iPad mini 4 and later, and iPod touch (7th generation)

Impact: An application may be able to execute arbitrary code with kernel privileges. Apple is aware of a report that this issue may have been actively exploited.

Description: An out-of-bounds write issue was addressed with improved bounds checking.

CVE-2022-32894: an anonymous researcher

Fuzzing Results

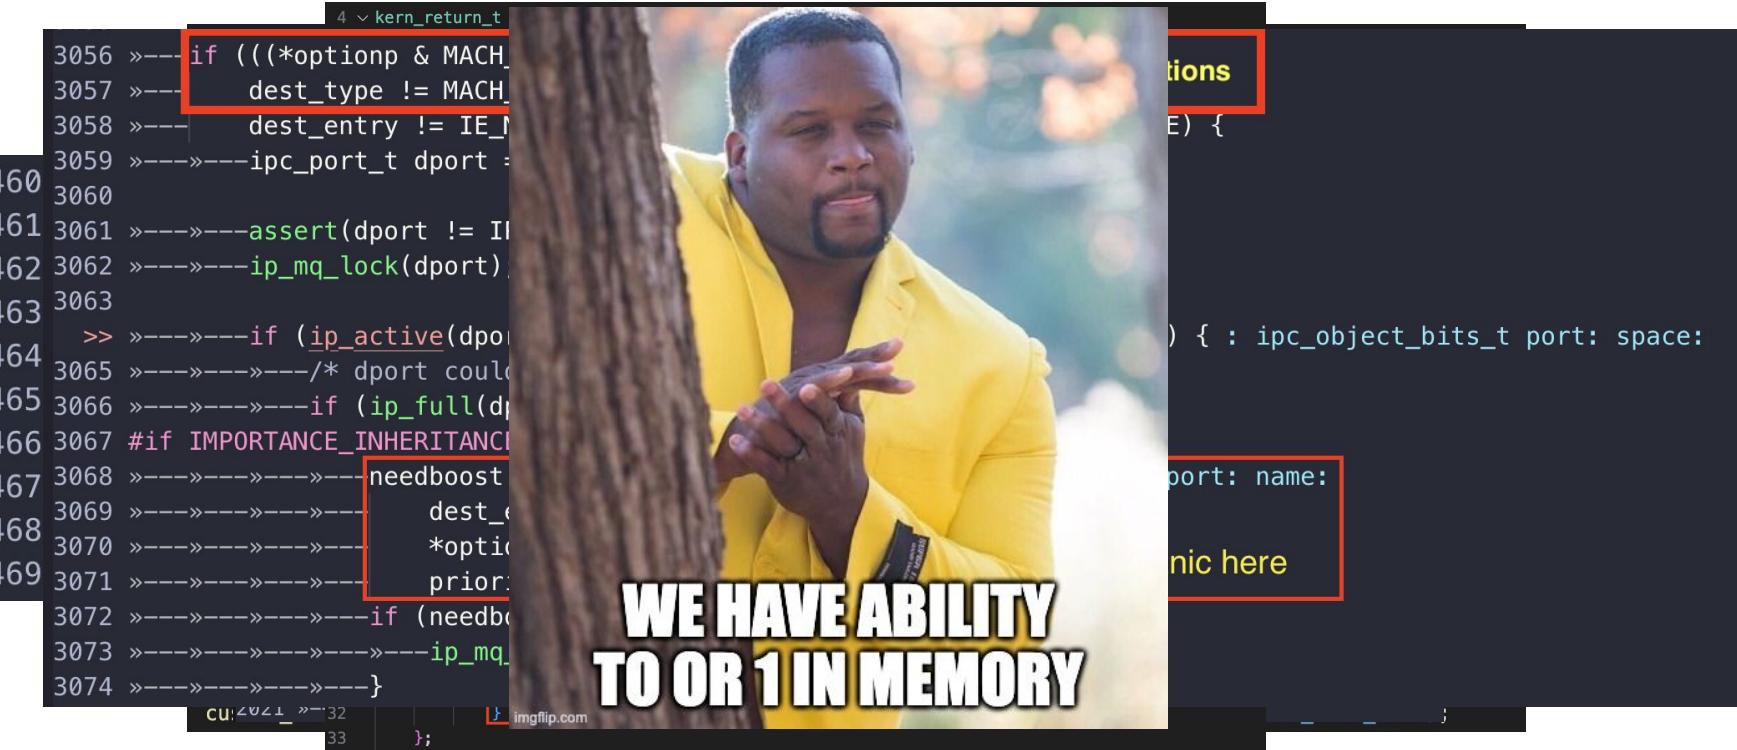
```
panic(cpu 2 caller 0xffffffff801b0b523b): ipc_port.c:461 Assertion failed: table != IPR_NULL
Panicked task 0xffffffff8542de96c8: 1 threads: pid 44699: table_null2
Backtrace (CPU 2), panicked thread: 0xffffffff85358f3c20, Frame : Return Address
0xffffffffd016dfb9e0 : 0xffffffff801a73168d mach_kernel : _handle_debugger_trap + 0x3cd
0xffffffffd016dfa30 : 0xffffffff801a8e9df7 mach_kernel : _kdp_i386_trap + 0x177
0xffffffffd016dfa70 : 0xffffffff801a8d7dae mach_kernel : _kernel_trap + 0x3de
0xffffffffd016dfa8e0 : 0xffffffff801a6c7a60 mach_kernel : trap_from_kernel + 0x26
0xffffffffd016dfb00 : 0xffffffff801a731aea mach_kernel : _DebuggerTrapWithState + 0xba
0xffffffffd016dfbc20 : 0xffffffff801a731139 mach_kernel : _panic_trap_to_debugger + 0x2e9
0xffffffffd016dfbc80 : 0xffffffff801b0b528f mach_kernel : _panic + 0x54
0xffffffffd016dfbcf0 : 0xffffffff801b0b523b mach_kernel : _panic
0xffffffffd016dfbd00 : 0xffffffff801a70e914 mach_kernel : _ipc_port_request_sparm + 0x124
0xffffffffd016dfbd20 : 0xffffffff801a708229 mach_kernel : _ipc_kmsg_copyin_from_user + 0x1f19
0xffffffffd016dfbe50 : 0xffffffff801a721a8d mach_kernel : _mach_msg_overwrite_trap + 0x2ad
0xffffffffd016dfbee0 : 0xffffffff801a8b8153 mach_kernel : _mach_call_munger64 + 0x273
0xffffffffd016dfbfa0 : 0xffffffff801a6c8246 mach_kernel : _hdl Mach_scall64 + 0x16

Process name corresponding to current thread (0xffffffff85358f3c20): table_null2
Boot args: -v keepsyms=1 tlbto_us=0 -zc zlog1=default.kalloc.768 zlog2=default.kalloc.256 vti=9 debug=0x814e kext-dev-mode=1 kcsuffix=development kdp_match_name=serial

Mac OS version:
21A559

Kernel version:
Darwin Kernel Version 21.1.0: Wed Oct 13 17:33:22 PDT 2021; root:xnu-8019.41.5~1/DEVELOPMENT_X86_64
Kernel UUID: 212C1F60-0EC2-3E01-B533-F035E8B52DD7
```

Root Cause (1/3)



Root Cause (2/3)

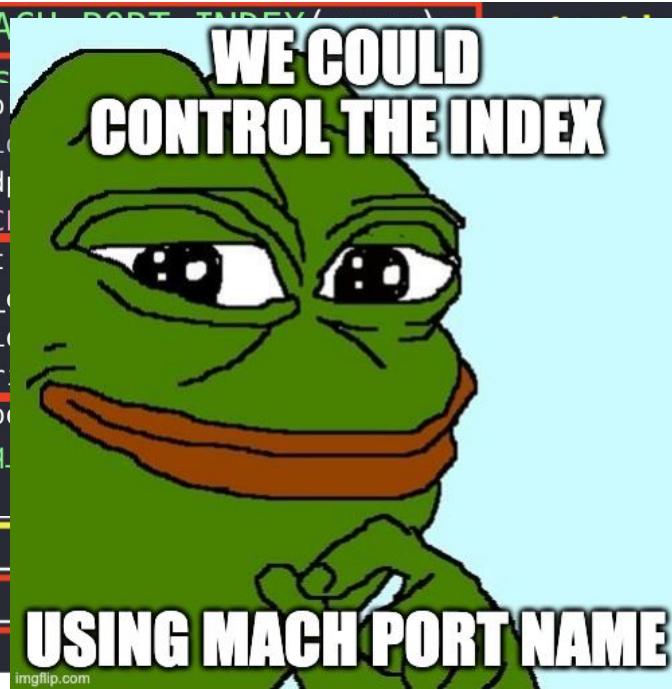
```
111 struct ipc_entry {  
112     union {  
113         struct ipc_object *XN  
30 114         struct ipc_object *XN  
30 115     };  
30 116     union {  
30 117         ipc_entry_bits_t ie_  
30 118         ipc_entry_num_t ie_  
30 119     };  
30 120     uint32_t             ie_di  
30 121     mach_port_index_t ie_in  
30 122     union {  
30 123         mach_port_index_t ie_  
124         ipc_table_index_t ie_  
125     };  
126 };
```



```
) ie_object;  
) volatile ie_volatile_object;  
:  
  
is union field with ie_next  
    ↗  
    ↗
```

Root Cause (3/3)

```
240 434 >--- index = MACH_PORT_INDEX(  
2416 > 435 >--- table = is  
  >> >--->---if (ip_active(dport  
3065 >>> >---/* dport could  
3066 >>> >---if (ip_full(dport  
3067 #if IMPORTANCE_INHERITANCE  
3068 >>> >--->--->---needboost  
3069 >>> >--->--->---dest_c  
3070 >>> >--->--->---*option  
3071 >>> >--->--->---priorit  
3072 >>> >--->--->---if (needbo  
3073 >>> >--->--->--->---ip_mq  
3074 >>> >--->--->---}  
2429 >  
2430 > 444 >---entry->ie_<br/>  
2431 > 445 >---table->ie_<br/>  
        > 446 >---space->is
```



WE COULD
CONTROL THE INDEX

index from mach port name

el)) { : ipc_object_bits_t port: space:

e, port: name:

entry->ie_request is index

ex is first free entry

Mach Port Spraying

```
71 "build" : "iPhone OS 15.6 (19G71)",  
92 "product" : "iPhone12,1",  
93 "socId" : "0x00008030",  
94 "kernel" : "Darwin Kernel Version 21.6.  
95 "incident" : "78F0E711-B5BC-4CAF-950E-2  
96 "crashReporterKey" : "d0e5ad8a18032e9bf  
97 "date" : "2023-08-12 23:34:35.15 +0800"  
98 "panicString" : "panic(cpu 5 caller 0xf  
99 ffee8a673820)\n\t x0: 0xffffffe4327b3  
0  x5: 0xffffffe8a673c60 x6: 0xfffff  
000000808 x11: 0x0000000000000000\n\t x  
16: 0xffffffe8ab98000 x17: 0xffffffe8  
100 0a907 x22: 0xffffffe4318c0a18 x23: 0x  
101 0000000000000000\n\t x28: 0xffffffe8ab9  
c0 cpsr: 0x60400204 esr: 0x9600  
OS version: 19G71\nKernel version: Darw  
el UUID: DFEEB758-B116-3C50-8913-51E34A  
0000000a83c000\nKernel text base: 0xf  
a5 0x000bec5a\nSleep : 0x00000000 0  
: 0xffffffe130d48000 - 0xffffffe730d480  
000\n . GEN0 : 0xffffffe264078000 - 0  
000\n . GEN1 : 0xffffffe430d40000 - 0  
000\n . GEN2 : 0xffffffe430d40000 - 0  
000
```



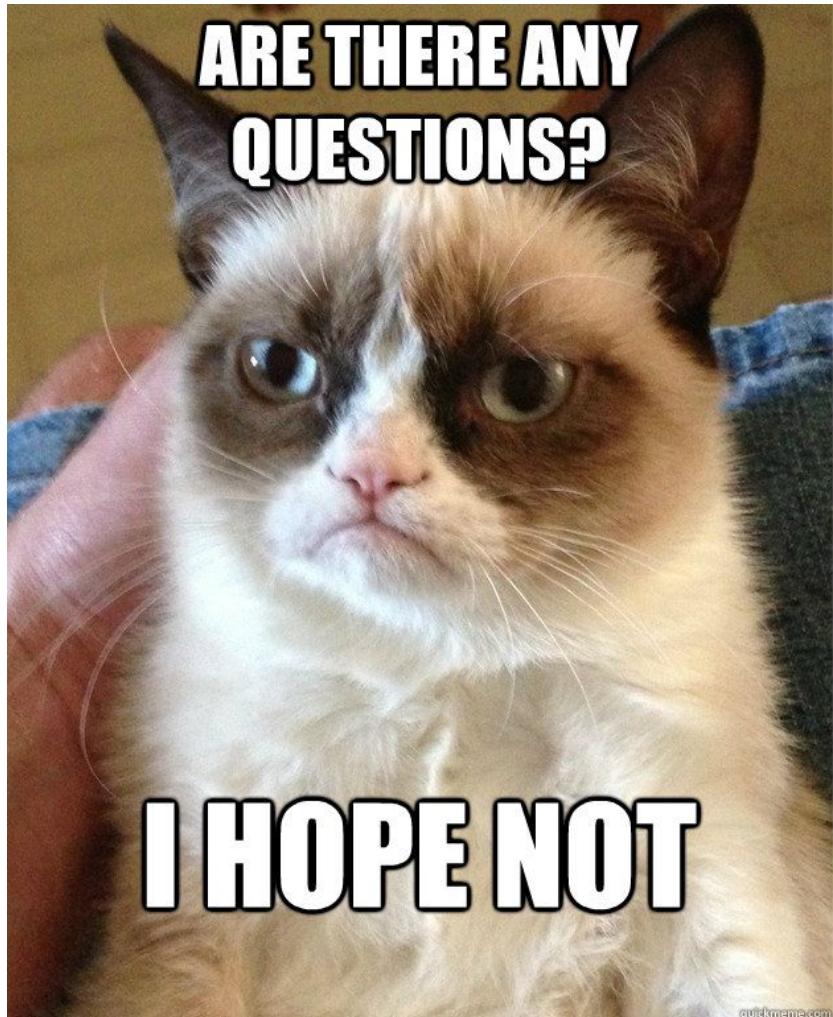
```
0.41~4~/RELEASE_ARM64_T8030",  
ff01238a8c0, lr 0x8ec1e2701238318c (saved state: 0xffff  
41 x3: 0x0000000000000000b1\n\t x4: 0x0000000000000000  
fffffe43042ad30 x9: 0x0000000000000000 x10: 0x00000000  
x14: 0xffffffe130e45540 x15: 0x000000000003e18\n\t x  
18c0a18\n\t x20: 0xffffffe4327b3690 x21: 0x000000000000  
0x0000000000000000 x26: 0xffffffe4318c0a88 x27: 0x00  
18c sp: 0xffffffe8a673b70\n\t pc: 0xfffffff01238a8  
message: panic\nMemory ID: 0x6\nOS release type: User\n  
2022; root:xnu-8020.140.41~4~/RELEASE_ARM64_T8030\nKern  
boot?: YES\nPaniclog version: 13\nKernel slide: 0x0  
Epoch Time: sec usec\nBoot : 0x64d7a1  
calendar: 0x64d7a671 0x0009fe79\n\nZone info:\nZone map  
ac000\n.R0 : 0xffffffe2173ac000 - 0xffffffe264078  
- 0xffffffe430d40000\n.GEN2 : 0xffffffe430d40000 -
```

Key takeaways

1. Mach port is very complex and it use a lot of “**union**” in their **structures** which could lead to some **side effects**.
2. With ability to generate test-case from custom template, my fuzzer able to explore some **potential code paths** which **haven't reach before**.



Q&A



References

- [Mach Overview](#)
- [CVE-2020-27932.](#)
- <https://www.fortinet.com/blog/threat-research/inspecting-mach-messages-in-macos-kernel-mode--part-i-sniffing->
- <https://medium.com/@ali.pourhadi/ipc-mach-message-cab64ff1b569>
- [MOSEC 2021: Exploitation of XNU Port Type Confusions](#)
- <https://googleprojectzero.blogspot.com/2022/04/cve-2021-1782-ios-in-wild-vulnerability.html>
- [Understanding Mach IPC from MOSEC 2022](#)
- [CVE-2021-1870 Analysis](#)
- [https://github.com/kholia/OSX-KVM.](https://github.com/kholia/OSX-KVM)

THANK YOU !